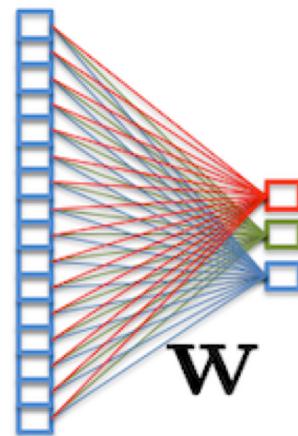
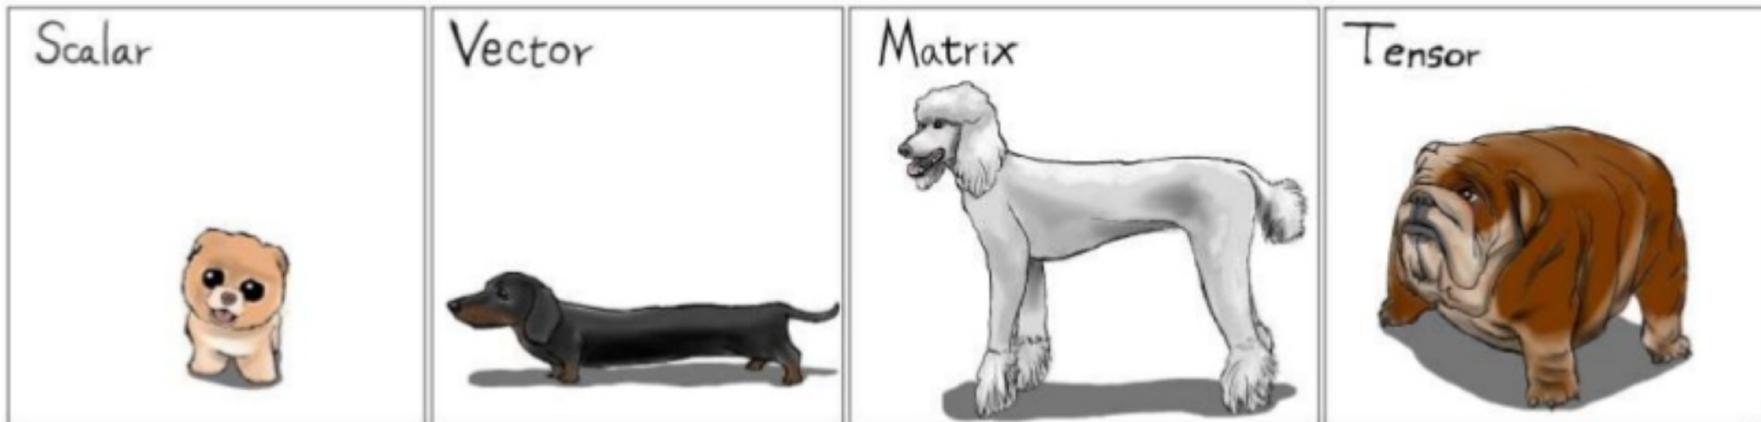


# *5IF - Deep Learning and Differentiable Programming*

## 2.1 Frameworks and tensors

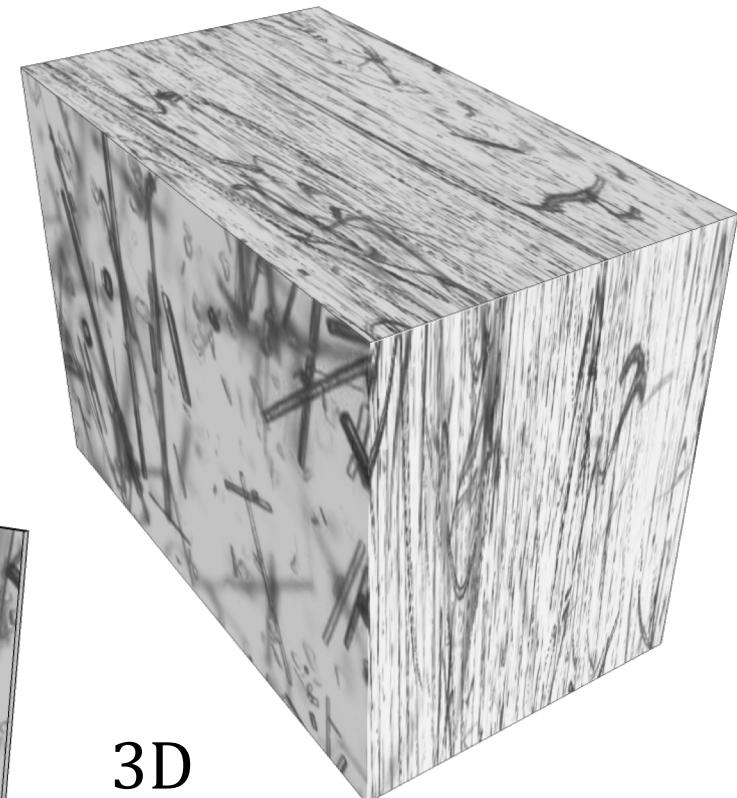
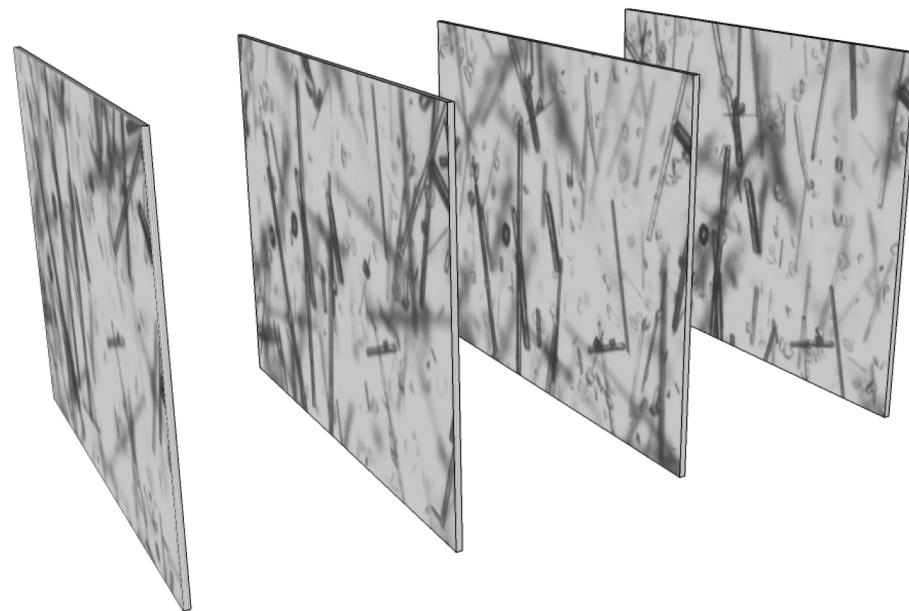
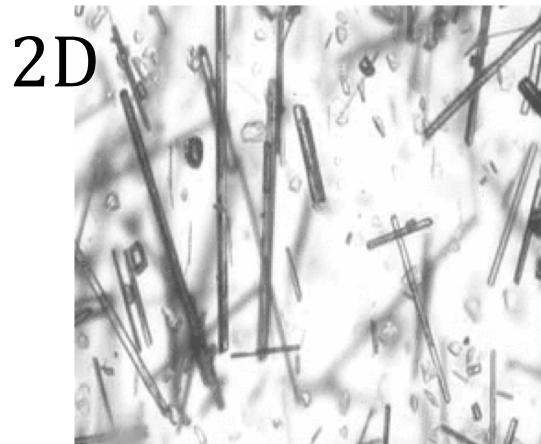


# We manipulate tensors



[Figure: Anima Anandkumar]

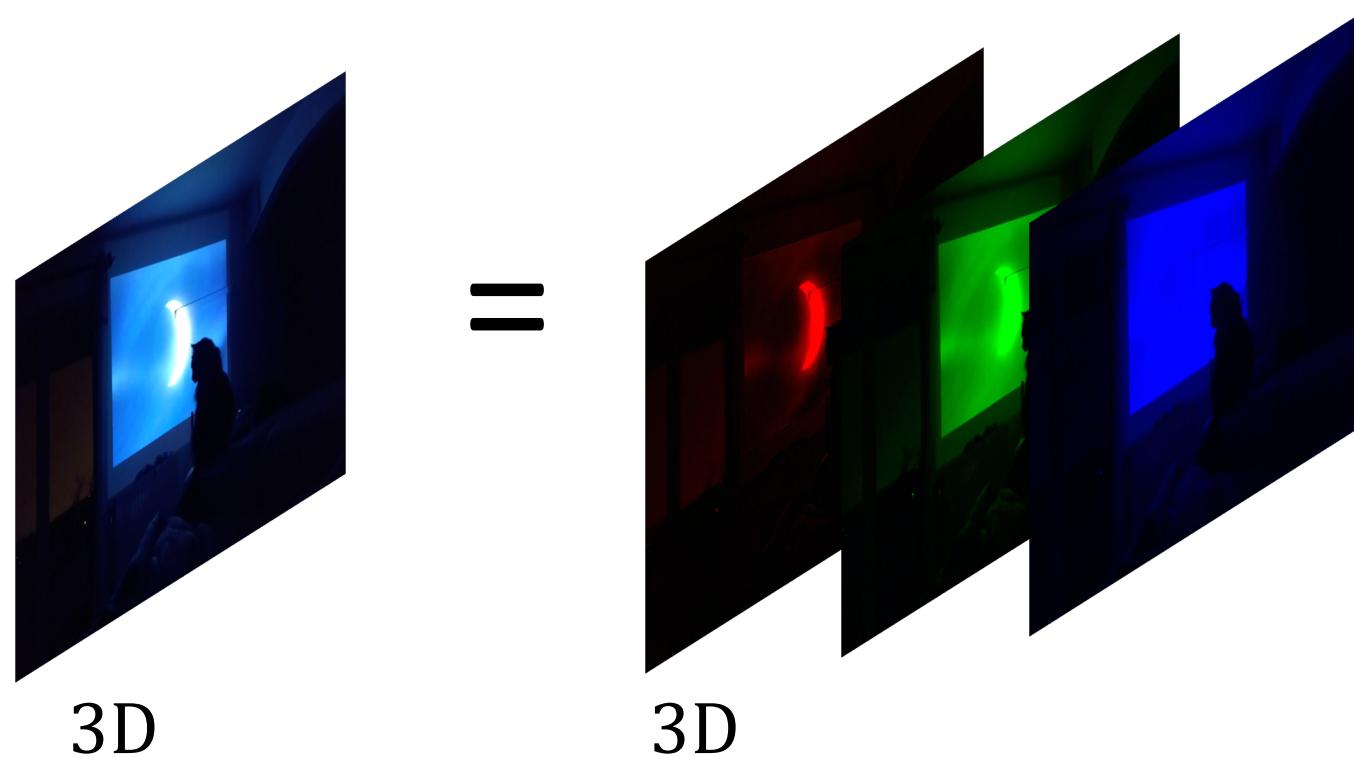
# High dimensional tensors



[Data: LAGEP]

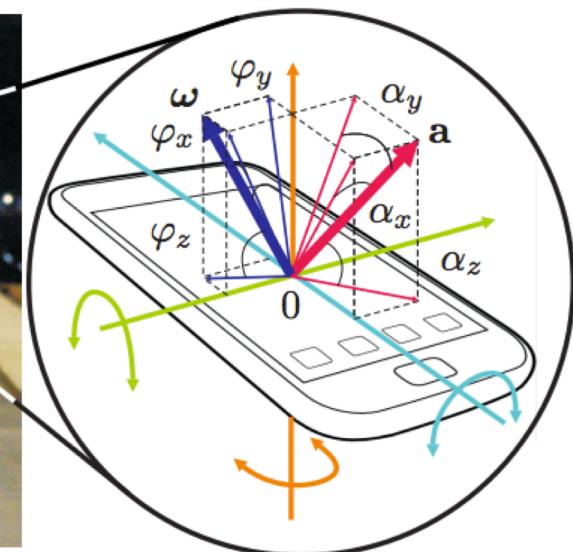
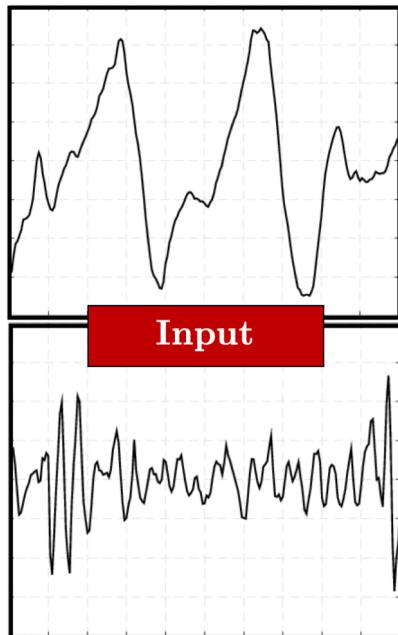
# Images as tensors

A color image has 3 color channels (red, green, blue) and is therefore a 3D tensor.



# (Multiple) 1D signals

- 1500 volunteers, 1500 Nexus 5 smartphones
- Several months of natural daily usage, 27.6 TB of data
- Multiple sensors: camera, touchscreen, GPS, bluetooth, wifi, cell antenna, inertial, magnetometer
- This work: inertial sensors only, recorded at 200Hz



:

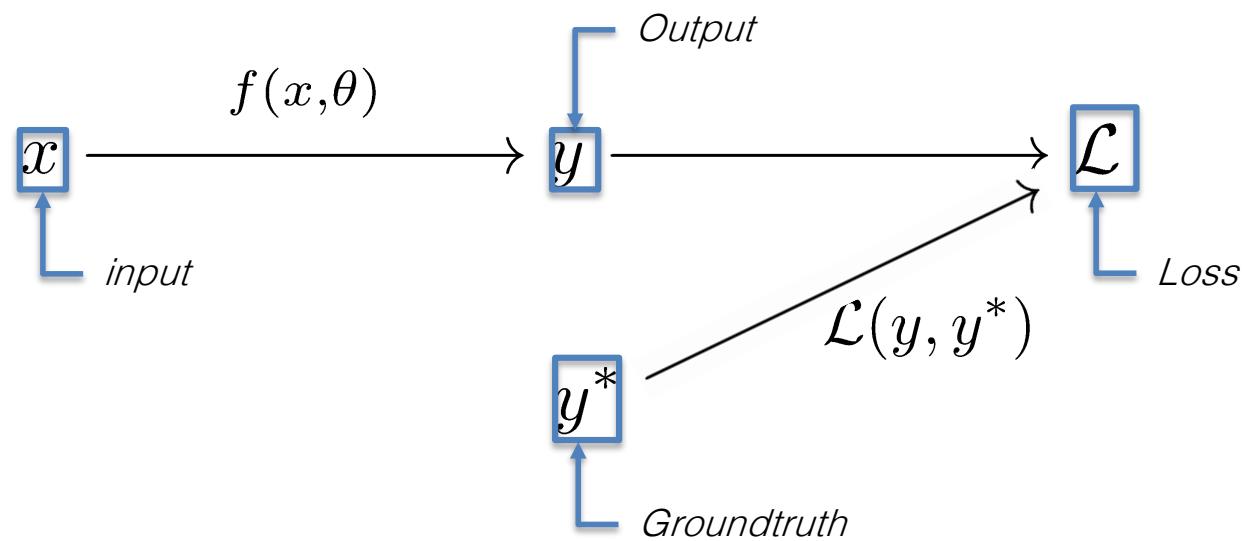
Google

[Neverova, Wolf, Lacey, Fridmann, Chandra, Barbelli, Taylor, *IEEE Access 2016*]

# Tensors: examples

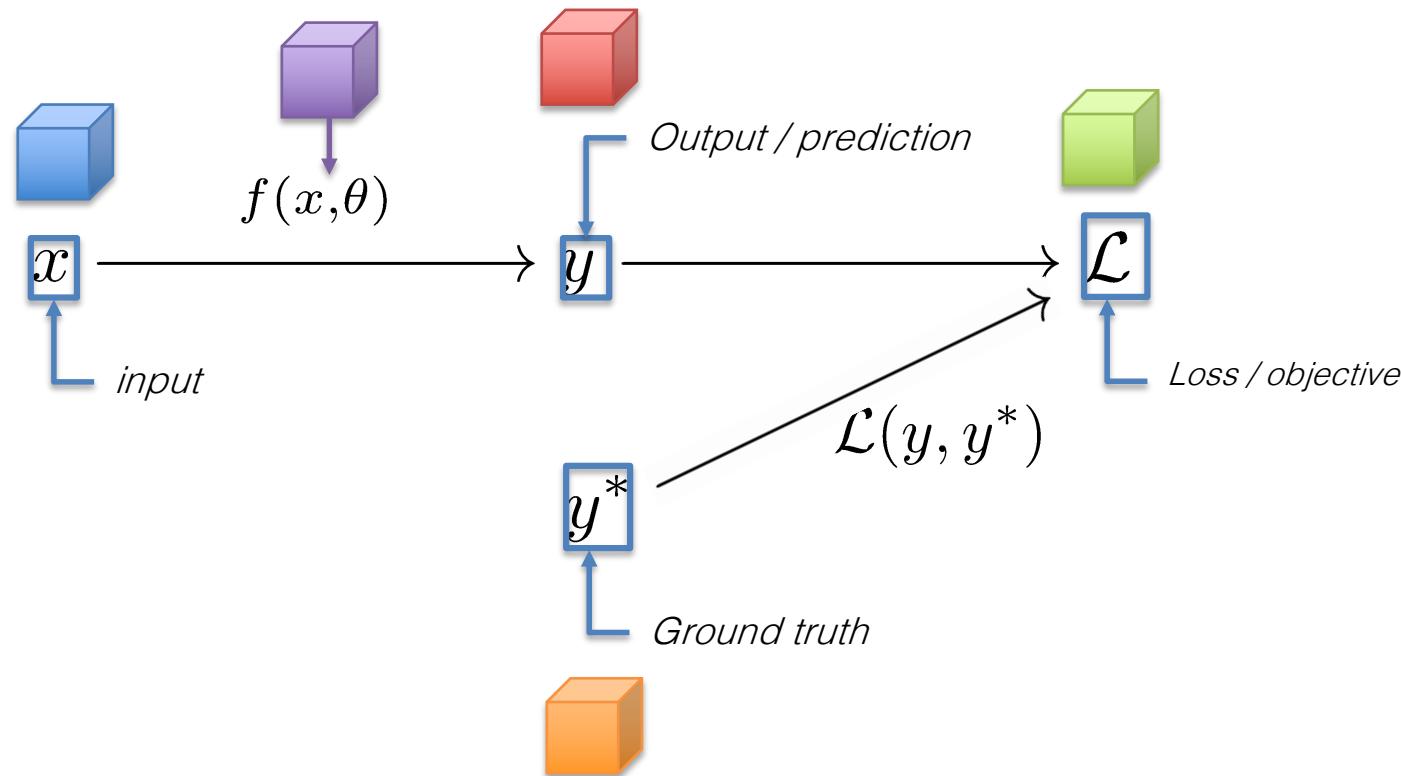
- Example of a tensor of dim 2 (input data, 1D signal)
  - Batch dimension (multiple samples)
  - Signal dimension
- Example of a tensor of dim 2 (output data, classification)
  - Batch dimension (multiple samples)
  - Prediction for different classes
- Example of a tensor of dim 3 (layer activation, 1D signal)
  - Batch dimension (multiple samples)
  - Signal dimension
  - Feature dimension
- Example of a tensor of dim 4 (layer activation, 2D image)
  - Batch dimension (multiple samples)
  - Spatial X dimension
  - Spatial Y dimension
  - Feature dimension
- Example of a tensor of dim 5 (input data, 2D+t video)
  - Batch dimension (multiple samples)
  - Spatial X dimension
  - Spatial Y dimension
  - Color channel dimension
  - Time dimension

# *Functional mappings*



# How do we code all this?

Inputs, outputs, layer activations, weights are tensors of different dimensions.



# Deep Learning Frameworks

Before 2013

2002



The wild west. Models are handcoded in Matlab, C++ by a small minority of people doing deep learning.

2010

theano

Initial release of [Torch](#) by IDIAP Research Institute, later Facebook AI Research. Interfaces: Lua, C++. It uses Autograd (automatic dynamic differentiation) [Replaced by PyTorch.]

2013

Caffe

[Theano](#) introduced by Montreal Institute for Learning Algorithms (MILA), University of Montreal. Autograd. [Discontinued]

[CAFFE](#) introduced by University of California, Berkeley. Interfaces: C++ and shell+text files.

2015



[Tensorflow](#) introduced by Google, with immediate success. Interfaces: Python, C++ (less supported). Uses a static calculation graph, not Autograd.

2015



[MXNet](#) (Appache with researchers from CMU, NYU, NUS, MIT)

# Deep Learning Frameworks

2015



[Keras](#) is a meta-language developed by Google Engineer François Chollet, simplifying coding, initially developed to run on top of Theano. Now also runs on Tensorflow, CNTK. Part of Tensorflow now.

2016



[PyTorch](#) introduced by Facebook, with immediate success. Interaces: Python, C++ (less supported). Uses a static calculation graph, not Autograd.

2016



[CNTK](#) by Microsoft

2017



[Tensorflow](#) introduces eager mode, a dynamic graph calculation mode based on Autograd to respond to the success of PyTorch & Co.

2019

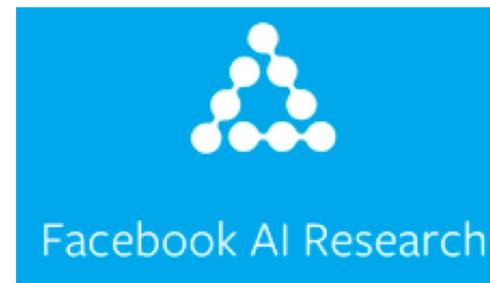


[Tensorflow](#) makes the eager mode the default mode in version 2.0.

# Main frameworks



Google

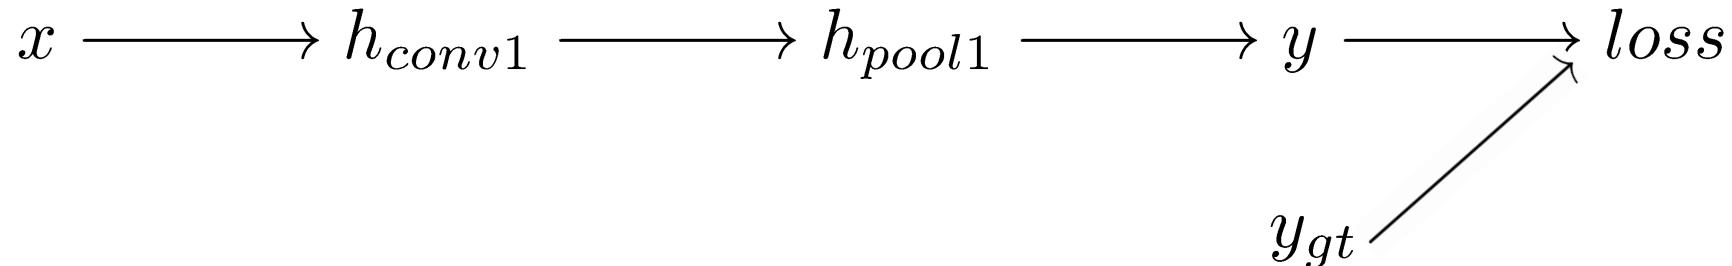


Tensorflow

PyTorch

- Both support execution and training on CPUs, GPUs, TPUs (google's machine learning hardware)
- Both use python.
- Tensorflow also supports C++, Swift.

# Tensorflow (Google) : example



Definition of a symbolic graph:

```
1 # The input tensors
2 x = tf.placeholder(tf.float32, [None, 224, 224, 1])
3
4 # The target labels
5 y_gt = tf.placeholder(tf.float32, [None, NO_CLASSES])
6
7 # Layer 1 : convolutional
8 W_conv1 = weight_variable([5, 5, 1, 32])
9 b_conv1 = bias_variable([32])
10 h_conv1 = tf.nn.relu(tf.nn.conv2d(x, W_conv1, strides=[1, 1, 1, 1], padding='SAME') + b_conv1)
11 h_pool1 = tf.nn.max_pool(h_conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
12
13 # Layer 2 (output) : fully connected
14 W_fc2 = weight_variable([N, NO_CLASSES])
15 b_fc2 = bias_variable([NO_CLASSES])
16 y = tf.matmul(h_pool1, W_fc2) + b_fc2
17
18 # The loss function
19 loss = tf.reduce_sum(tf.nn.softmax_cross_entropy_with_logits(y, y_gt))
```

Automatic calculation of the gradients:

```
22 # Define gradients of loss with respect to some weights
23 grads_input = tf.gradients(loss, W_conv1)[0]
24
25 # Alternatively, define gradients of loss with respect to the input
26 grads_input = tf.gradients(loss, x)[0]
```

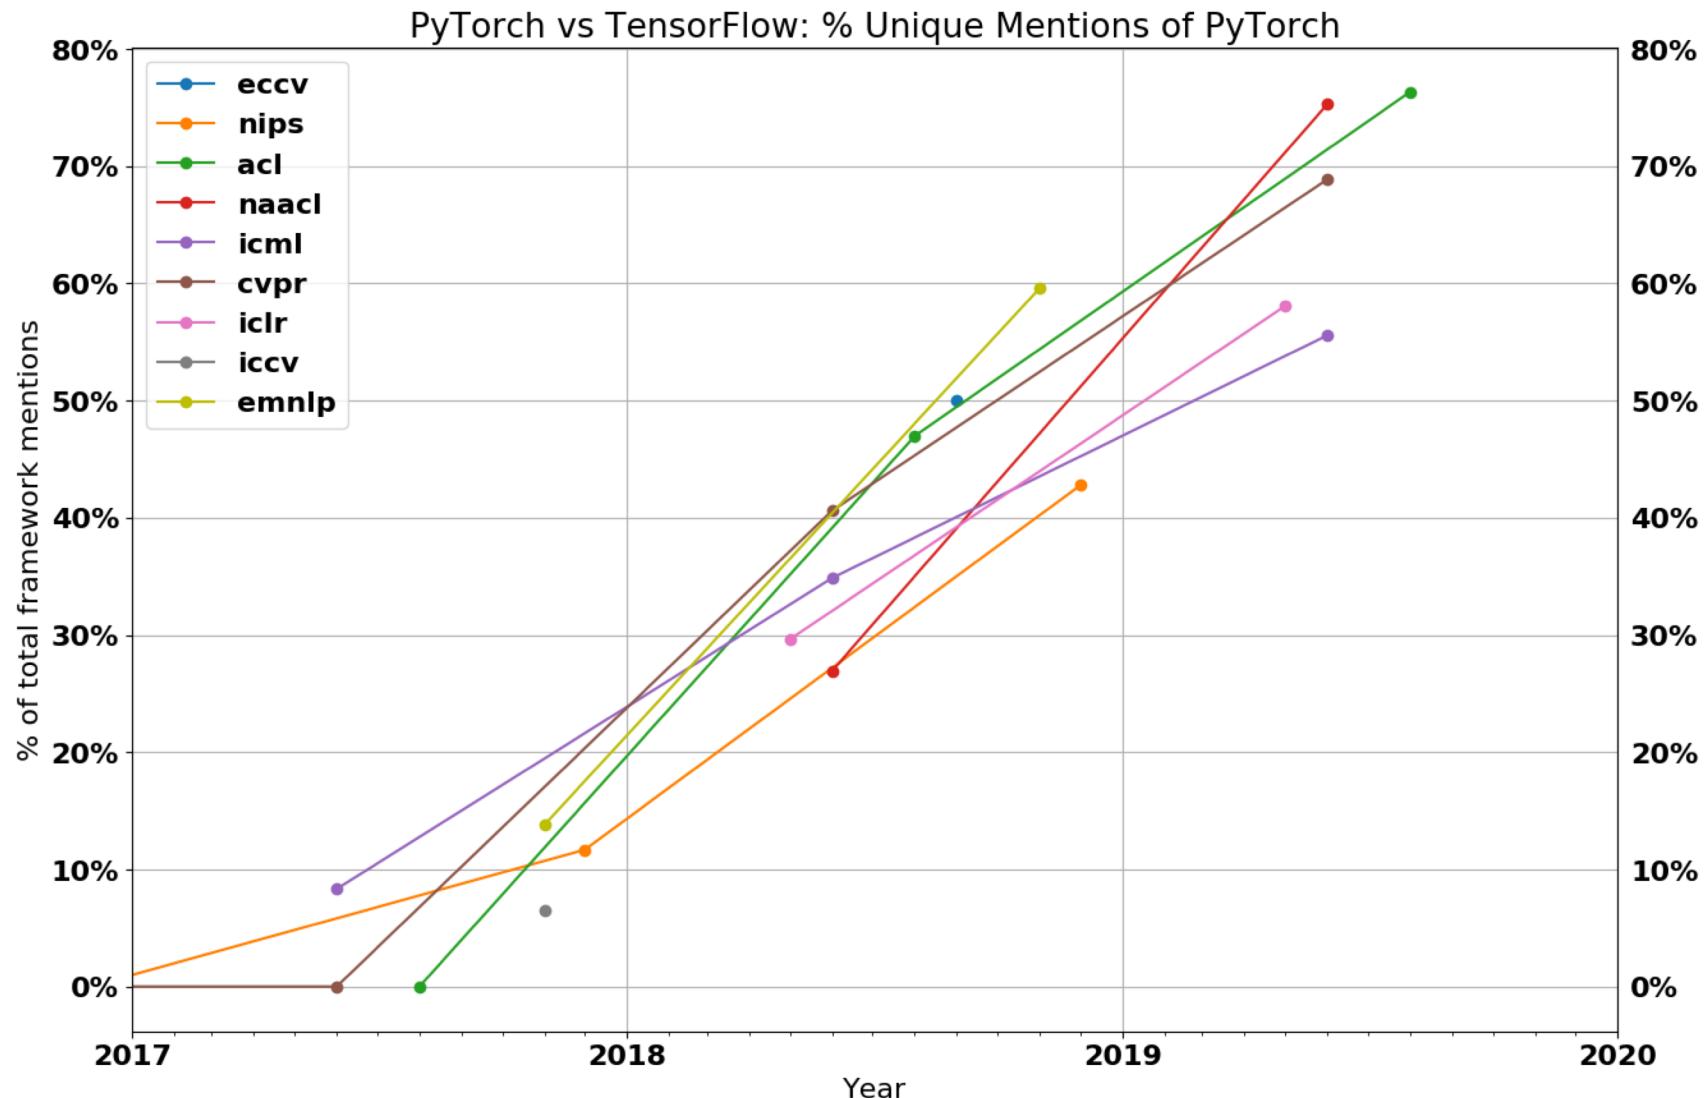
# PyTorch (Facebook) : example

Imperative calculations like in standard python. Easy debugging.

```
1 ▼ class Net(nn.Module):
2
3 ▼   . . . def __init__(self):
4       . . . . super(Net, self).__init__()
5
6       . . . . # Set up my layers
7       . . . . self.fc1 = nn.Linear(7*7*64, 1024)
8       . . . . self.fc2 = nn.Linear(1024, 10)
9
10      . . . . # Set up loss function and optimizer
11      . . . . crossentropy = nn.CrossEntropyLoss()
12      . . . . optimizer = optim.SGD(net.parameters(), lr=0.001, mo
13
14 ▼   . . . def train(self, xinput, ylabel):
15
16       . . . . x = self.fc1(xinput)
17       . . . . x = F.relu(x)
18       . . . . youtput = self.fc2(x)
19
20       . . . . loss = crossentropy(youtput, ylabel)
21
22       . . . . print "Debug:", xinput, youtput, x
23
24       . . . . loss.backward() ←
25
26       . . . . optimizer.step()
```

The gradients are calculated by backtracking through the dynamically created execution graph.

# PyTorch vs. Tensorflow



# PyTorch

## Tensor cheat sheet

# PyTorch Live Install Party



<https://pytorch.org/get-started/locally/>

# Installing PyTorch

Create a virtual environment where we are "safe". Use Python 3.7.

```
1 cd  
2 virtualenv --system-site-packages -p /usr/local/bin/  
    python3.7 pytorch1.2
```

Activate the environment:

```
1 source ~/pytorch1.2/bin/activate
```

Now all installs will be done in the virtual env.

Install PyTorch itself and the vision package

```
1 pip3 install torch torchvision  
2 pip3 install scikit-image
```

# Creating tensors

```
1 # loading PyTorch
2 import torch
3
4 # create with given shape
5 torch.full((shape), value)
6 torch.full_like(other_tensor, value)
7
8 # create with given values
9 torch.tensor((values))
10 torch.tensor((values), dtype=torch.int16)
11
12 # create from numpy array
13 torch.from_numpy(numpyArray)
14
15 # Create zeros or ones
16 torch.zeros((shape))
17 torch.zeros_like(other_tensor)
18 torch.ones((shape))
19 torch.ones_like(other_tensor)
```

# Creating tensors, tensor I/O

```
1 # Random tensors
2 torch.randn(3, 4)
3
4 # Tensor I/O
5 A = torch.load ("A.tensor")
6 torch.save (A, "A.tensor")
```

# Tensor slicing

Slicing is similar to python (NumPy) Slicing or Matlab notation.

Example for a 2D tensor:

```
1 A[1,5]                      # access an element (row, col)
2 A[:,5]                       # column access
3 A[1,:]                       # row access
4
5 A[1:6,:]     A[1:,: ]       # range access (1:6 = 1,2,3,4,5)
6 A[:,0:-1]                   # Negative index count backwards
7                           # -1 = last col/row
8
9 A==3                         # provides a tensor of logical
10                          # results
11
12 A[4:17,3] = B               # replace a slice
13
14 A[B==3]=4                   # Set values in A to 4 at pos
15                           # where there is a 3 in B
```

# Manipulating tensors

```
1 # concatenate tensors
2 torch.cat((tensors), axis)
3
4 # split tensors into chunks of equal size
5 torch.split(tensor, splitSize, dim=0)
6
7 # reshape tensor w/o changing the data
8 torch.view(tensor, shape)
9
10 # Repeat along a given dimension
11 X.repeat(4,2)
12
13 # transpose tensor
14 torch.t(tensor) # 1D and 2D tensors
15 torch.transpose(tensor, dim0, dim1)
16
17 # Sorting
18 torch.sort(input, dim=-1)
```

# Tensor math

```
1 # Overloaded operators
2 x = A+Y*Z-B                      # * is elementwise mul
3
4 # Sum, product, min, max of all elements
5 torch.sum(tensor)      torch.min(tensor)
6 torch.prod(tensor)     torch.max(tensor)
7
8 # Linear algebra
9 torch.mm(A, B)                  # Matrix multiplication
10 torch.inverse(tensor)    # Matrix inversion
11 torch.det(tensor)        # Determinant
```

# Elementwise operations

```
1 torch.exp(tensor)      torch.log(tensor)
2 torch.cos(tensor)      torch.cosh(tensor)
3 torch.sin(tensor)      torch.sinh(tensor)
4 torch.tan(tensor)      torch.tanh(tensor)
5
6 torch.add(tensor, tensor2) # or tensor+scalar
7 torch.div(tensor, tensor2) # or tensor/scalar
8 torch.mult(tensor, tensor2) # or tensor*scalar
9 torch.sub(tensor, tensor2) # or tensor-scalar
```

# Broadcasting

If a PyTorch operation supports broadcast, then its Tensor arguments can be automatically expanded to be of equal sizes (without making copies of the data).

```
1 x=torch.empty(5,1,4,1)
2 y=torch.empty( 3,1,1)
3 (x+y).size()
```

```
1 torch.Size([5, 3, 4, 1])
```

```
1 x=torch.empty(5,2,4,1)
2 y=torch.empty(3,1,1)
3 (x+y).size()
```

```
1 RuntimeError: The size of tensor a (2) must match the size
   of tensor b (3) at non-singleton dimension 1
```

<https://pytorch.org/docs/stable/notes/broadcasting.html>

# Broadcasting: rules

Two tensors are “broadcastable” if the following rules hold:

- Each tensor has at least one dimension.
- When iterating over the dimension sizes, starting at the trailing dimension, the dimension sizes must either be equal, one of them is 1, or one of them does not exist.

If two tensors  $x, y$  are “broadcastable”, the resulting tensor size is calculated as follows:

- If the number of dimensions of  $x$  and  $y$  are not equal, prepend 1 to the dimensions of the tensor with fewer dimensions to make them equal length.
- Then, for each dimension size, the resulting dimension size is the max of the sizes of  $x$  and  $y$  along that dimension.

<https://pytorch.org/docs/stable/notes/broadcasting.html>

# Images as tensors

```
1 pip install scikit-image
```

```
1 import torch
2 from skimage import io
3
4 image = io.imread("blue_tigrou.jpg")
5
6 io.imsave("tigre_saved.jpg", image)
7
8 G=image[:, :, 0]
9 G=image[:, :, 1]
10 B=image[:, :, 2]
11 io.imsave("tigre_r.jpg", R)
12 io.imsave("tigre_g.jpg", G)
13 io.imsave("tigre_b.jpg", B)
```



# Tensors from .csv files

```
1 from numpy import genfromtxt  
2  
3 # Import the text file into a numpy array  
4 n = genfromtxt('file.csv', delimiter=';')  
5  
6 # Convert to torch tensor  
7 D = torch.tensor(n, dtype=torch.float32)
```