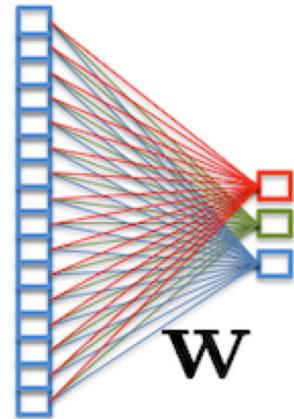


5IF - Deep Learning et Programmation Différentielle

2.6 Stochastic Gradient Descent



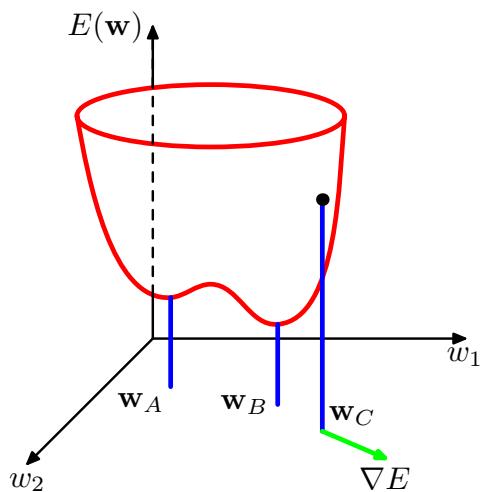
Christian Wolf

Learning by gradient descent

Iterative minimisation through **gradient descent**:

$$\theta^{[t+1]} = \theta^{[t]} + \nu \nabla \mathcal{L}(h(x, \theta), y^*)$$

↑
Learning rate



Can be blocked in a local minimum (not that it matters much ...)

[Figure: C. Bishop, 2006]

Stochastic Gradient Descent

The vanilla version of SGD:

$$\theta^{[t+1]} = \theta^{[t]} - \nu \nabla$$

where ν is the learning rate (a hyper-parameter).

The learning rate has a big impact on convergence and convergence speed:

- Low ν : slow convergence
- High ν : overshoot target

⇒ decay learning rate during learning, e.g. divide by two every X epochs.

The following slides are partially inspired by

SGD with Momentum

Momentum tends to maintains gradient direction between updates:

$$\mathbf{v}^{[t+1]} = \mu \mathbf{v}^{[t]} - \nu \nabla$$

$$\theta^{[t+1]} = \theta^{[t]} + \mathbf{v}^{[t+1]}$$

where μ is a new hyper-parameter.

Typical values: $\mu = 0.5, 0.9, 0.95, 0.99$.

Nesterov's Accelerated Momentum

Nesterov's Accelerated Momentum calculates the gradient ∇ at the position $\tilde{\theta}^{[t+1]}$ at which momentum alone would have brought it:

$$\tilde{\theta}^{[t+1]} = \theta^{[t+1]} + \mu \mathbf{v}^{[t]}$$

$$\mathbf{v}^{[t+1]} = \mu \mathbf{v}^{[t]} - \nu \nabla \tilde{\theta}^{[t+1]}$$

$$\theta^{[t+1]} = \theta^{[t]} + \mathbf{v}^{[t+1]}$$

Y. Nesterov. A method of solving a convex programming problem with convergence rate $O(1/\text{sqr}(k))$. Soviet Mathematics Doklady, 1983

Adaptive learning rates: Adagrad

Adagrad keeps a variable vector \mathbf{c} holding sums of squared derivatives, per gradient element:

$$\mathbf{c}^{[t+1]} = \mathbf{c}^{[t]} + \nabla^2$$

$$\theta^{[t+1]} = \theta^{[t]} - \nu \frac{\nabla}{\sqrt{\mathbf{c}^{[t+1]}} + \epsilon}$$

where ϵ is small and ensures numerical stability.

Effect: a large gradient value will lead to lower effective learning rate for a given parameter.

J. Duchi, E. Hazan, Y. Singer, Adaptive Subgradient Methods for Online Learning and Stochastic Optimization, JMLR, 2011.

Adaptive learning rates: RMSProp

RMSProp keeps a running average instead of accumulated gradients:

$$\mathbf{c}^{[t+1]} = \beta \mathbf{c}^{[t]} + (1 - \beta) \nabla^2$$

$$\theta^{[t+1]} = \theta^{[t]} - \nu \frac{\nabla}{\sqrt{\mathbf{c}^{[t+1]}} + \epsilon}$$

G. Hinton, unpublished.

Adaptive learning rates: ADAM

The **ADAM update rule** is similar to RMSProp, but smoothes the momentum term:

$$\mathbf{m}^{[t+1]} = \beta_1 * \mathbf{m}^{[t]} + (1 - \beta_1) \nabla$$

$$\mathbf{v}^{[t+1]} = \beta_2 * \mathbf{v}^{[t]} + (1 - \beta_2) \nabla^2$$

$$\theta^{[t+1]} = \theta^{[t]} - \nu \frac{\mathbf{m}^{[t+1]}}{\sqrt{\mathbf{v}^{[t+1]}} + \epsilon}$$

Typical values of the hyper-parameters:

$$\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 1e-08$$

D.P. Kingma, J. Ba, Adam: A method for stochastic optimization. Machine Learning, 2014

Adaptive learning rates: ADAM

The **ADAM update rule** with bias correction decreases the effect of initialization to zero (bias):

$$\tilde{\mathbf{m}}^{[t+1]} = \beta_1 * \mathbf{m}^{[t]} + (1 - \beta_1) \nabla$$

$$\mathbf{m}^{[t+1]} = \frac{\tilde{\mathbf{m}}^{[t+1]}}{1 - \beta_1^t}$$

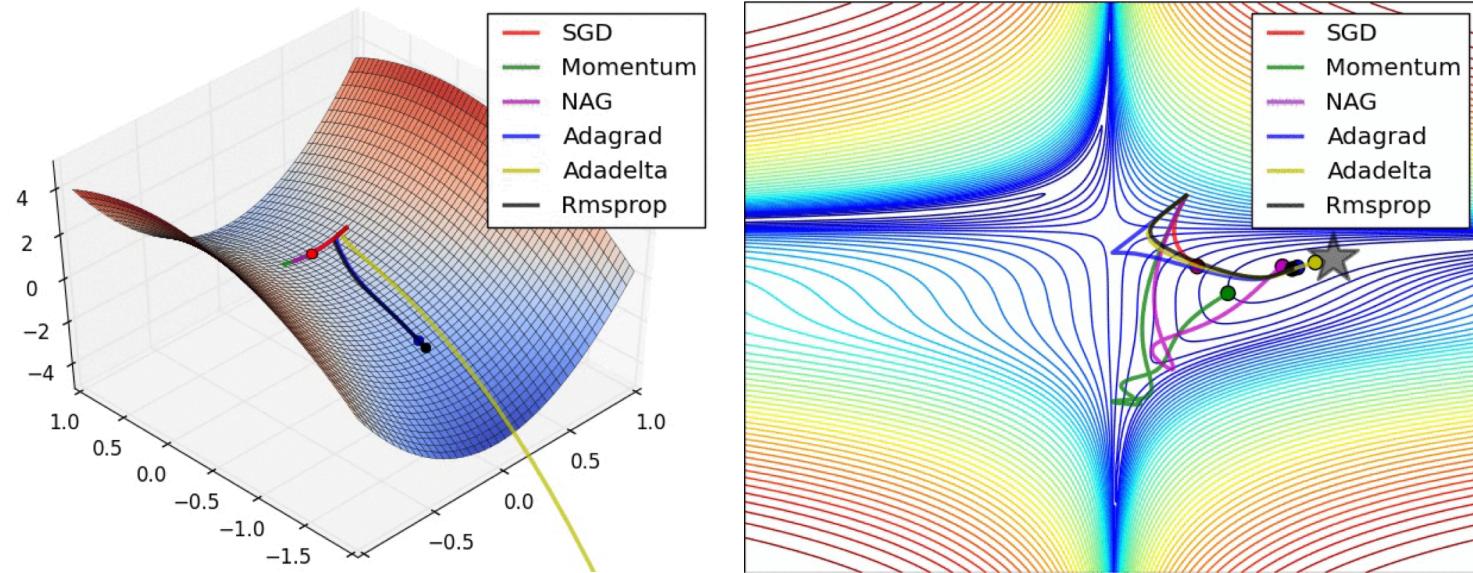
$$\tilde{\mathbf{v}}^{[t+1]} = \beta_2 * \mathbf{v}^{[t]} + (1 - \beta_2) \nabla^2$$

$$\mathbf{v}^{[t+1]} = \frac{\tilde{\mathbf{v}}^{[t+1]}}{1 - \beta_2^t}$$

$$\theta^{[t+1]} = \theta^{[t]} - \nu \frac{\mathbf{m}^{[t+1]}}{\sqrt{\mathbf{C}^{[t+1]}} + \epsilon}$$

Remark: $x^{[t]}$ indexes iteration t ; x^t denotes x to the power of t .

Visualization



[Animations: Alex Radford, Open-AI]

<http://cs231n.github.io/neural-networks-3/>

Learning rates

If you are unsure, use ADAM but also try SGD.

Even the adaptive methods use global learning rates, which need to be set.

Recall:

- Low ν : slow convergence
- High ν : overshoot target

⇒ decay learning rate during learning, e.g. divide by two every X epochs.

Experiments

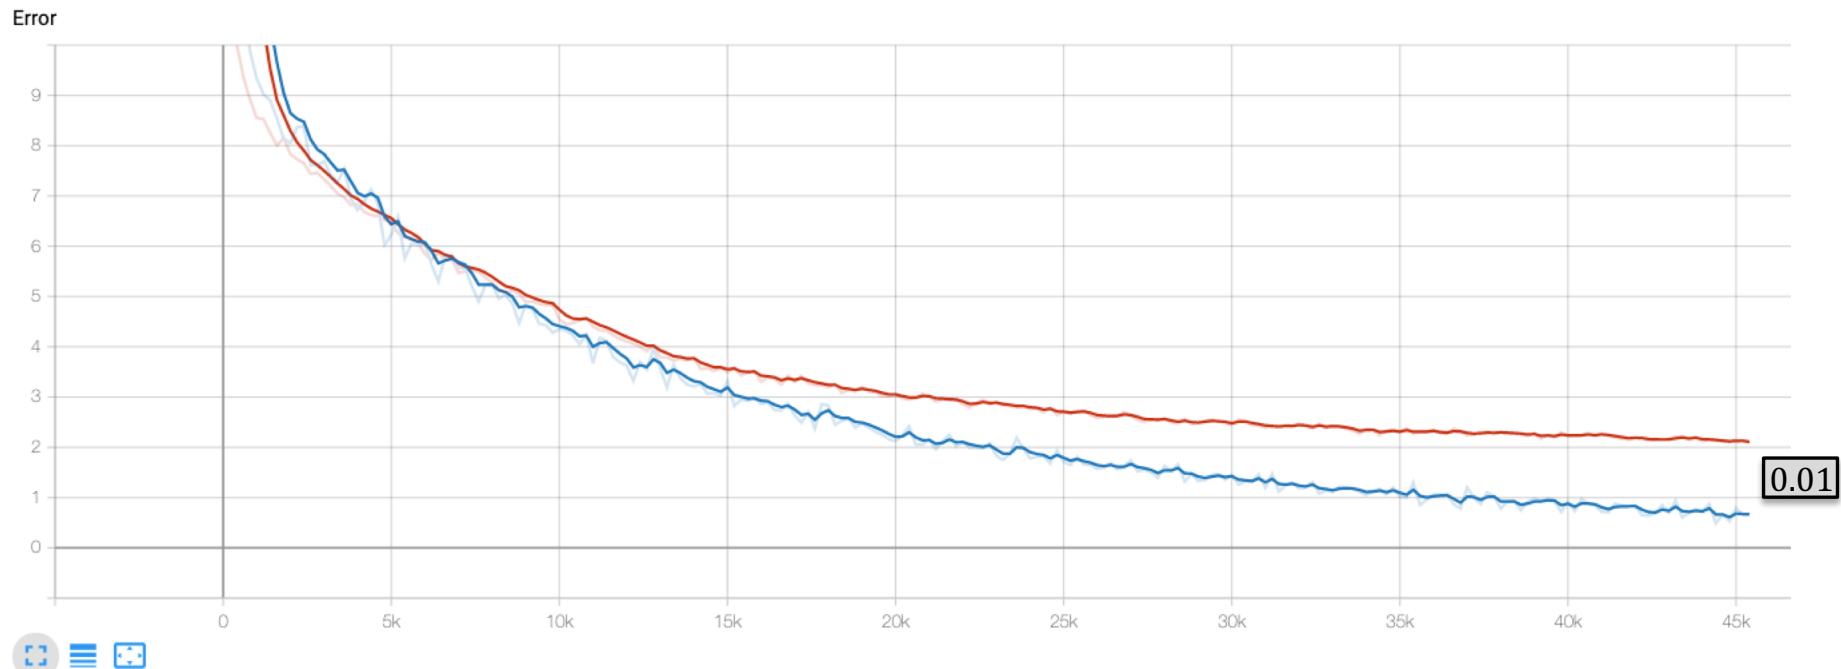
We will reuse our 2 layer MLP with 2000 hidden units and ReLU activation and optimize it with SGD and different learning rates on MNIST:

```
1 class MLP(torch.nn.Module):
2     def __init__(self):
3         super(MLP, self).__init__()
4         self.fc1 = torch.nn.Linear(28*28, 200)
5         self.fc2 = torch.nn.Linear(200, 10)
6
7     def forward(self, x):
8         x = x.view(-1, 28*28)
9         x = F.relu(self.fc1(x))
10    return self.fc2(x)
11
12 model = MLP()
13 crossentropy = torch.nn.CrossEntropyLoss()
14 optimizer = torch.optim.SGD(model.parameters(),
15     lr=0.01)
```

Learning rate

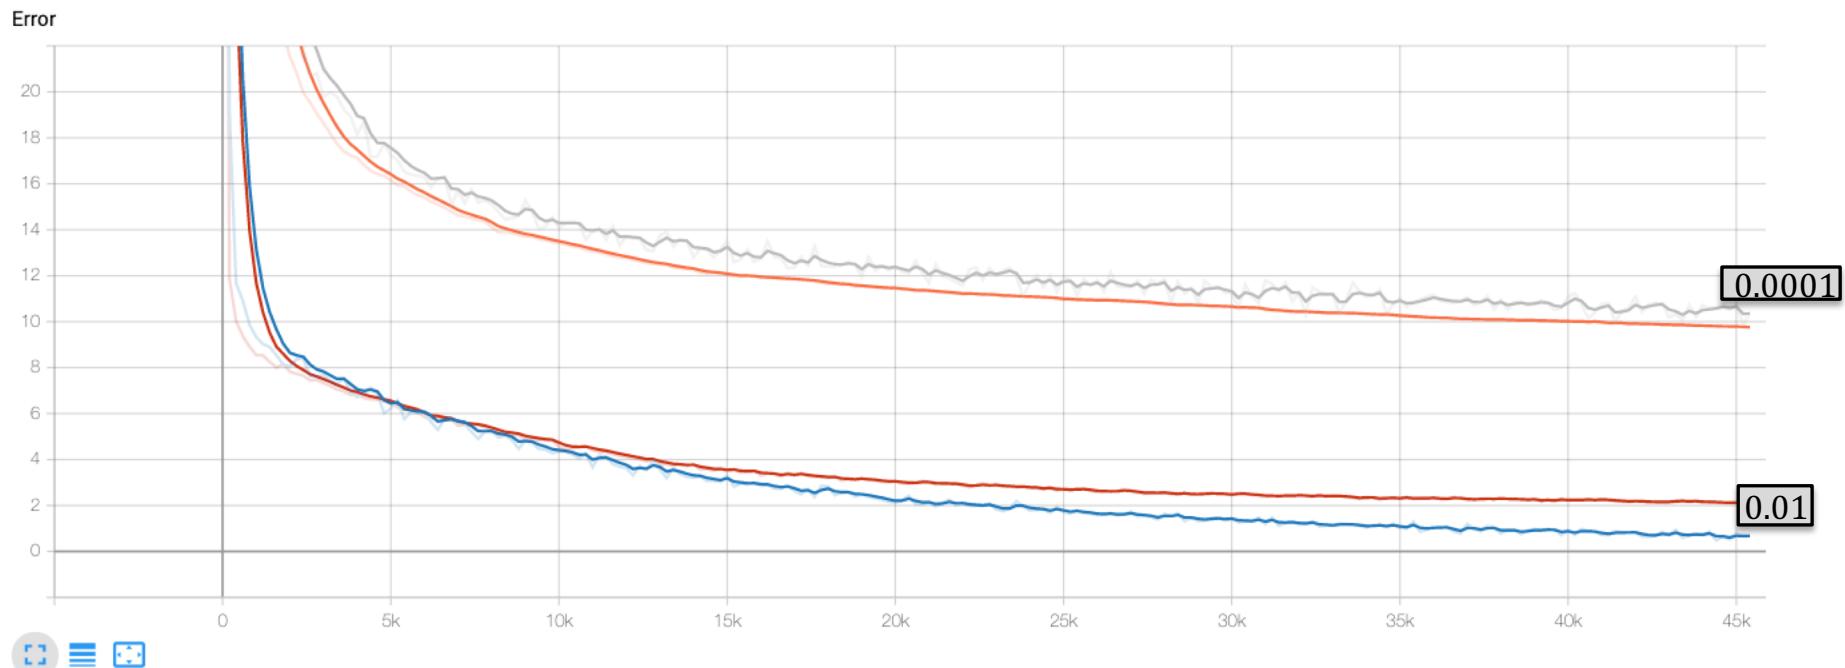
The impact of learning rates

A well chosen learning rate of 0.01:



The impact of learning rates

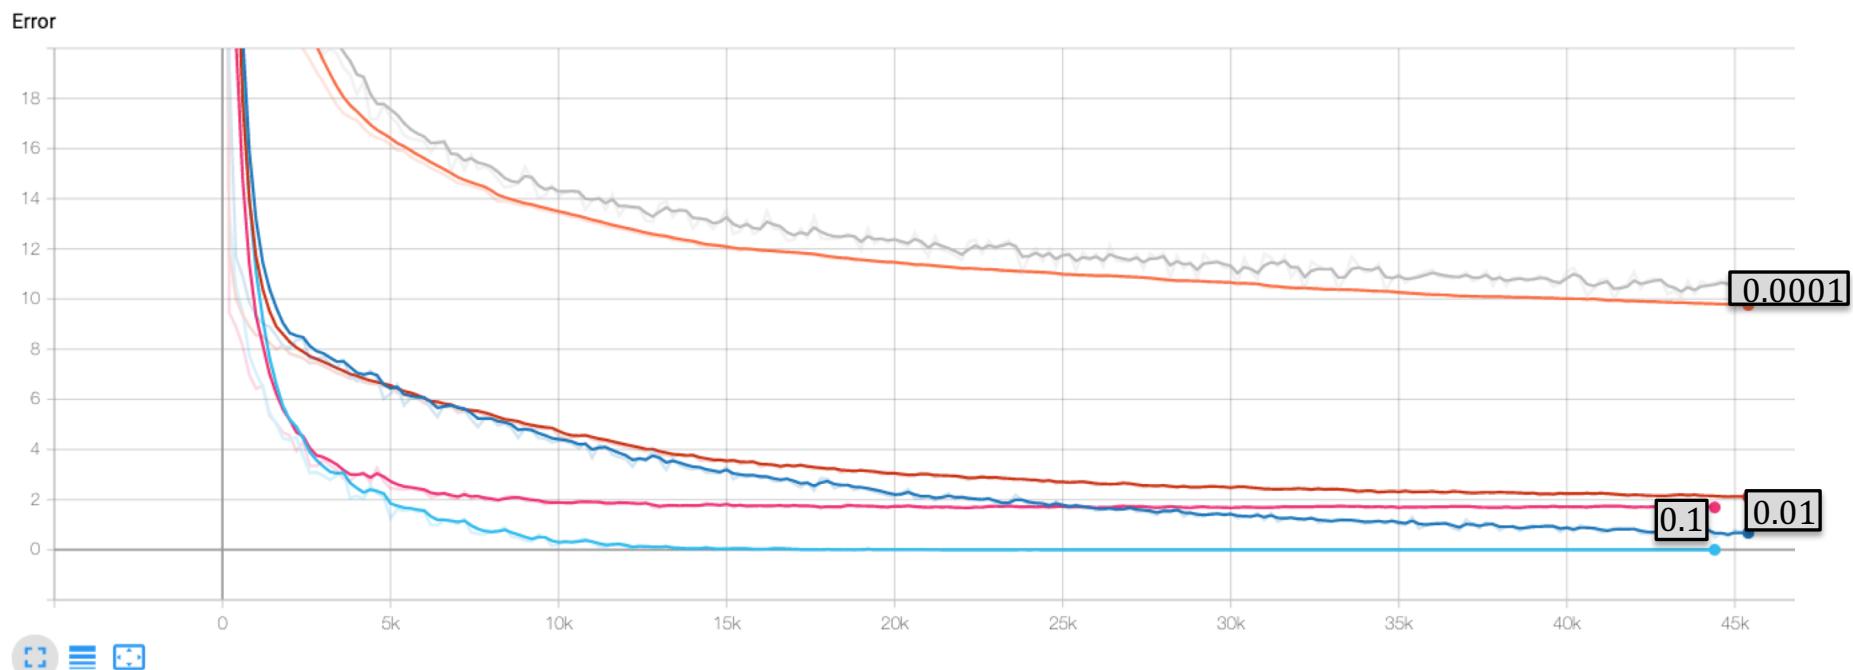
We add the curves for a low learning rate 0.0001:
Slow convergence.



The impact of learning rates

We add the curves for a high learning rate 0.1:

Convergence is fast at the beginning but fails to find a good optimum at the end.



The impact of learning rates

We add the curves for a ridiculously high learning rate 1:
Oscillations start to appear (convergence problems).

