

# *Lecture: Deep Learning and Differential Programming*

## 3.1 Computer Vision

**1** Image classification

**2** Object detection and recognition

**3** Semantic segmentation

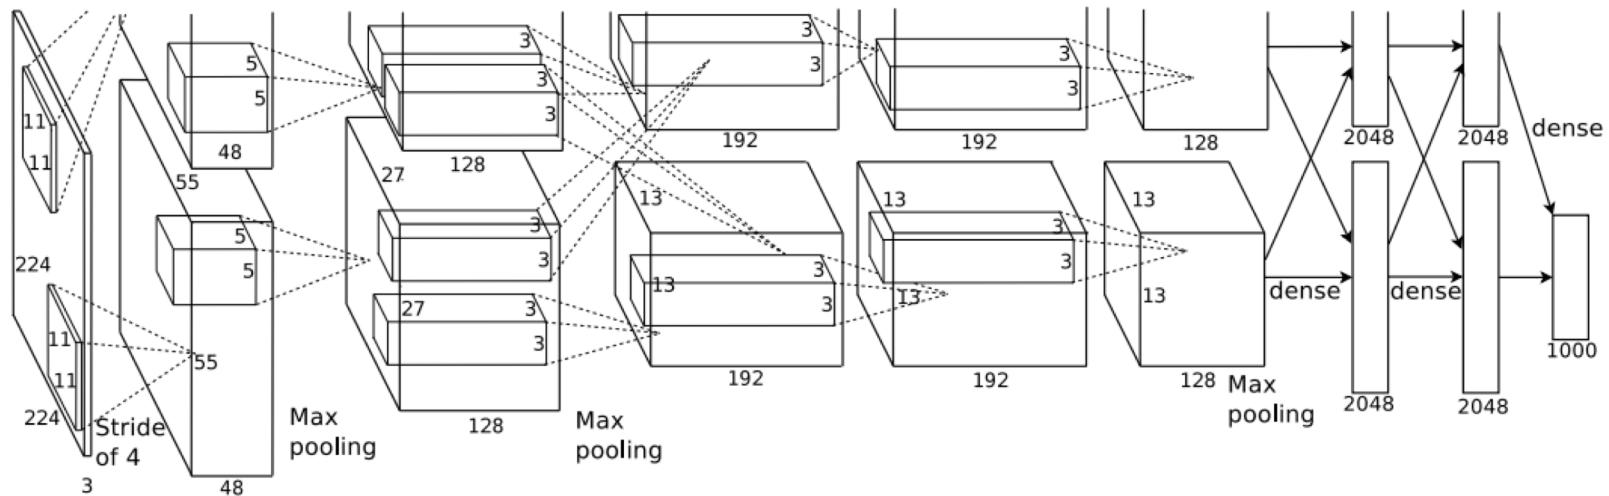
**4** Instance segmentation

# AlexNet

The model which made deep learning hugely popular by winning the Imagenet competition in 2012.

8 trainable layers (5 convolutions, 3 fully connected).

Contribution: ReLU activation, dropout, multi-GPU training



# VGG 16

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input ( $224 \times 224$ RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

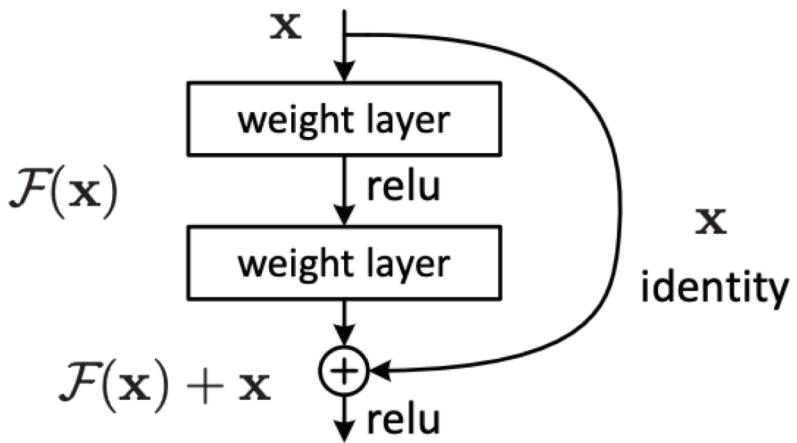
Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- 16 layers
- Only 3x3 convolutions

# ResNet

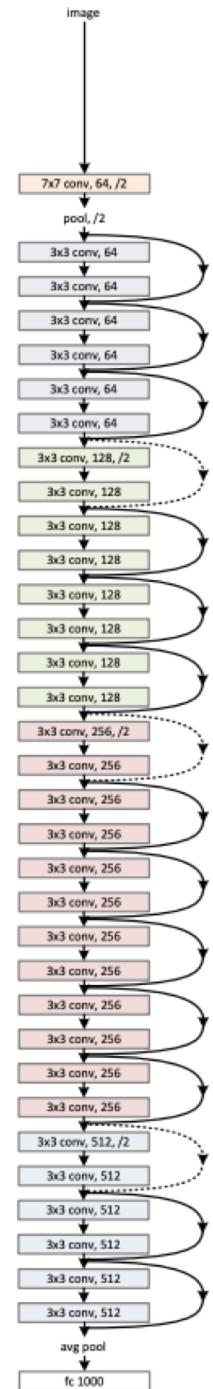
Wins ImageNet 2015 competition.

Novelty: residual blocks – a block predicts the difference to its input.



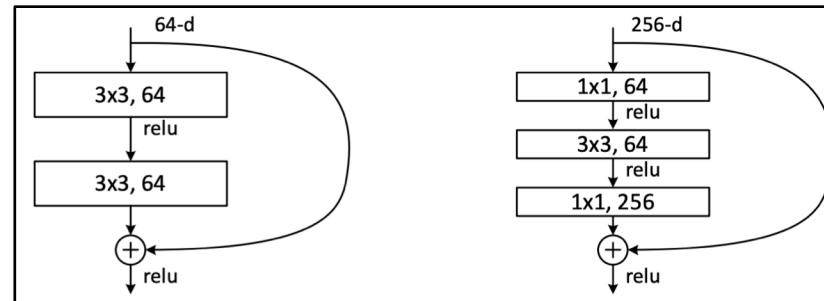
[He, Zhang, Ren, Sun,  
CVPR 2016]

Resnet 34



# ResNet: variants

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56			3×3 max pool, stride 2		
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1			average pool, 1000-d fc, softmax		
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$



**1** Image classification

**2** Object detection and recognition

**3** Semantic segmentation

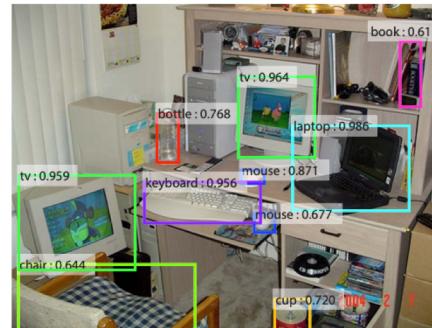
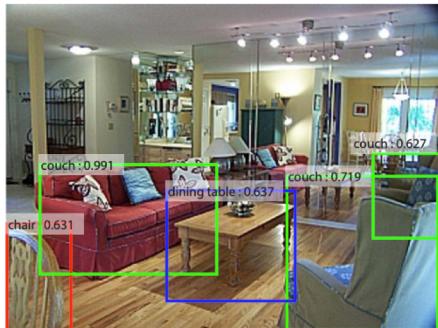
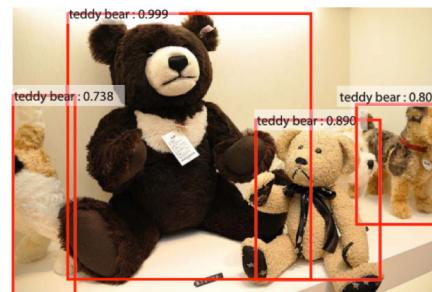
**4** Instance segmentation

# Detection, localization, recognition

A more complex problem. We need to:

- Detect whether an object exists
- Localize it (regress its bounding box coordinates)
- Recognize its class

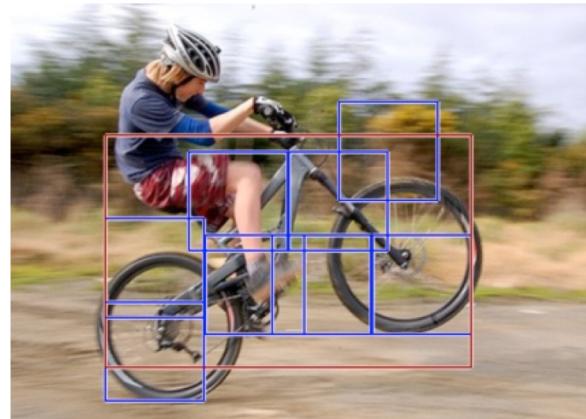
Multiple instances are possible



[Figure: Faster R-CNN, Ren, He, Girshick, Sun, NIPS 2015]

# Before Deep Learning: Deformable Parts Models

- Model an object/human/activity as a collection of local parts
- Learn a filter for each part
- Learn an anchor position and deformation coefficients for each part
- Test each image pixel whether it can be the center of the object (sliding window):
- For each possible center, optimize over (latent) local part positions

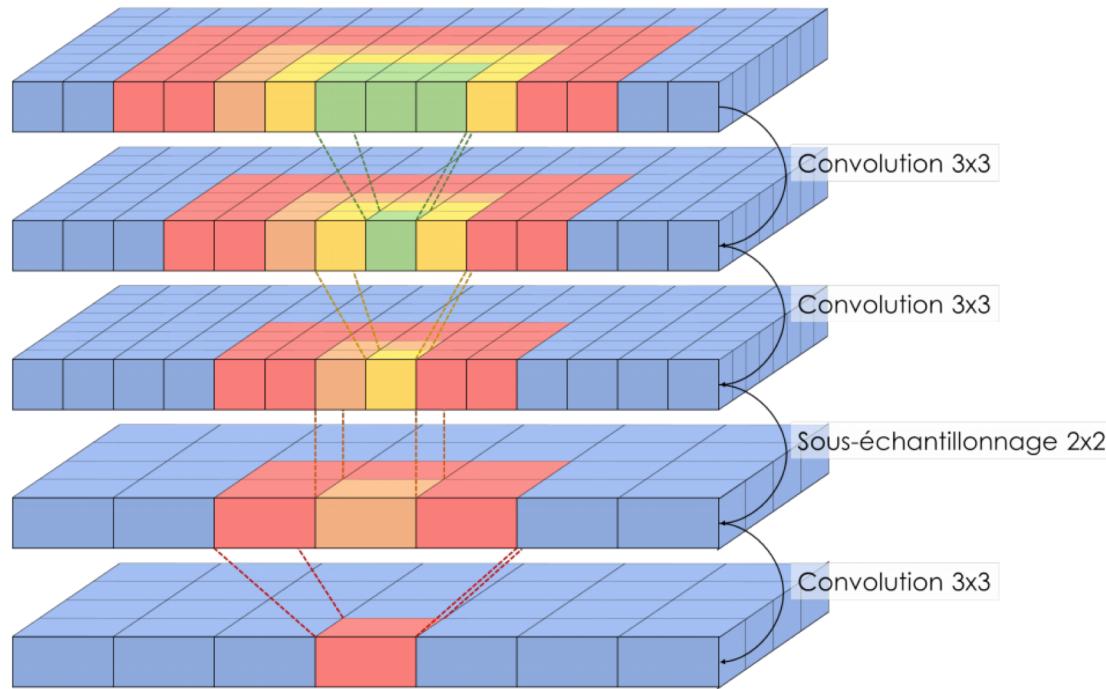


$$\sum_{i=0}^n F'_i \cdot \phi(H, p_i) - \sum_{i=1}^n d_i \cdot \phi_d(dx_i, dy_i) + b,$$

Local appearance      Deformation

# Spatial feature maps

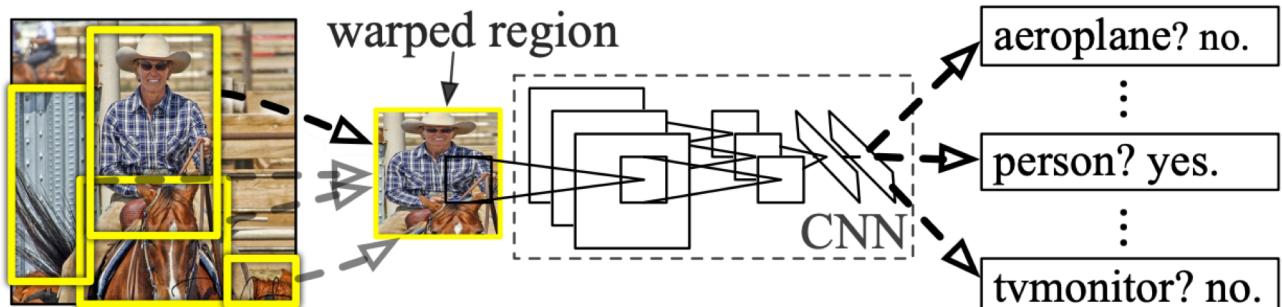
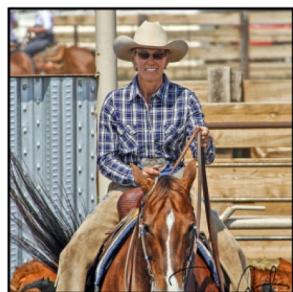
Most vision methods requiring localization exploit the fact that in convolutional neural networks (CNNs) intermediate layer activations have a spatial meaning: each cell corresponds to a rectangular area in the input image (“**receptive field**”).



[Figure: Damien Fourure,  
PhD thesis, 2017]

# R-CNN

- Detect a large number of candidate regions (« **region proposals** ») with some heuristic method
- Feed each candidate region into a convolutional neural network for recognition



**1. Input image**

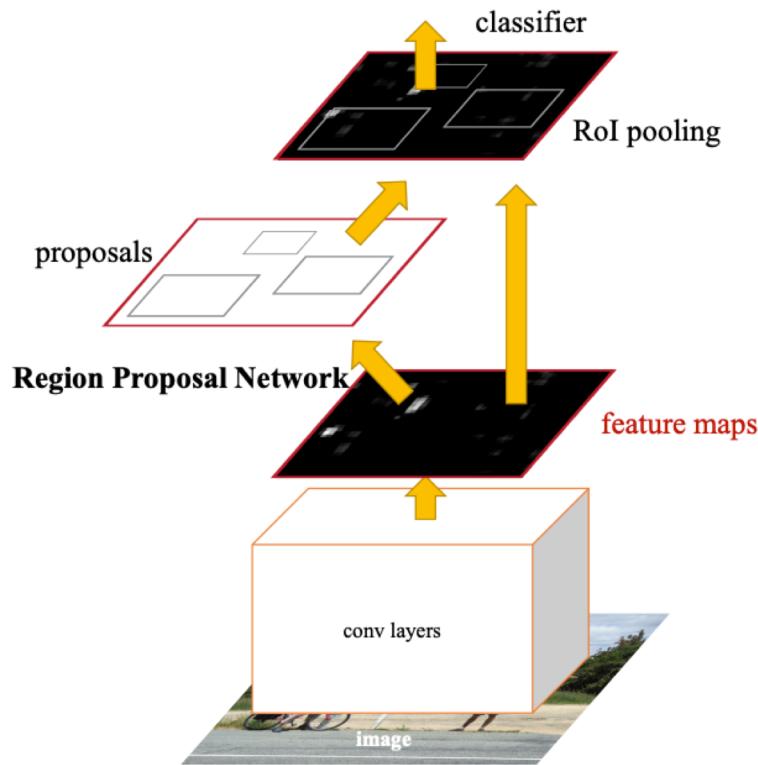
**2. Extract region proposals (~2k)**

**3. Compute CNN features**

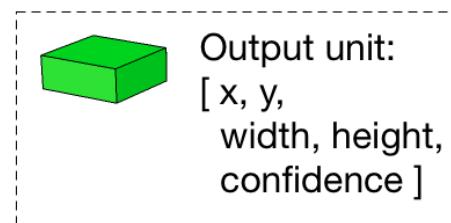
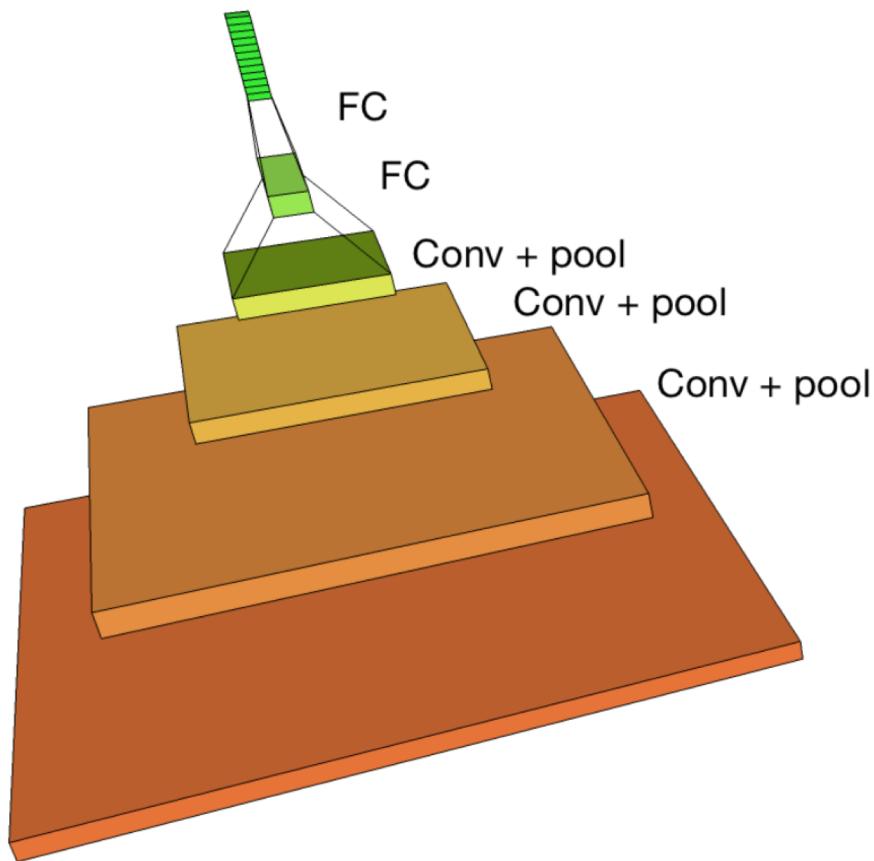
**4. Classify regions**

# Faster R-CNN

- The region proposals are now predicted by a neural network which is part of the full network
- For each proposal, features are collected **from the bounding box of the proposal** and fed to a classifier.



# Real time detectors



[Erhan et al.,2014]

« Multibox »

[Redmon et al.,2016]

« YOLO »

[Erhan et al.,2016]

« SSD »

**1** Image classification

**2** Object detection and recognition

**3** Semantic segmentation

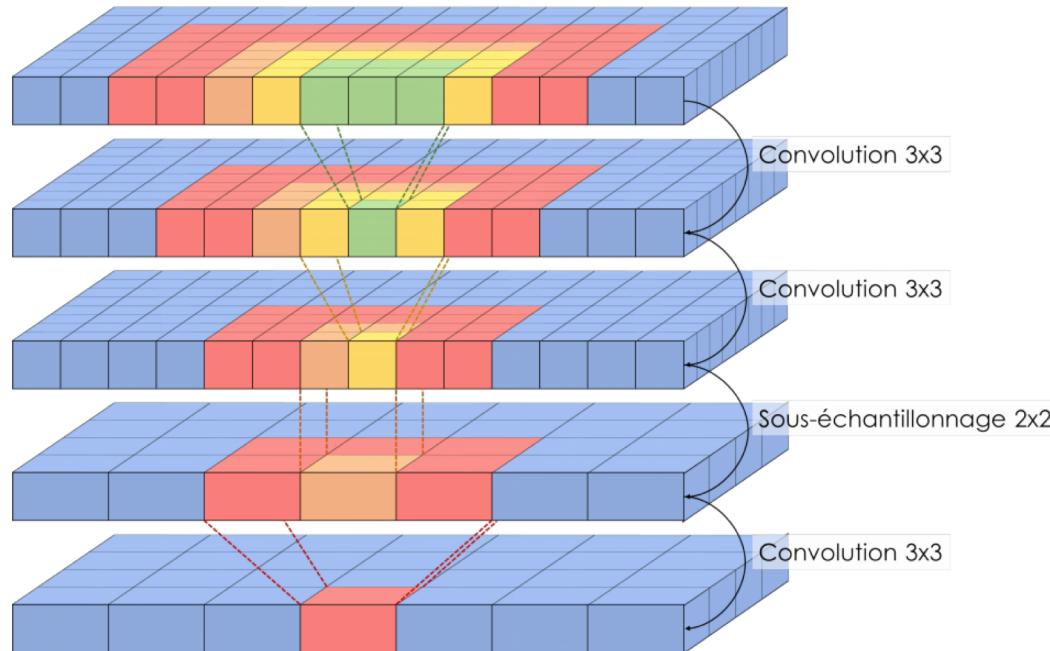
**4** Instance segmentation



# Recall: receptive Fields

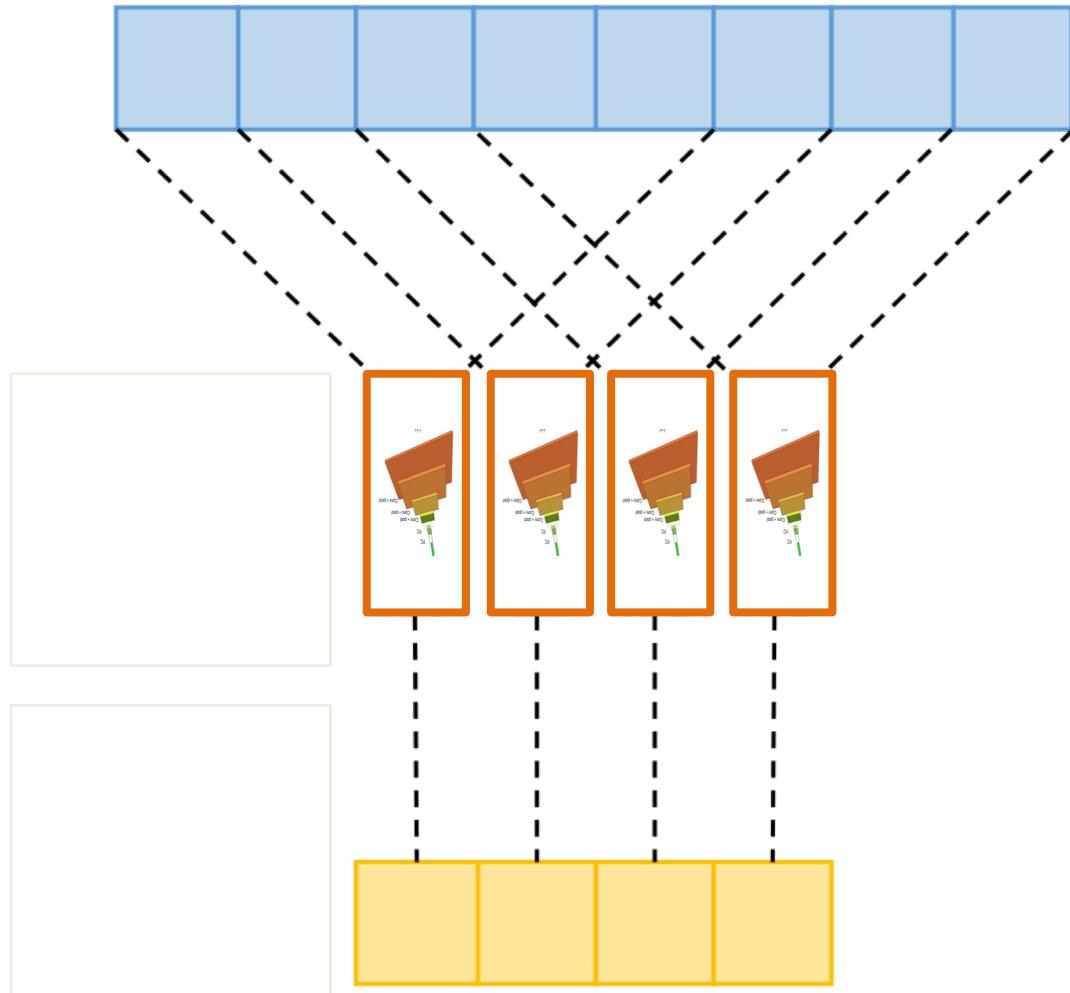
Pooling layers reduce the spatial resolution of intermediate activations and output layers.

How can we create dense predictions while keeping the output resolution equal to the input resolution?



[Figure: Damien Fourure,  
PhD thesis, 2017]

# Patchwise processing



A direct, simple  
and early  
method.

Not used  
anymore

Feed each pixel  
+ neighborhood  
as input into a  
network.

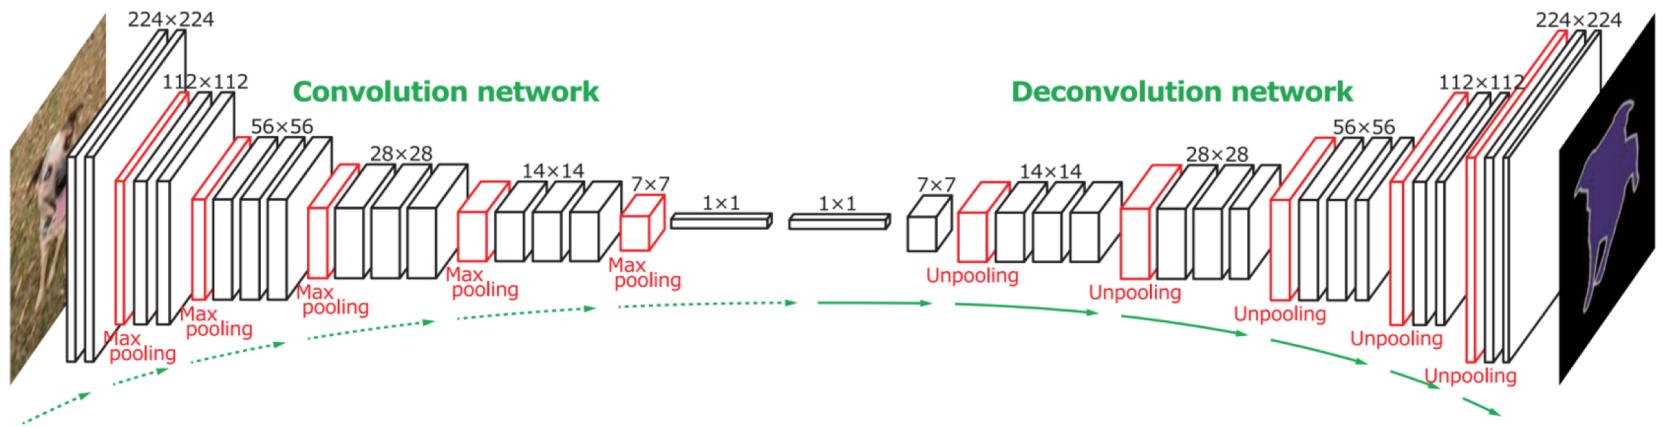
Slow!

# Conv-Deconv Networks

A classical "encoder network" produces a vectorial representation through convolutions and pooling.

A second network ("decoder") decodes this into an output image with the initial resolution.

All information must pass through the bottleneck layer!

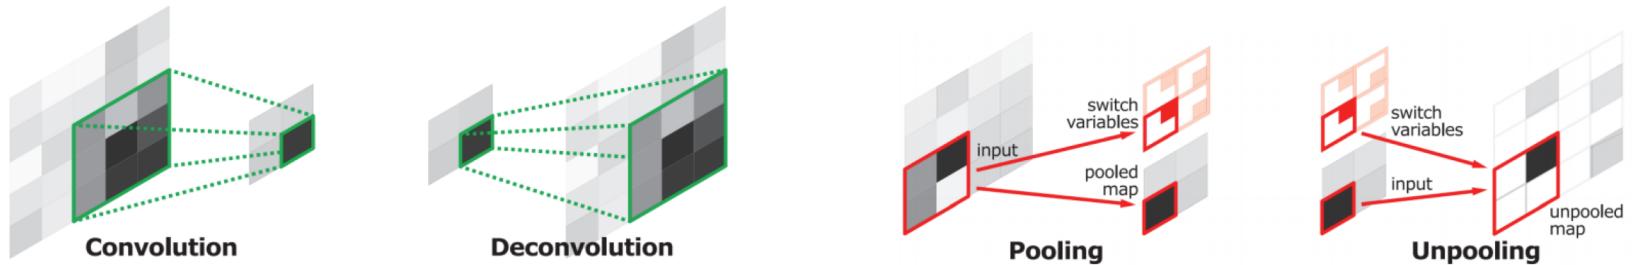


[Noh et al., 2015]

# Conv-Deconv Networks

In the decoder:

- convolutions are replaced with “deconvolutions” or “transposed convolutions”.
- Pooling is replaced with unpooling (switch variables keep the arg max location of the pooling layer)

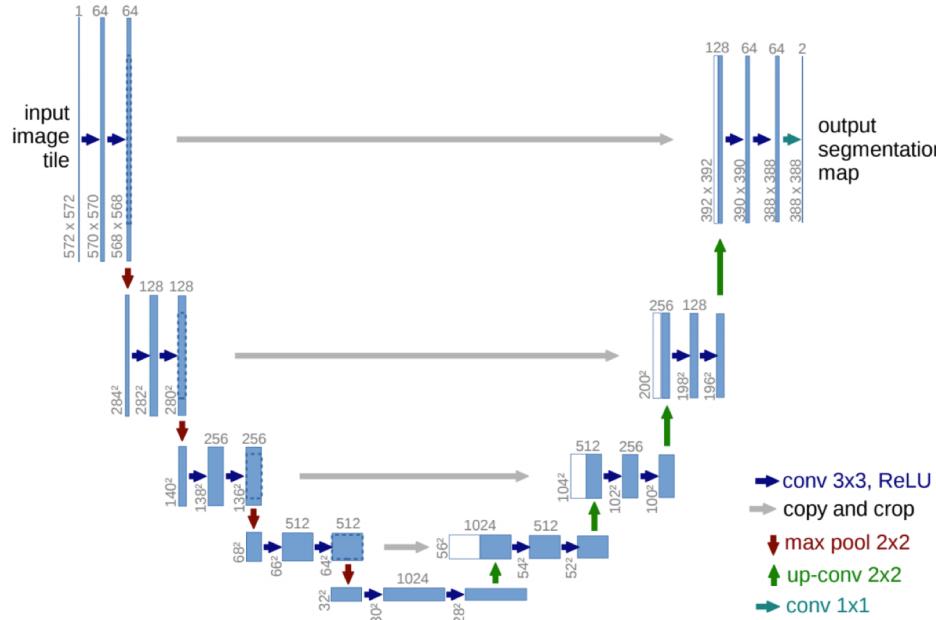


[Noh et al., 2015]

# U-Nets

Conv-Deconv: all information passes through the bottleneck layer.

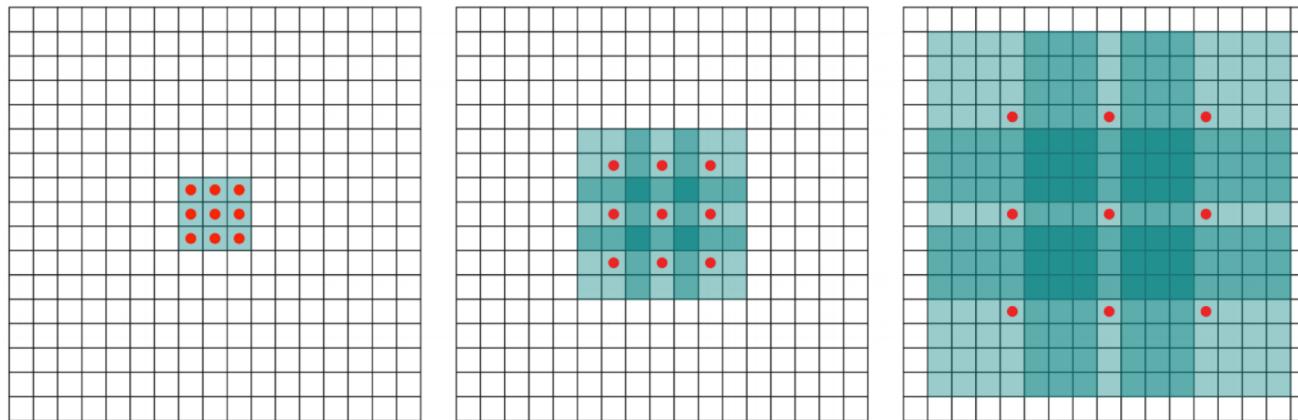
U-Nets add additional skip connections for low level information, close to pixels.



# Dilated Convolutions

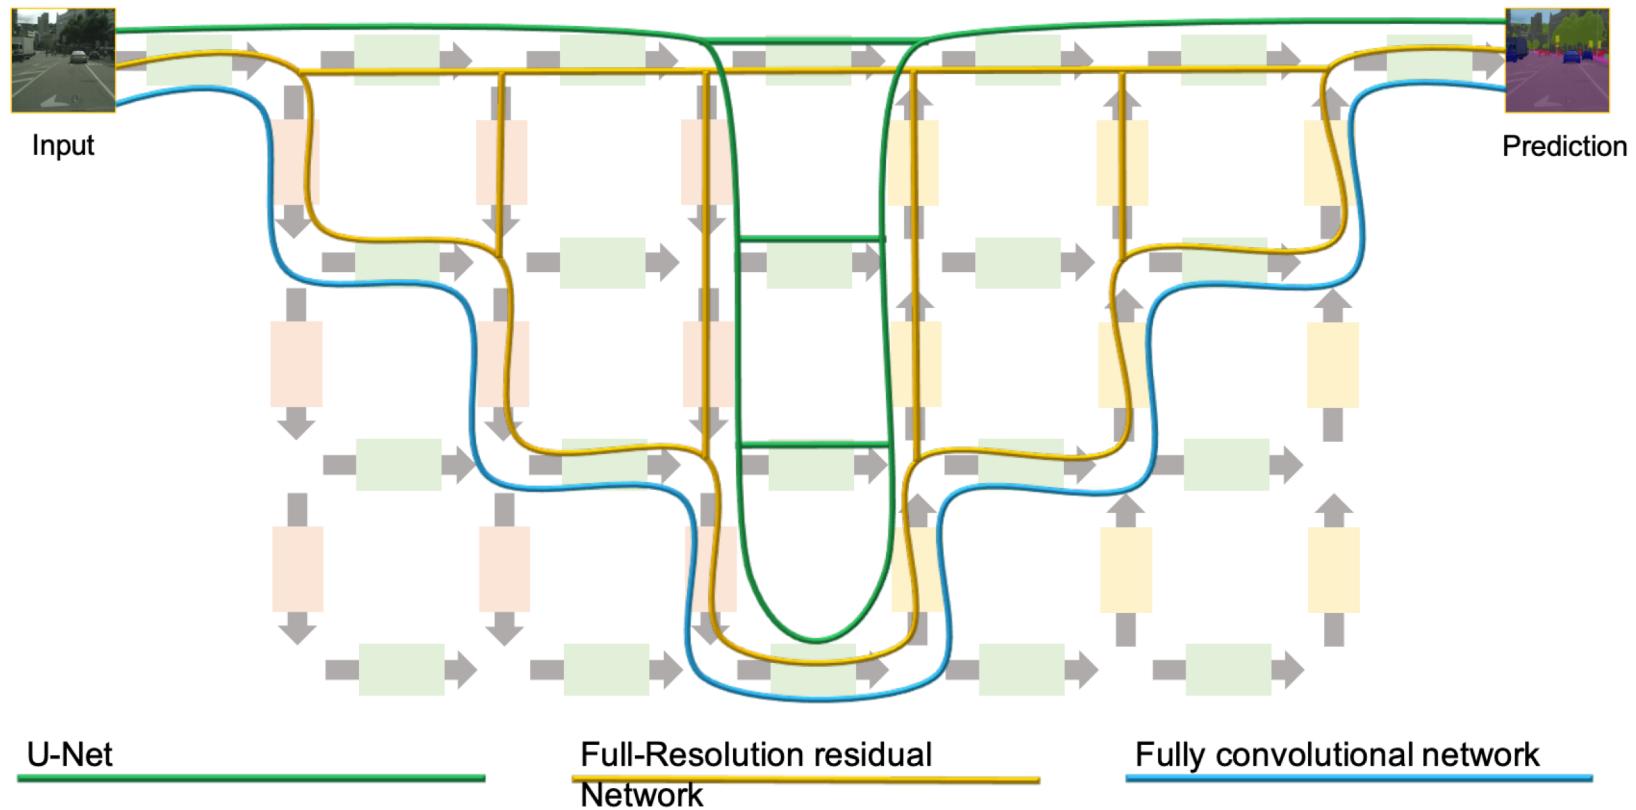
Alternative: do not use any pooling, keep spatial resolution throughout the network.

To increase the receptive field, change the size of the filters ... w/o augmenting the number of parameters!

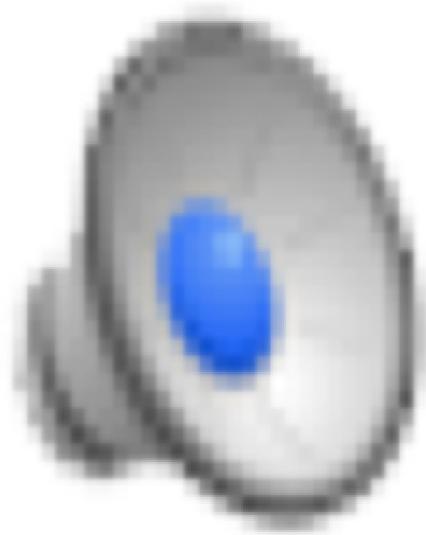


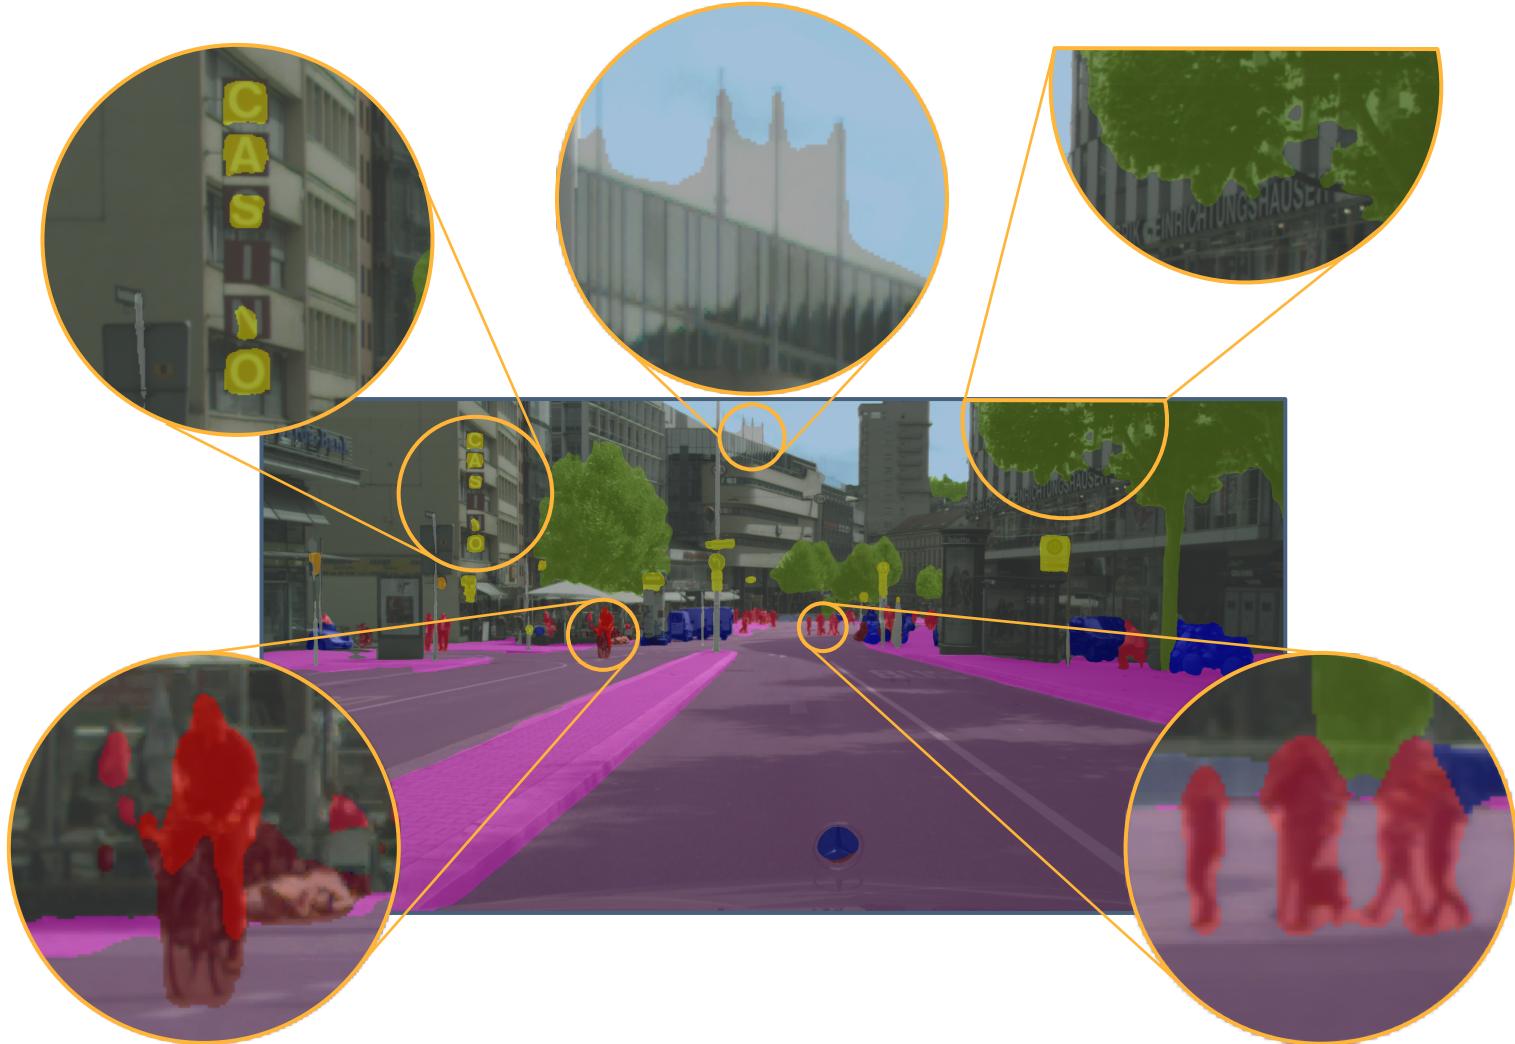
# Grid Networks

Grid Networks generalize a large number of networks.



# Grid Networks





[Fourure, Emonet, Fromont, Muselet, Tremeau, Wolf, BMVC 2017]

**1** Image classification

**2** Object detection and recognition

**3** Semantic segmentation

**4** Instance segmentation

# Instance segmentation

Bridges the gap between semantic segmentation and object detection:

- One region = 1 object instance
- Pixelwise boundaries instead of bounding boxes



Semantic Segmentation



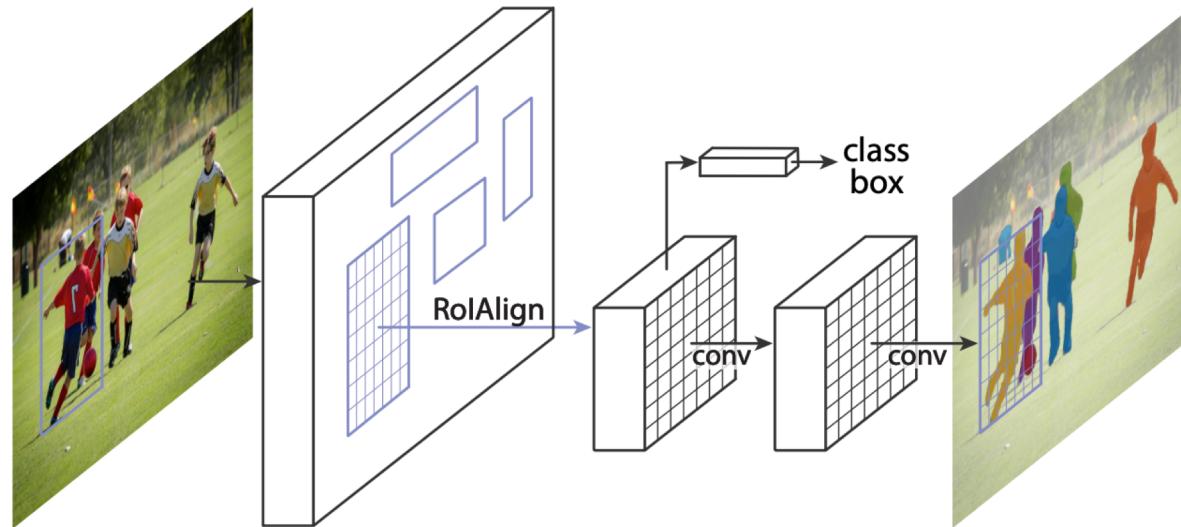
Instance Segmentation

# Mask R-CNN

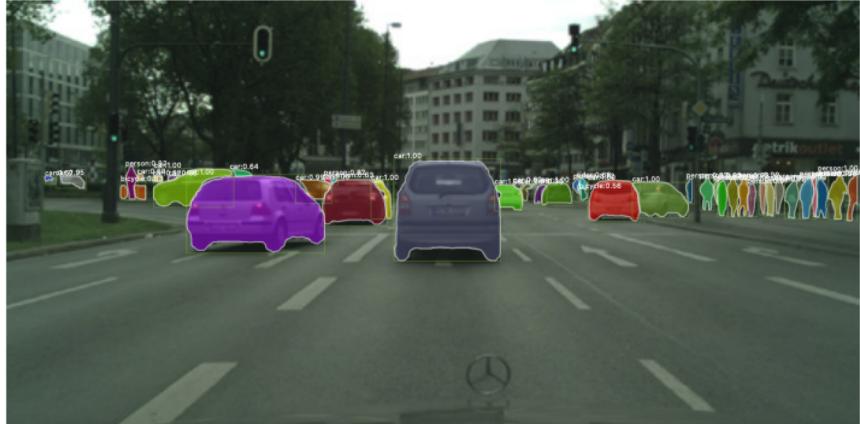
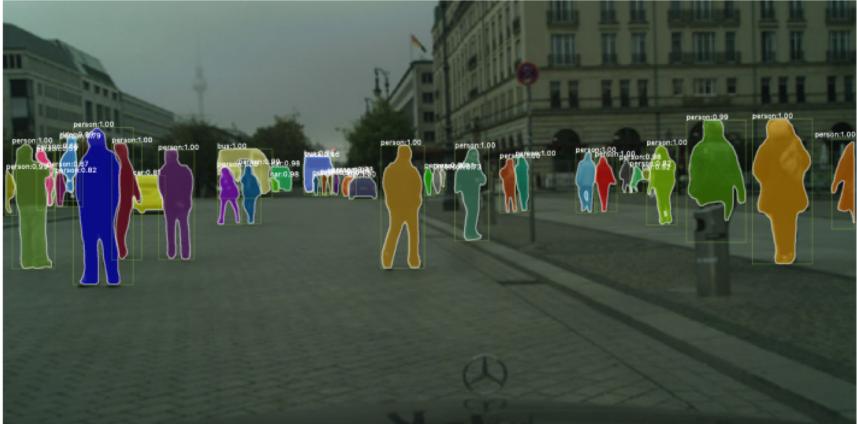
Region proposals as in object detection

Conv-Deconv network like in semantic segmentation ... but  
for each region proposal.

ResNet-101 Backbone.



# Mask R-CNN

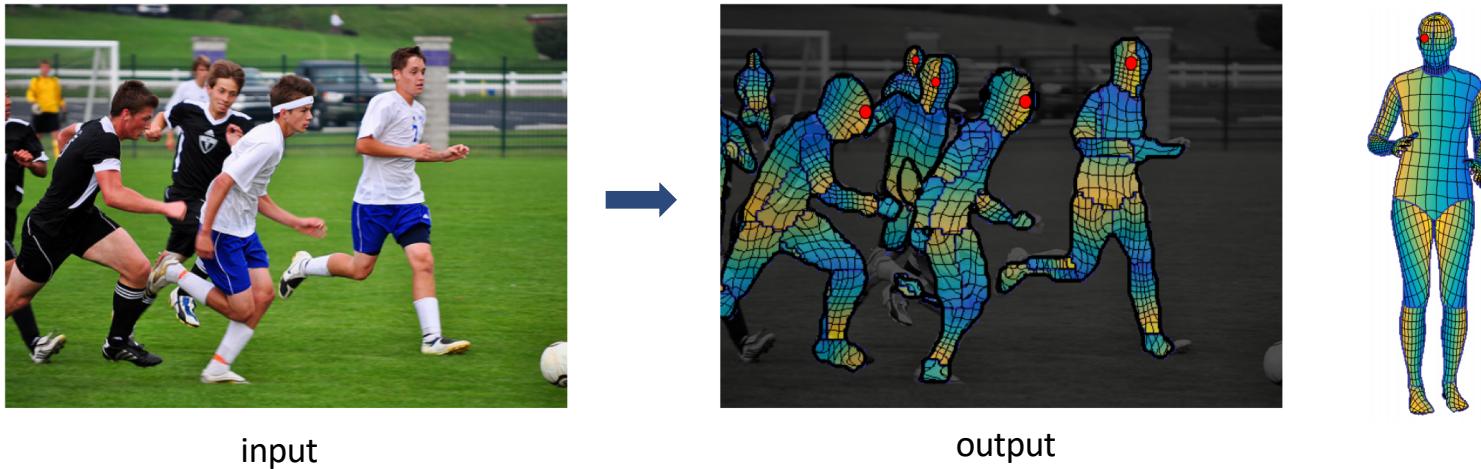


[Hen, Gkioxari, Dollar,  
Girshick, ICCV 2017]

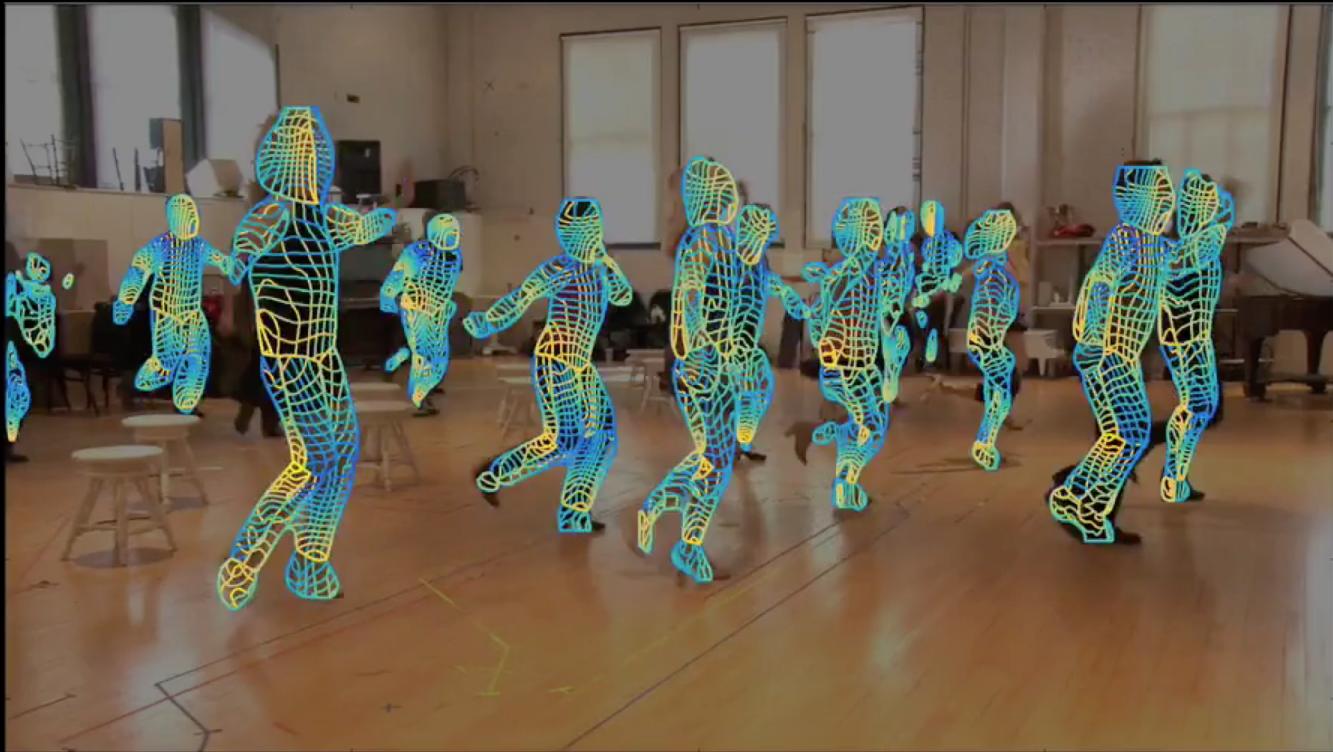
# Dense Pose

Densepose estimates human articulated pose in a dense manner:

One estimate per image pixel, corresponding to two parameters ( $u, v$ ) on the human body.

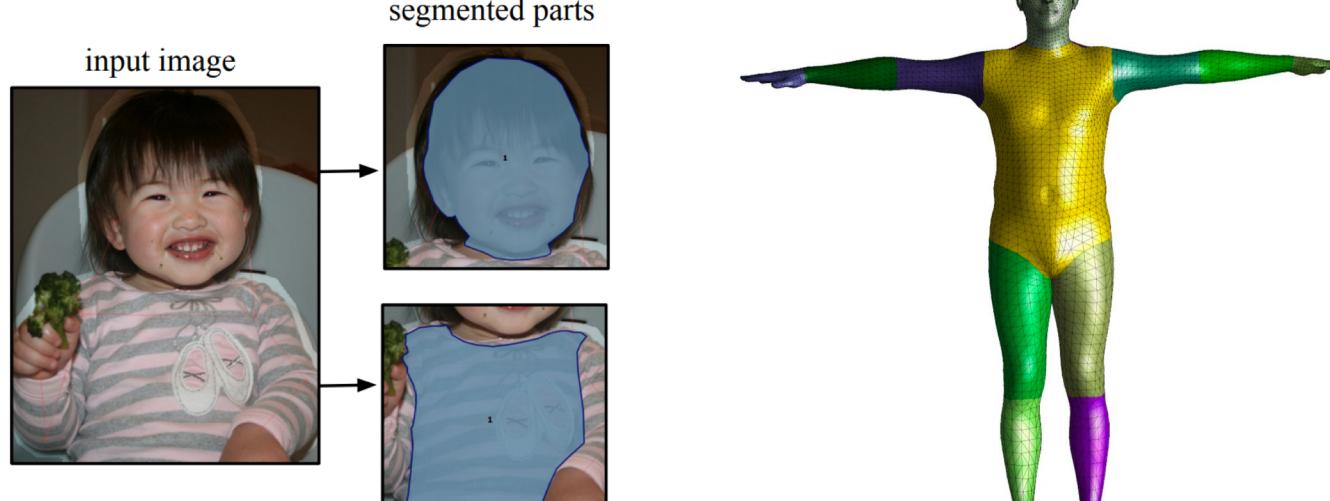


# PoseTrack dataset: Visual results



# The Densepose dataset

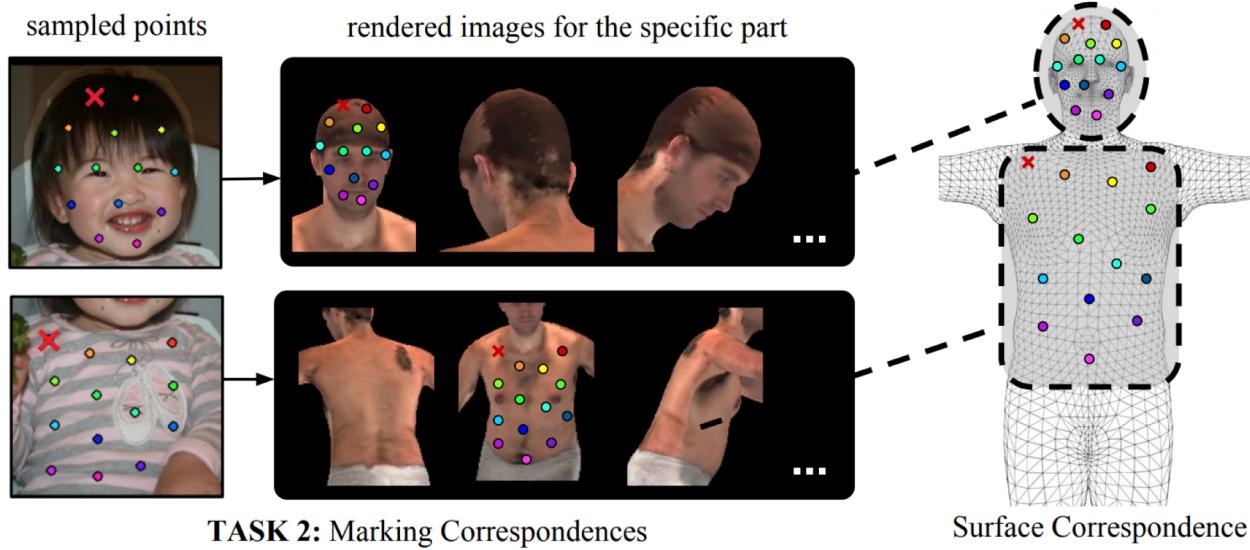
Annotation task 1: body part segmentation



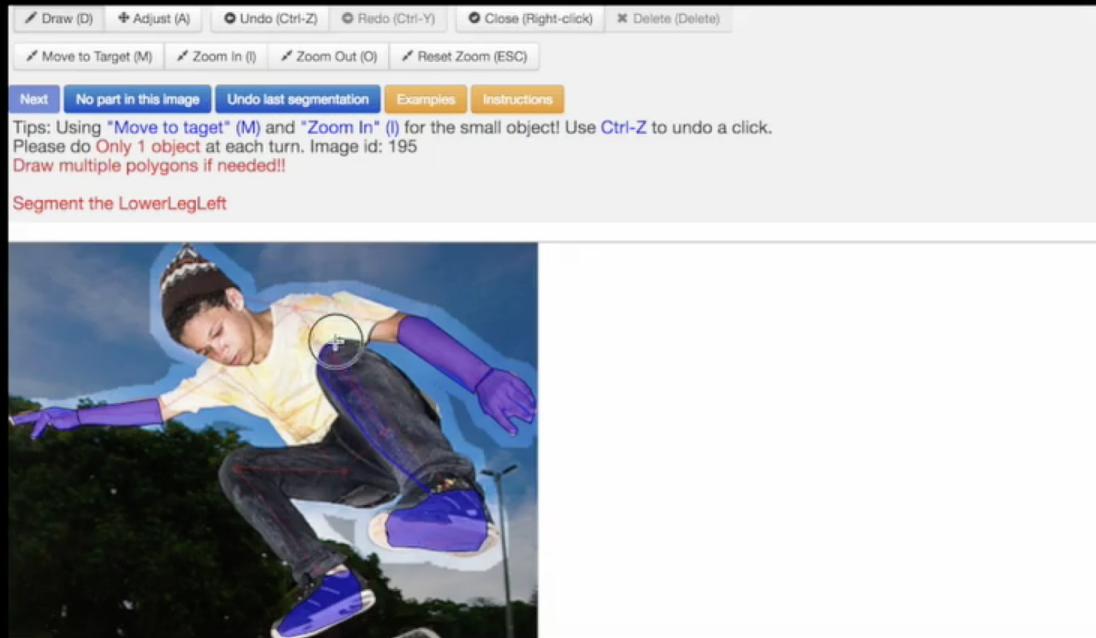
**TASK 1:** Part Segmentation

# The Densepose dataset

Annotation task 2: marking sparse correspondences



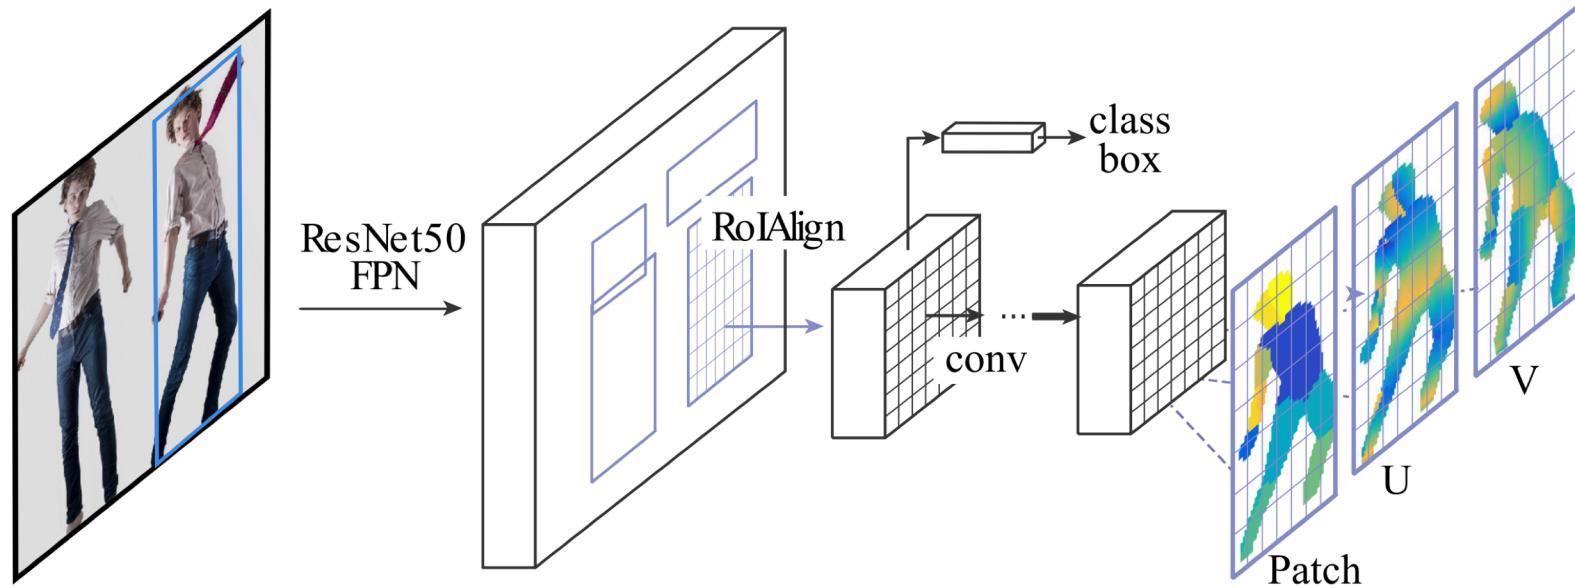
# COCO DensePose: Collecting Data



## Task - 1 Part Segmentation

# The neural architecture

Densepose builds on the Marks R-CNN architecture.



# Model Zoos

Different libraries propose « model zoos » containing well-known network architectures, sometimes with pre-trained parameters learned from standard datasets.

For PyTorch: `torchvision.models`:

- [AlexNet](#)
- [VGG](#)
- [ResNet](#)
- [SqueezeNet](#)
- [DenseNet](#)
- [Inception v3](#)
- [GoogLeNet](#)
- [ShuffleNet v2](#)
- [MobileNet v2](#)
- [ResNeXt](#)
- [Wide ResNet](#)
- [MNASNet](#)

# The torchvision model zoo

Construct models with random weights:

```
1 import torchvision.models as models  
2 resnet18 = models.resnet18()  
3 alexnet = models.alexnet()  
4 vgg16 = models.vgg16()
```

Construct models with pre-trained weights (on ImageNet):

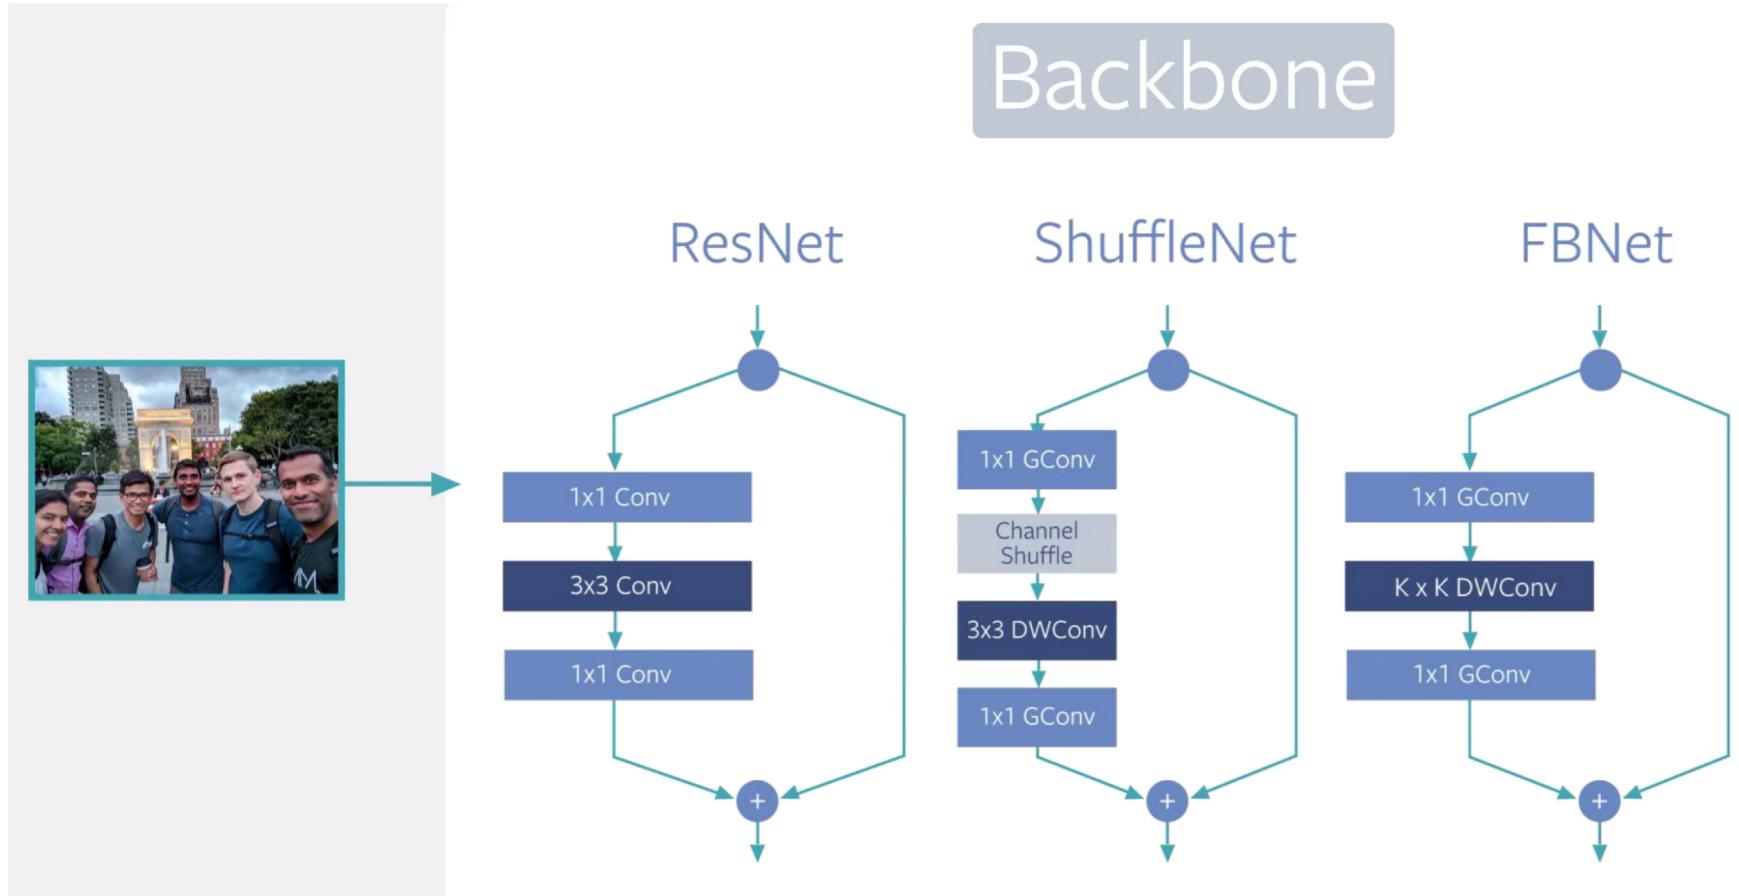
```
1 import torchvision.models as models  
2 resnet18 = models.resnet18(pretrained=True)  
3 alexnet = models.alexnet(pretrained=True)  
4 vgg16 = models.vgg16(pretrained=True)
```

# Facebook Detectron2

A single method capable of creating different predictions of different granularity using the same choice of “backbone” network (=network calculating features).

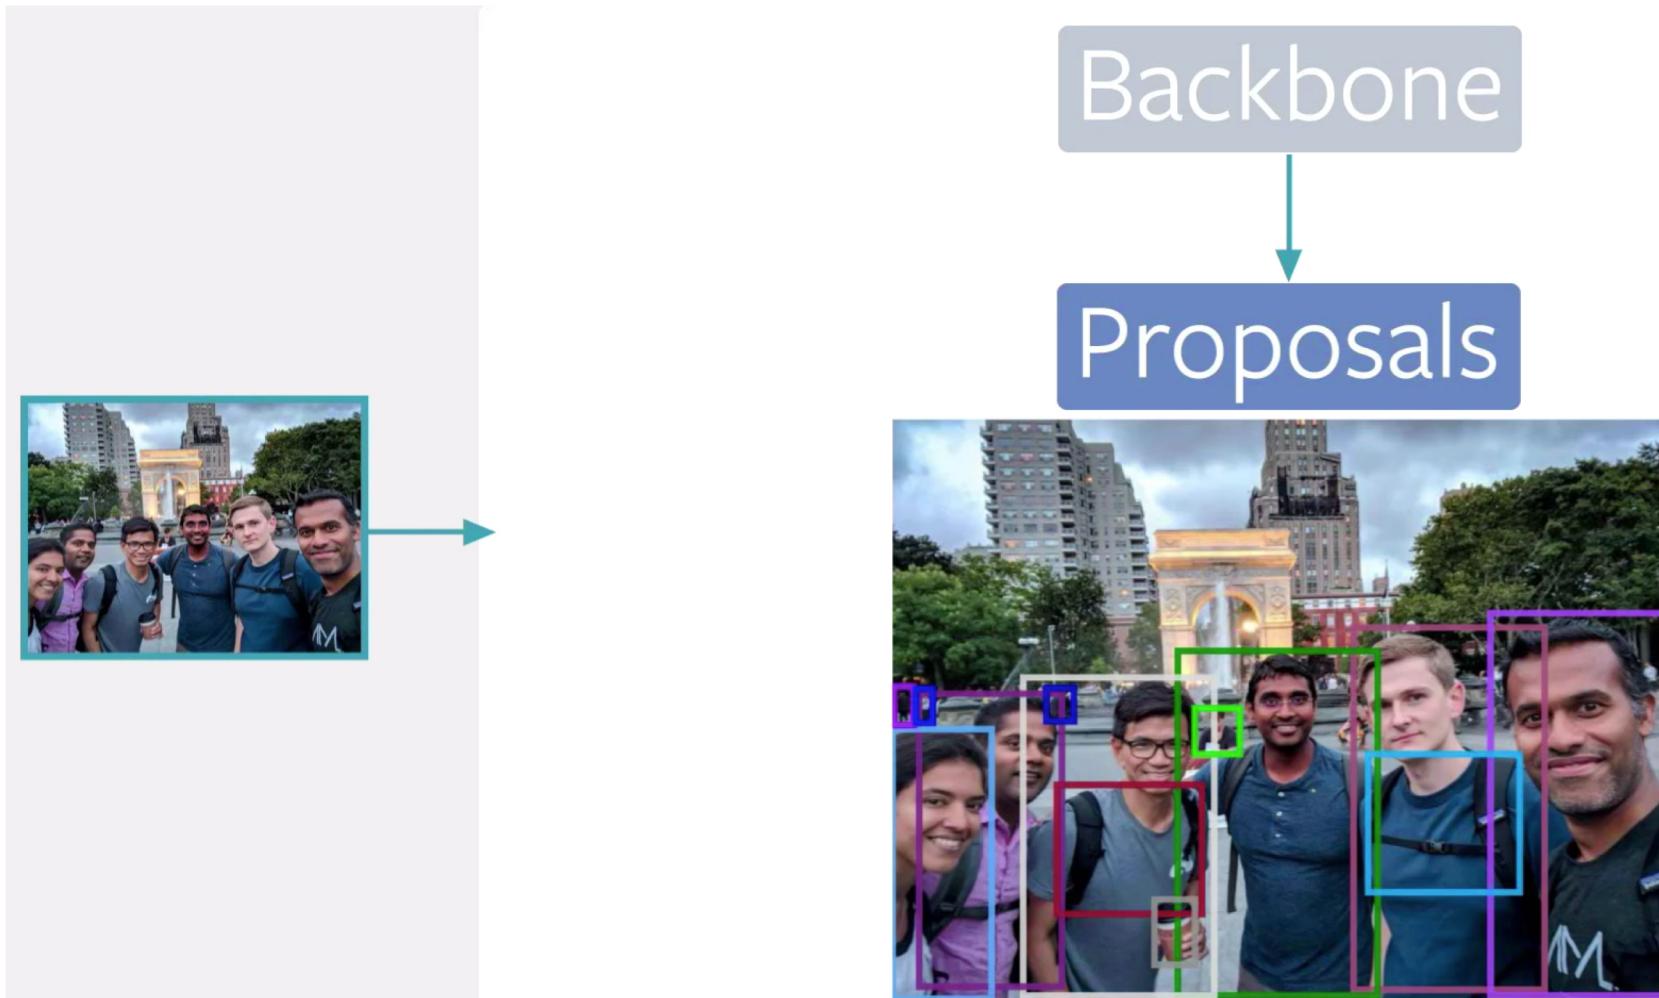


# Detectron backbones



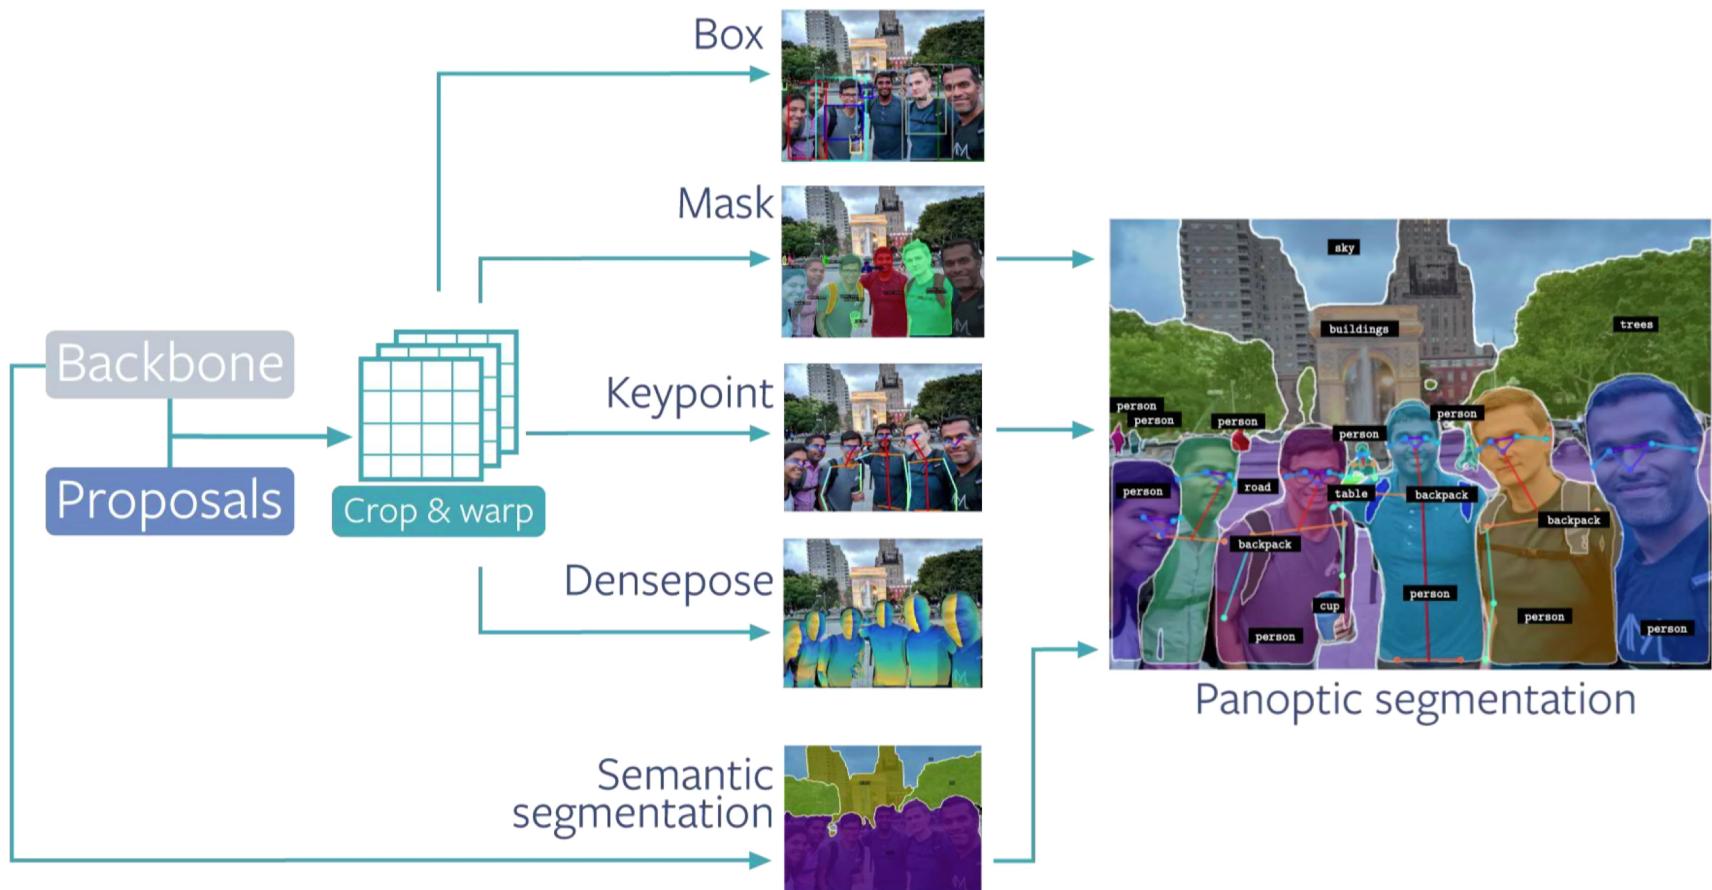
[Figure: Facebook]

# Detectron: proposals



[Figure: Facebook]

# Detectron: heads



[Figure: Facebook]