

# *Lecture: Deep Learning and Differential Programming*

## 3.3 GPUs - Software

# APIs

GPUs can be programmed on multiple different levels of abstraction.

We will (very briefly) study two cases:

CUDA/C++: direct low-level GPU programming

```
1 __global__  
2 void mult_kernel_simple(int mxWidth, float *mx1, float *mx2, float *output)  
3 {  
4     int c = blockIdx.x*blockDim.x + threadIdx.x;
```

PyTorch/Python: high-level of abstraction

```
31 model = LeNet().to("cuda:0")  
32  
33 for batch_idx, (data, labels) in enumerate(train_loader):  
34     data = data.to("cuda:0")
```

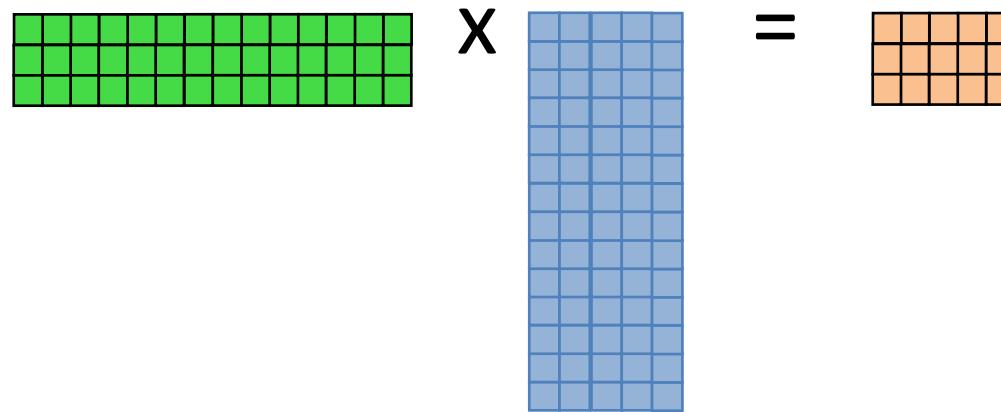
# 1

CUDA – low level GPU programming

# 2

GPUs & PyTorch

# Example : matrix multiplication



The classical sequential solution:

```
1   for( int r = 0; r < h; r++)
2   {
3       for( int c = 0; c < w; c++)
4       {
5           m1xm2[r*w + c] = 0;
6
7           for( int k = 0; k < w; k++)
8               m1xm2[r*w + c] += m1[r*w + k] * m2[k*w + c];
9
10    }
```

# The parallel solution

Parallel execution of a function called for each individual result value of the result matrix, with arguments being the indices of the value.

```
1 void mult_kernel_simple(int c, int r)
2 {
3     output[r*mxWidth + c] = 0.0f;
4     for( int k = 0; k < mxWidth; k++)
5         output[r*mxWidth + c] += mx1[r*mxWidth + k] * mx2[k*mxWidth + c];
6 }
```



Called in parallel for all c,r

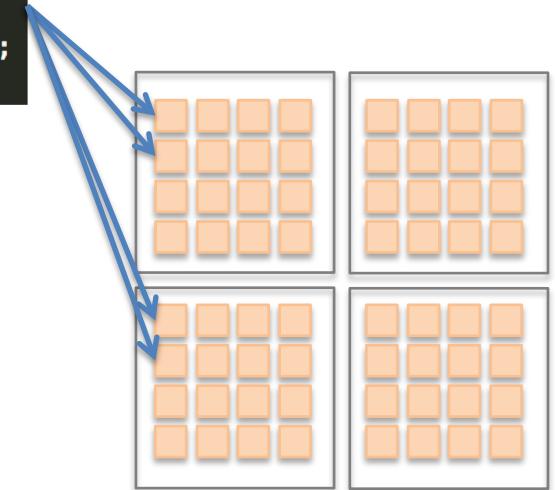
In CUDA, this function is called a *kernel*.

Each tuple (c,r) corresponds to a *thread*.

# Organisation into *blocks*

- Threads are organized into *blocks*
- Faster local memory can be shared by threads of the same block.

```
1 void mult_kernel_simple(int c, int r)
2 {
3     output[r*mxWidth + c] = 0.0f;
4     for( int k = 0; k < mxWidth; k++)
5         output[r*mxWidth + c] += mx1[r*mxWidth + k] * mx2[k*mxWidth + c];
6 }
```



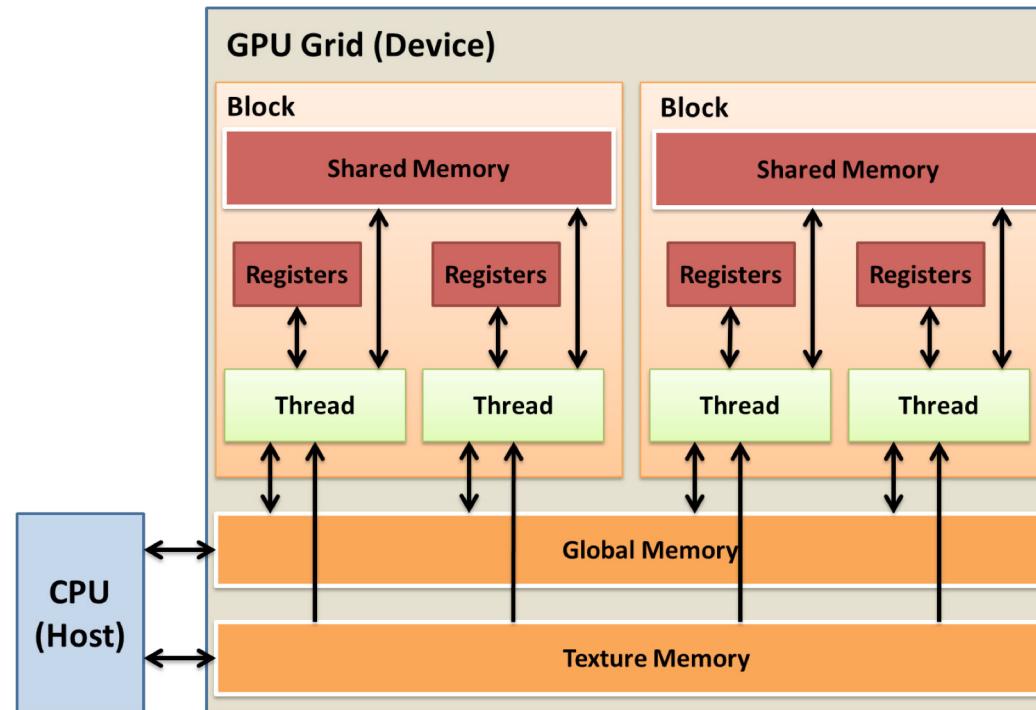
# SM

A block will be sent to a SM (*streaming multi-processor*)



# The classical CUDA sequence

- The CPU allocates memory on the GPU
- The CPU copies the data to GPU global memory
- The CPU launches the kernel on the GPU
- The GPU executes the kernel in parallel
- The CPU copies the result data back to CPU host memory



# The CUDA syntax of the kernel

Key word declares the *kernel*

```
1 __global__
2 void mult_kernel_simple(int mxWidth, float *mx1, float *mx2, float *output)
3 {
4     int c = blockIdx.x*blockDim.x + threadIdx.x;
5     int r = blockIdx.y*blockDim.y + threadIdx.y;
6
7     output[r*mxWidth + c] = 0.0;
8     for( int k = 0; k < mxWidth; k++)
9         output[r*mxWidth + c] += mx1[r*mxWidth + k] * mx2[k*mxWidth + c];
10 }
```

Index of the block in the grid

Index of the thread in the block

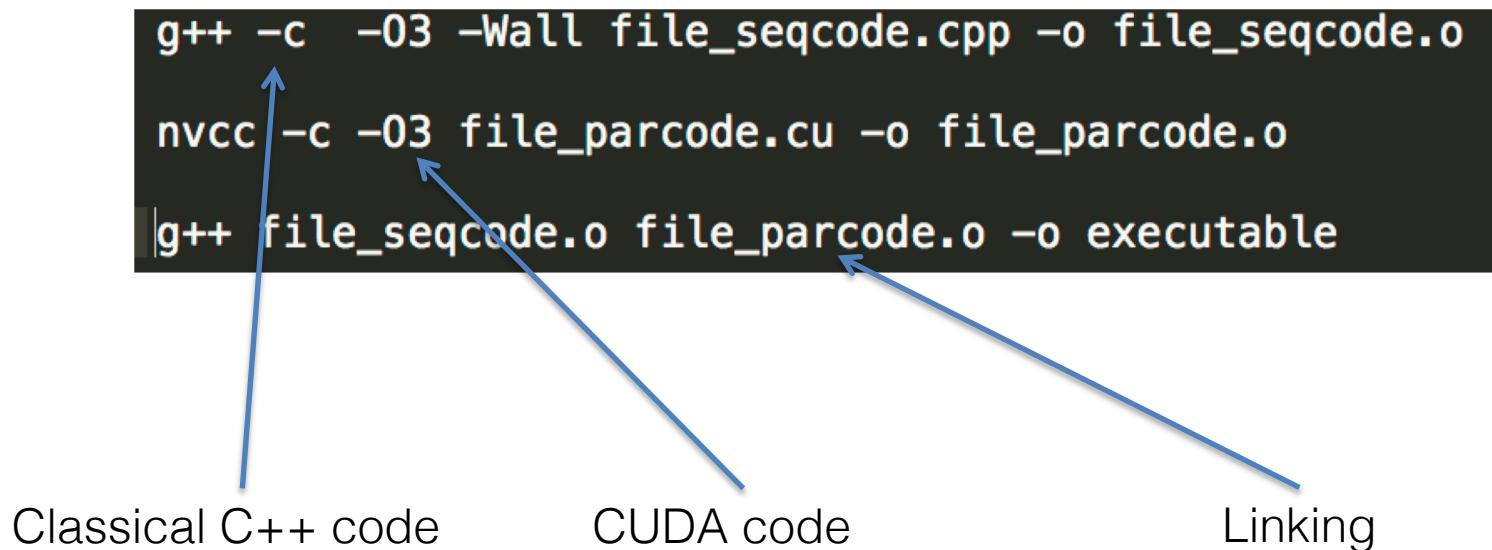
# Calling the kernel

```
1 int main(void)
2 {
3     ...
4     // allocate memory on GPU
5     cudaMalloc( (void**) &gpuMatrix1, matrixSizeInBytes);
6     ...
7
8     // copy data from CPU memory to GPU memory
9     cudaMemcpy(gpuMatrix1, matrix1, matrixSizeInBytes, cudaMemcpyHostToDevice);
10    ...
11
12    // Set grid and block size
13    dim3 dimBlock(32, 32);
14    dim3 dimGrid(matrixWidth/dimBlock.x, matrixWidth/dimBlock.y);
15
16    // run kernel
17    mult_kernel_simple<<<dimGrid, dimBlock>>>( matrixWidth, gpuMatrix1, gpuMatrix2, gpu0
18
19    // copy back results from GPU memory to CPU memory
20    cudaMemcpy( outputData, gpuOutput, matrixSizeInBytes, cudaMemcpyDeviceToHost);
21    ...
22 }
```

Call the *kernel*/in parallel for a set of threads

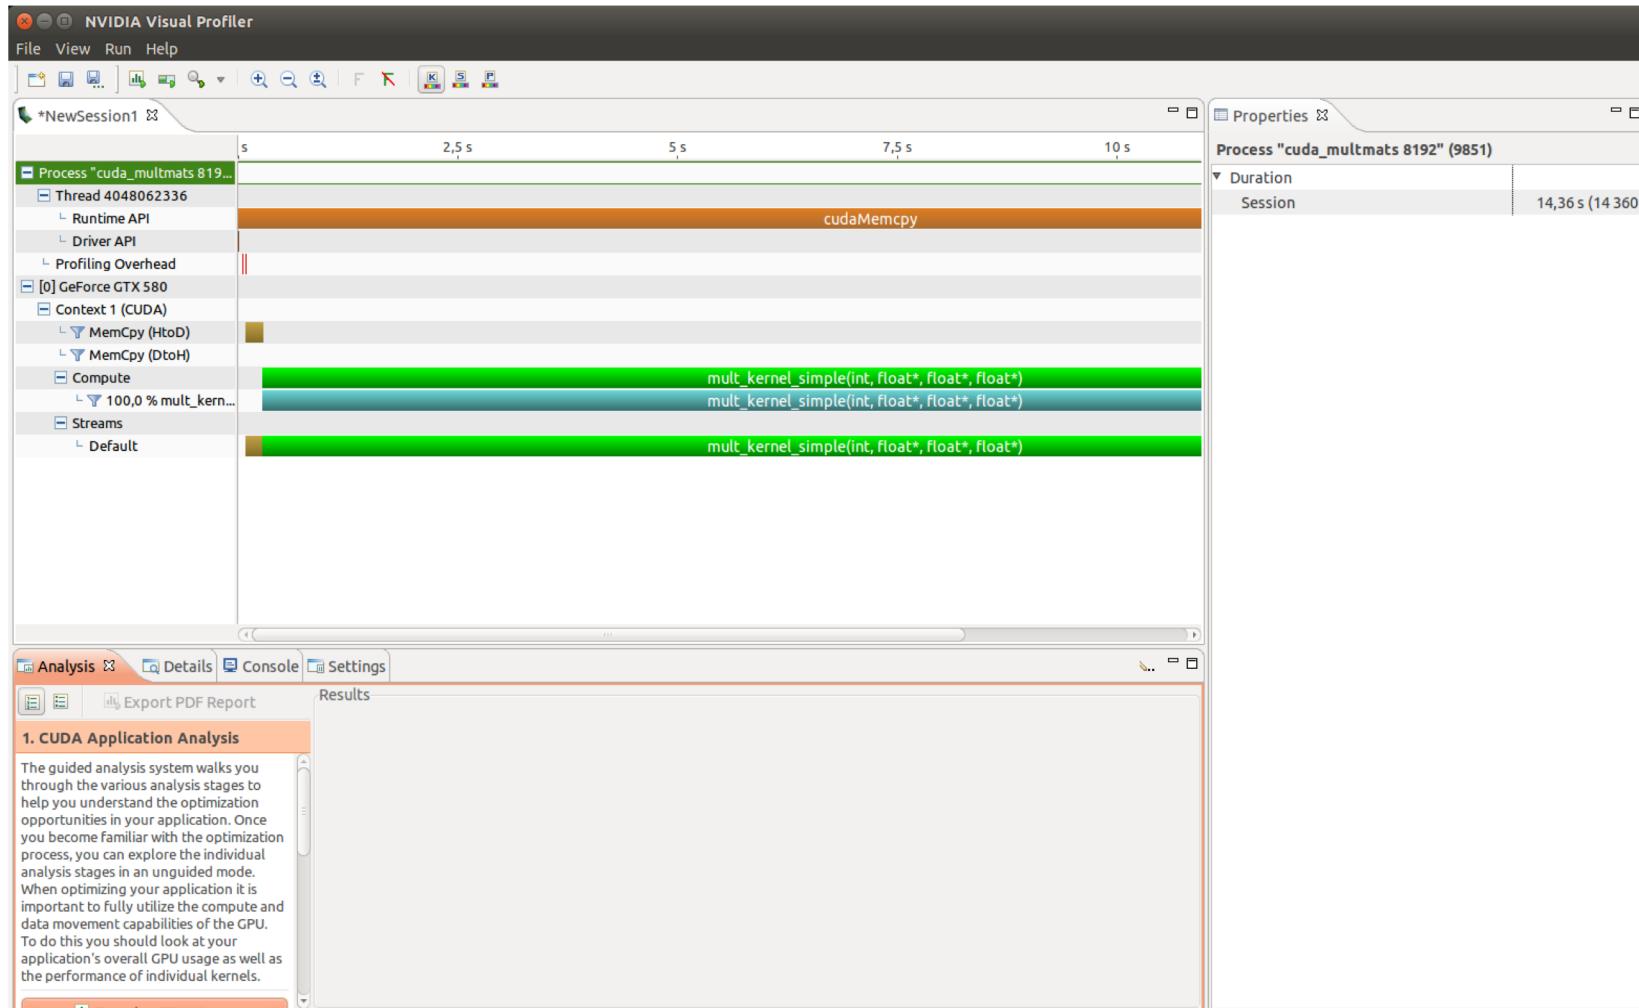
# Compilation

CUDA uses a specific compiler, which is based on a generic C++ compiler (gcc)



# Debugging and profiling

The nvvp profiler is part of the CUDA toolkit.



# 1 CUDA – low level GPU programming

## 2 GPUs & PyTorch

# The PyTorch GPU interface

The model is transferred to GPU memory with `.to(device)`

Device can be "cpu", "cuda:0", "cuda:1" etc.

```
1 model = LeNet()  
2 model = model.to("cuda:0")
```

We also send the data to GPU memory.

We get data back to the CPU with the `.cpu()` method:

```
1 # Cycle through batches  
2 for idx, (data, labels) in enumerate(train_loader):  
3     data = data.to("cuda:0")  
4     optimizer.zero_grad()  
5     y = model(data)  
6     loss = crossentropy(y, labels)  
7     loss.backward()  
8     running_loss += loss.cpu().item()  
9     optimizer.step()  
10  
11     _, predicted = torch.max(y.data.cpu(), 1)
```

# PyTorch vs. Cuda

- For standard functions (Linear, Conv2d, Pooling etc.), PyTorch ships GPU support.
- PyTorch allows to write custom neural network layers (requiring to specify the forward and and the backward pass)
- Custom layers require CUDA programming to run on GPUs.