


Lecture: Deep Learning and Differential Programming

5.2 Should we model or learn?




$$\ddot{\theta}_t = \frac{g \sin \theta_t + \cos \theta_t \left[\frac{-F_t - ml\dot{\theta}_t^2 \sin \theta_t + \mu_c \operatorname{sgn}(\dot{x}_t)}{m_c + m} \right] - \frac{\mu_p \dot{\theta}_t}{ml}}{l \left[\frac{4}{3} - \frac{m \cos^2 \theta_t}{m_c + m} \right]}$$

$$\ddot{x}_t = \frac{F_t + ml[\dot{\theta}_t^2 \sin \theta_t - \ddot{\theta}_t \cos \theta_t] - \mu_c \operatorname{sgn}(\dot{x}_t)}{m_c + m}$$









$$\ddot{\theta}_t = \frac{g \sin \theta_t + \cos \theta_t \left[\frac{-F_t - ml\dot{\theta}_t^2 \sin \theta_t + \mu_c \operatorname{sgn}(\dot{x}_t)}{m_c + m} \right] - \frac{\mu_p \dot{\theta}_t}{ml}}{l \left[\frac{4}{3} - \frac{m \cos^2 \theta_t}{m_c + m} \right]}$$

$$\ddot{x}_t = \frac{F_t + ml[\dot{\theta}_t^2 \sin \theta_t - \dot{\theta}_t \cos \theta_t] - \mu_c \operatorname{sgn}(\dot{x}_t)}{m_c + m}$$






Example:
Robot navigation

Purely geometrical
mapping + planning

Adding semantics



- + Real world solutions
- Simple tasks and reasoning (eg. waypoint navigation)


Purely learned policies
(Deep-RL)

Inductive biases:
Geometry, topology,
stability etc.




- + High-level reasoning
- + Discover tasks from reward
- Does not transfer to the real world

Learn how to create a map (SLAM)




[Tateno, Tombari, Laina, Navab,
CVPR 2017]


Shortest path problems: Dijkstra's algorithm




Shortest path problems: Dijkstra's algorithm




Shortest path problems: Dijkstra's algorithm




Shortest path problems: Dijkstra's algorithm




Shortest path problems: Dijkstra's algorithm



Shortest path problems: Dijkstra's algorithm



Shortest path problems: Dijkstra's algorithm




Active Neural Mapping

- Can we combine traditional path finding (Dijkstra, A*, front propagation etc.) with learning?
- Habitat AI Challenge (@CVPR 2019): Point Goal Task:
- Winner (Chapelot et al.):
 - A learned mapper predicts free space
 - A global policy is learned with RL (reward = coverage)
 - Planning with front-propagation
 - A local policy predicts navigation actions, learned with RL (reward= L_2 distance to global goal).
- No end to end training!
- Easy transfer from coverage objective to point goal (replace global policy by fixed one).

[Chapelot, Gupta, Gandhi, Gupta, Salakhutdinov, 2019 (unpublished)]

Active Neural Mapping


Deep networks trained with RL



[Chapelot, Gupta, Gandhi, Gupta,
Salakhutdinov, 2019 (unpublished)]

Classical planning: shortest
path, not trainable


Active Neural Mapping



[Chapelot, Gupta, Gandhi, Gupta,
Salakhutdinov, 2019 (unpublished)]

Cognitive Mapping and Planning

- Differentiable planner (value iteration networks)
- End to end training, but no RL (imitation learning)



[Gupta, Tolani, Davidson, Levine,
Sukthankar, Malik, CVPR 2017]

Purely geometrical
mapping + planning

Adding semantics




- + Real world solutions
- Simple tasks and reasoning (eg. waypoint navigation)


Purely learned policies
(Deep-RL)

Inductive biases:
Geometry, topology,
stability etc.




- + High-level reasoning
- + Discover tasks from reward
- Does not transfer to the real world





Control Theory



Reinforcement Learning


Computer vision from UAVs



Several problems:

- Computer vision for Control
- Vehicle Control
- Task related computer vision (eg. plant detection)

End to end learning?




Requires a training strategy:

- imitation from human trajectories?
- Reinforcement Learning

Very low explainability


The beauty and the beast



[Kaufmann, Gehrig, Foehn, Ranftl,
Dosovitskiy, Koltun, Scaramuzza¹, 2018]

The beauty and the beast

Vision / perception: DNNs for state estimation and obstacle / door detection.



[Kaufmann, Gehrig, Foehn, Ranftl,
Dosovitskiy, Koltun, Scaramuzza, 2018]

The beauty and the beast

Control: classical model predictive control

$$\begin{aligned} \min_{\mathbf{u}} \int_{t_0}^{t_f} & \left(\bar{\mathbf{x}}_t^\top(t) \mathbf{Q} \bar{\mathbf{x}}_t(t) + \bar{\mathbf{u}}_t^\top(t) \mathbf{R} \bar{\mathbf{u}}_t(t) \right) dt \\ \bar{\mathbf{x}}(t) = \mathbf{x}(t) - \mathbf{x}_r(t) \quad \bar{\mathbf{u}}(t) = \mathbf{u}(t) - \mathbf{u}_r(t) \\ \text{subject to} \quad \mathbf{r}(\mathbf{x}, \mathbf{u}) = 0 \quad & \mathbf{h}(\mathbf{x}, \mathbf{u}) \leq 0. \end{aligned}$$

[Kaufmann, Gehrig, Foehn, Ranftl,
Dosovitskiy, Koltun, Scaramuzza, 2018]

Neural Lander


- A drone is controlled with a PID controller based on a model of the drone
- Unknown disturbances (unsteady airflow) is learned by a deep neural network and supervised learning

$$\dot{\mathbf{p}} = \mathbf{v},$$

$$m\dot{\mathbf{v}} = m\mathbf{g} + R\mathbf{f}_u + \mathbf{f}_a, \quad (1a)$$


$$\dot{R} = RS(\boldsymbol{\omega}),$$

$$J\dot{\boldsymbol{\omega}} = J\boldsymbol{\omega} \times \boldsymbol{\omega} + \boldsymbol{\tau}_u + \boldsymbol{\tau}_a, \quad (1b)$$



Estimated from state and
control inputs

Neural Lander



[Shi et al., 2019]