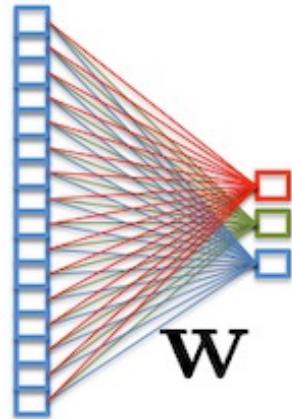


5IF - Deep Learning et Programmation Différentielle

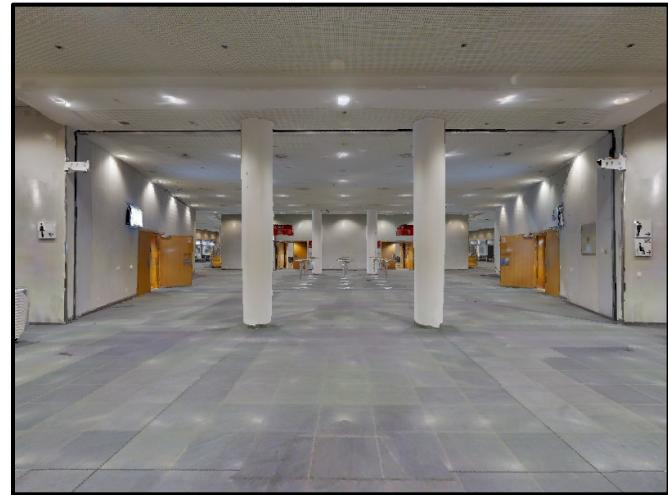
5.3 Robot Navigation



Robotics



Perception



Navigation



HCI



Compagnon robots

Visual navigation and spatial reasoning



Office space



Homes



Hospitals



Edward
Beeching



Jilles
Dibangoye



Olivier
Simonin



Christian
Wolf

Purely geometrical
mapping + planning

Adding semantics



- + *Real world solutions*
- *Simple tasks and reasoning (eg. waypoint navigation)*

Purely learned policies
(Deep-RL)

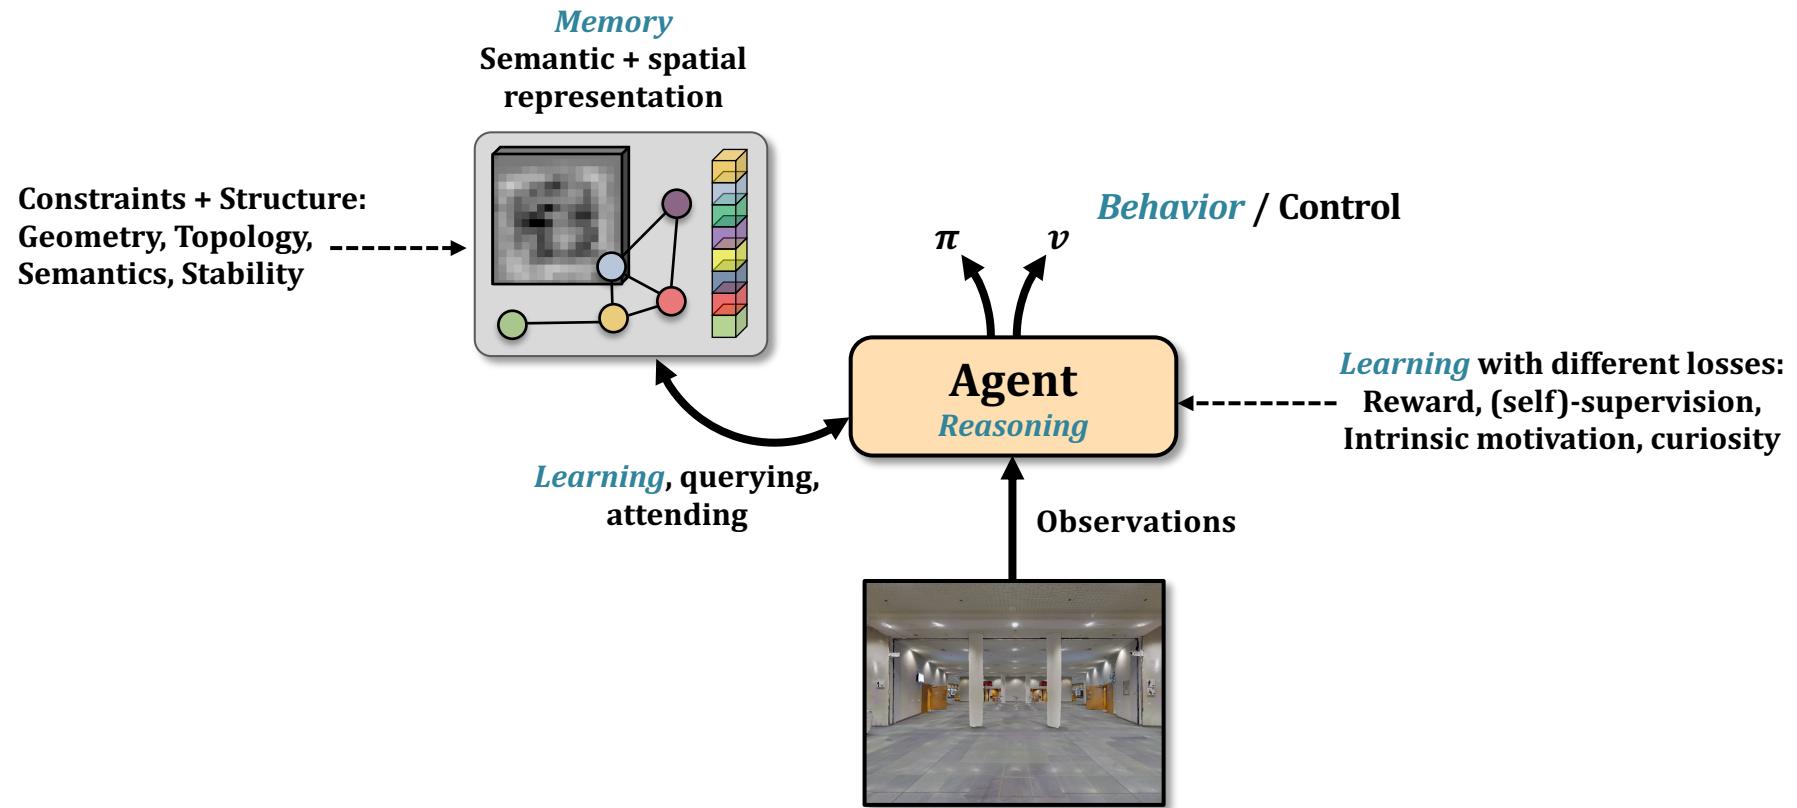
Inductive biases:
Geometry, topology,
stability etc.

- + *High-level reasoning*
- + *Discover tasks from reward*
- *Does not transfer to the real world*



Question:

What kind of inductive bias can we create for learning navigation?

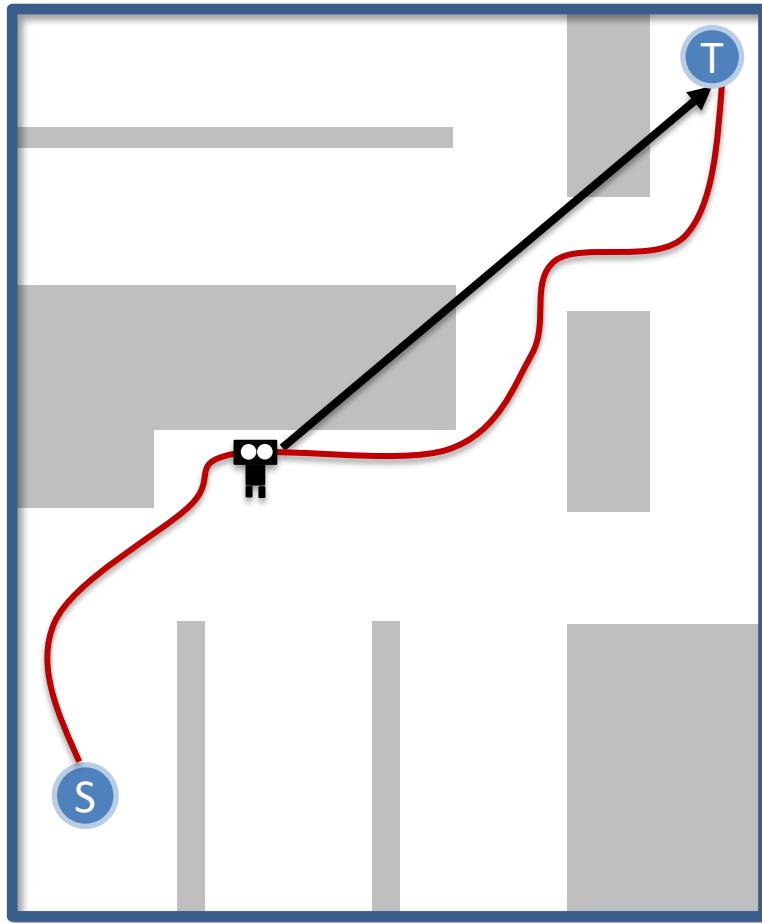


Find my keys!

Can we learn a spatial representation which is richer than navigational space?

Can we learn it from reward alone?

Task: PointGoal (+GPS)



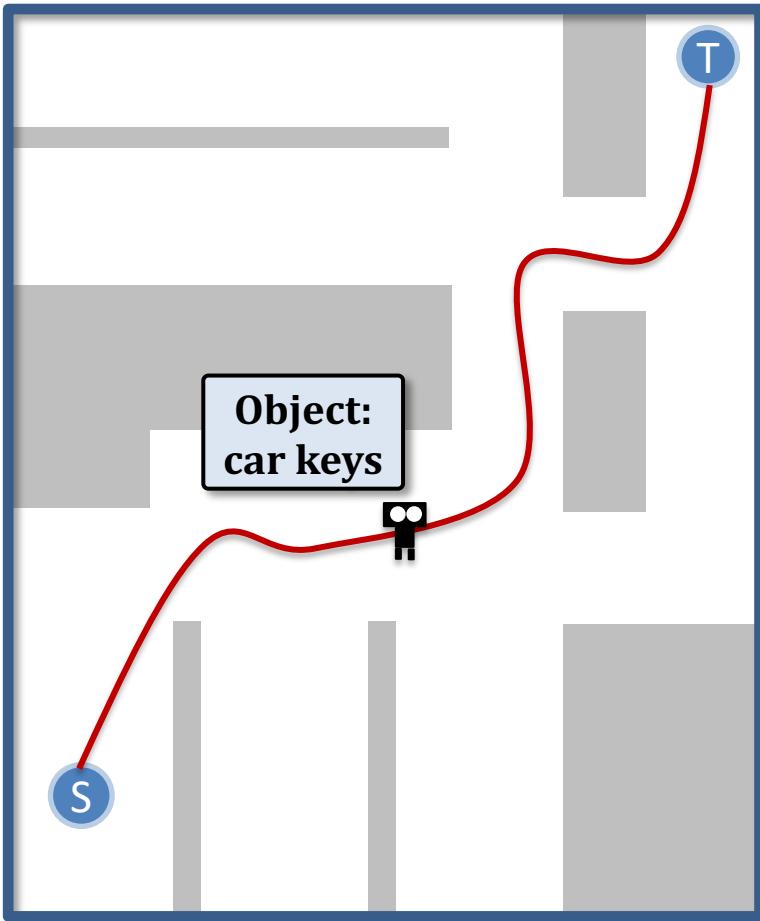
Task:

Find a target given coordinates,
receive a direction vector et each
step

Required reasoning:

- Recognize free navigation space
- Follow the direction vector
- Learn how to overcome
obstacles (difference between
Euclidean and geodesic path)

Task: ObjectNav



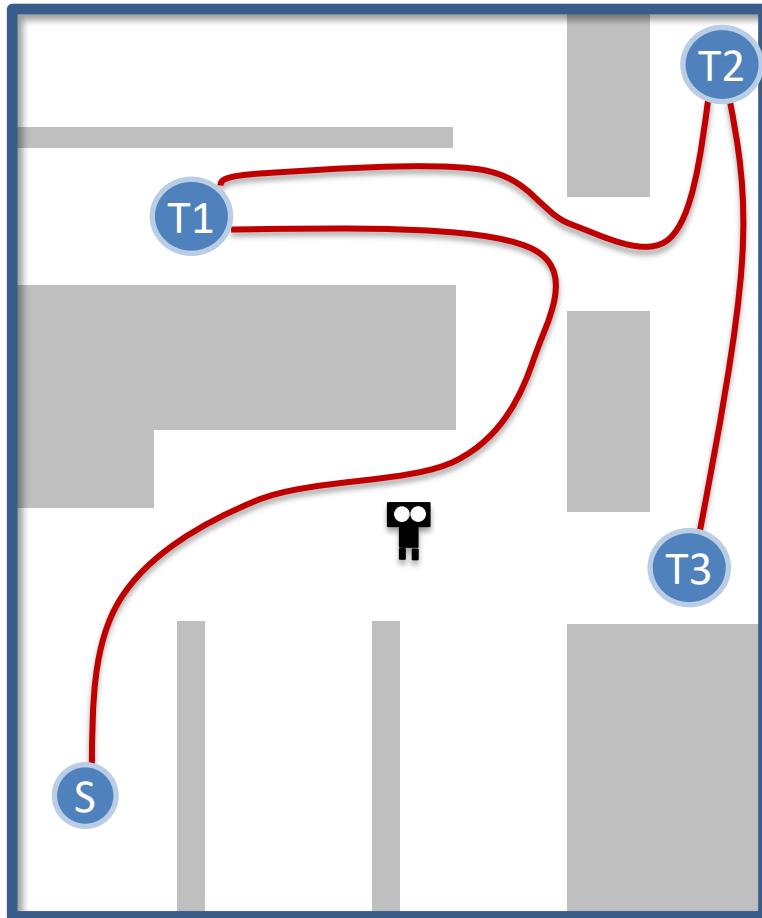
Task:

Find a target object given its object class.

Required reasoning:

- Recognize free navigation space
- *Explore the environment*
- *Recognize the object when seen and move towards it.*

Task: K-item scenario



Task:

Navigate to a list of objects sequentially in the right order.

Required reasoning:

- Recognize free navigation space
- Explore the environment
- *Map an object if I need to find it later (!)*
- Recognize the object when seen and move towards it.

[Beeching, Dibangoye, Simonin, Wolf, ECML-PKDD 2020]

[Beeching, Dibangoye, Simonin, Wolf, ICPR 2020]

What do we want the model to learn?

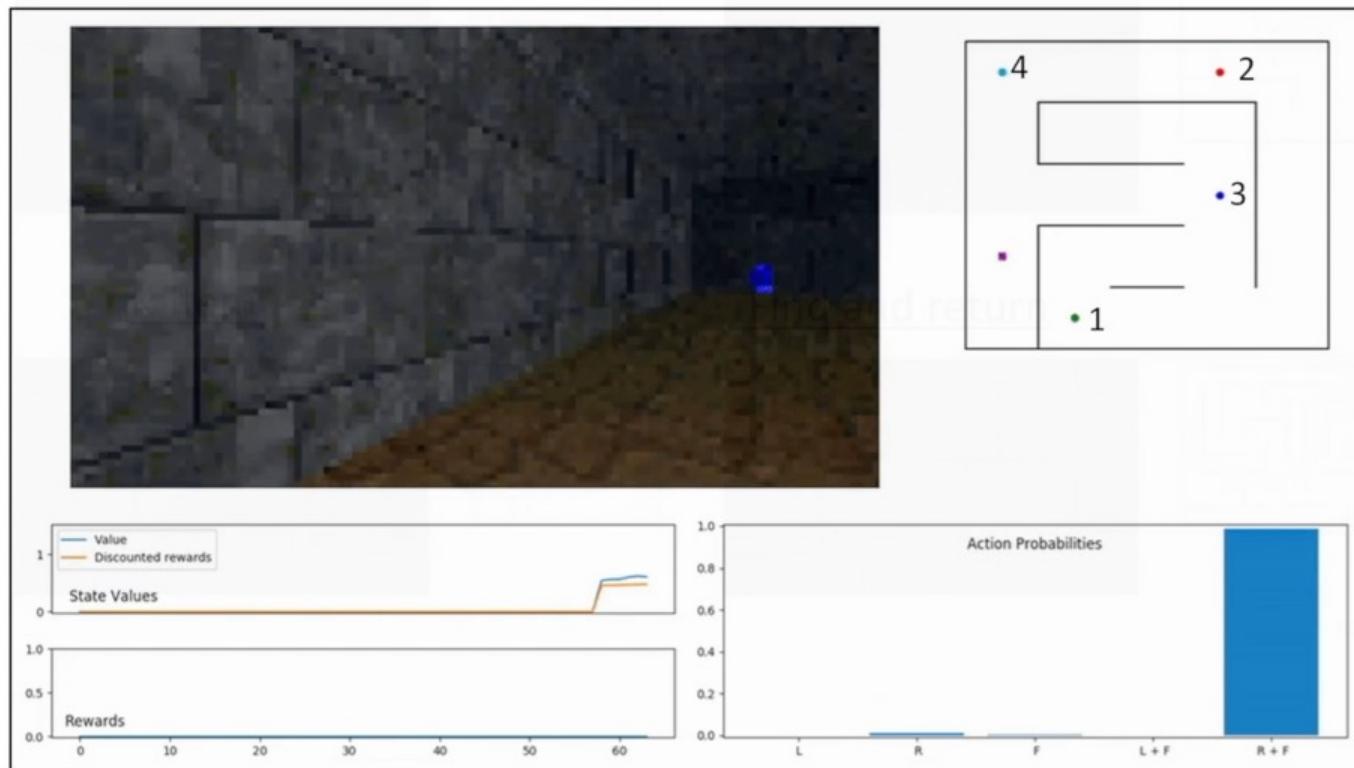
Generalization to new environments requires:

	Agent memory	Network parameters
Information on the (seen) environment, e.g. position of a couch	✓	✗
Positions of objects placed in the environment	✓	✗
Regularities of the environment (eg. bath tubs are in bathrooms; toilets are accessible from an aisle, not the living room)	✗	✓
Object affordances	✓	✓

The task formulation decides how learned information is stored!

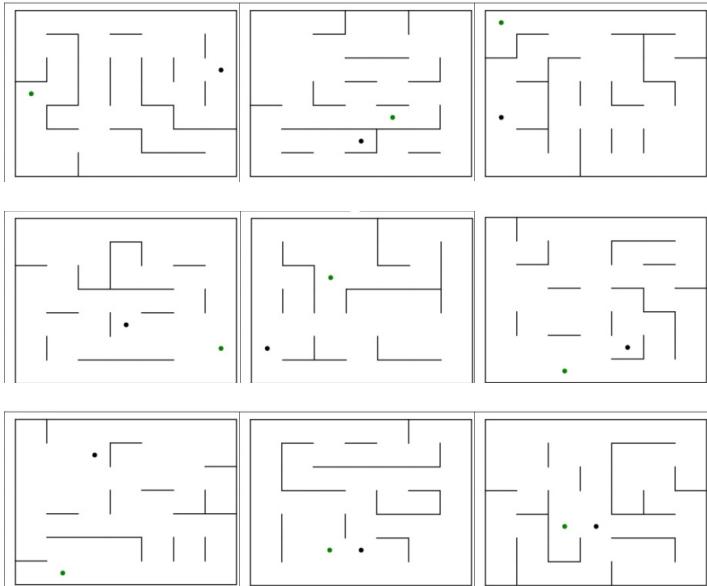
Scenarios: Ordered K-item

- Collect K items in a specific order
- Tests: Vision & memory

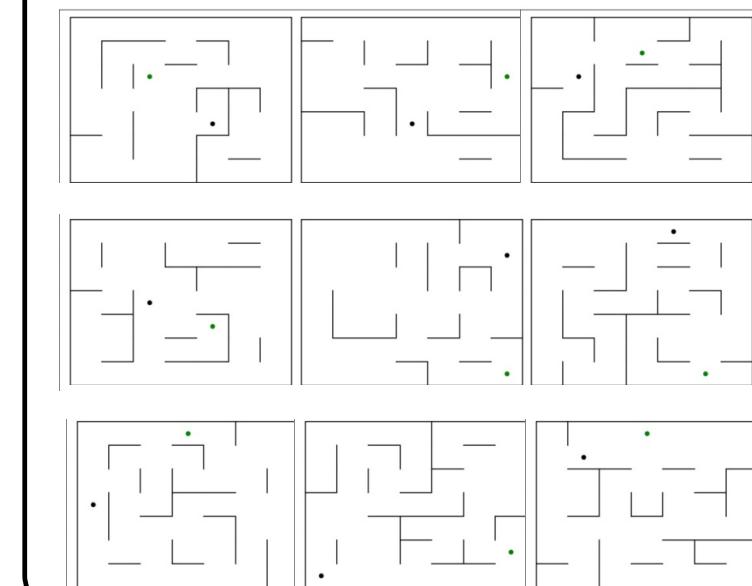


Generalization to unseen mazes

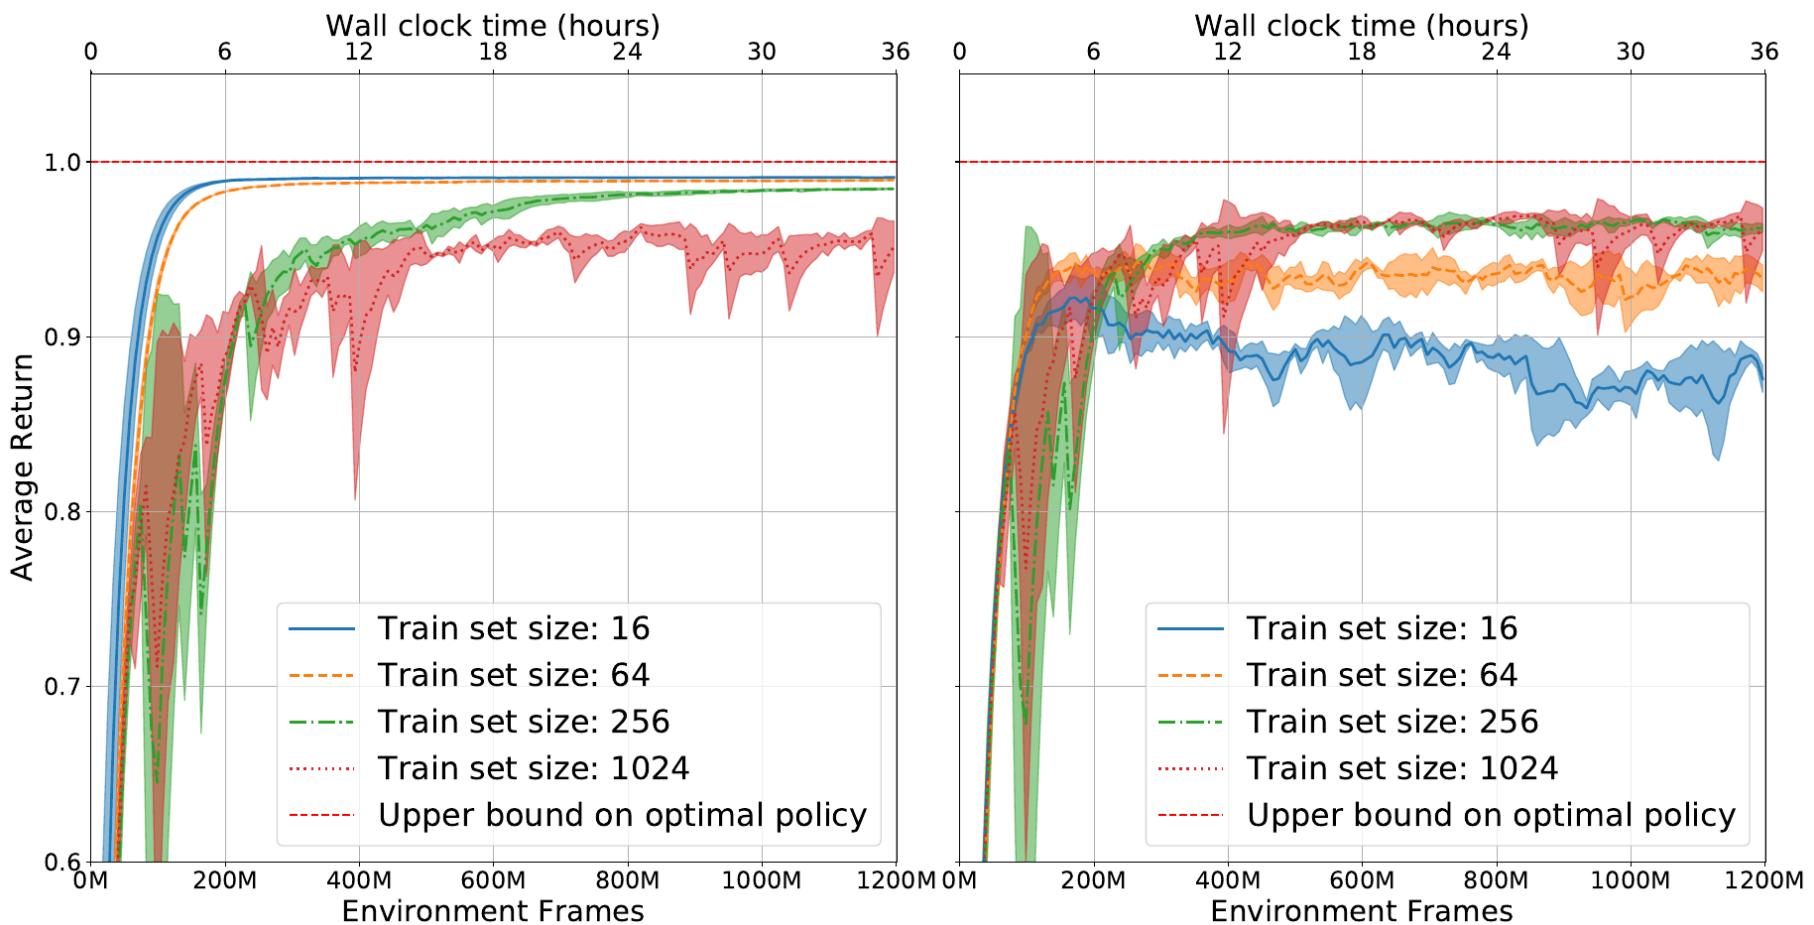
N - Training configurations



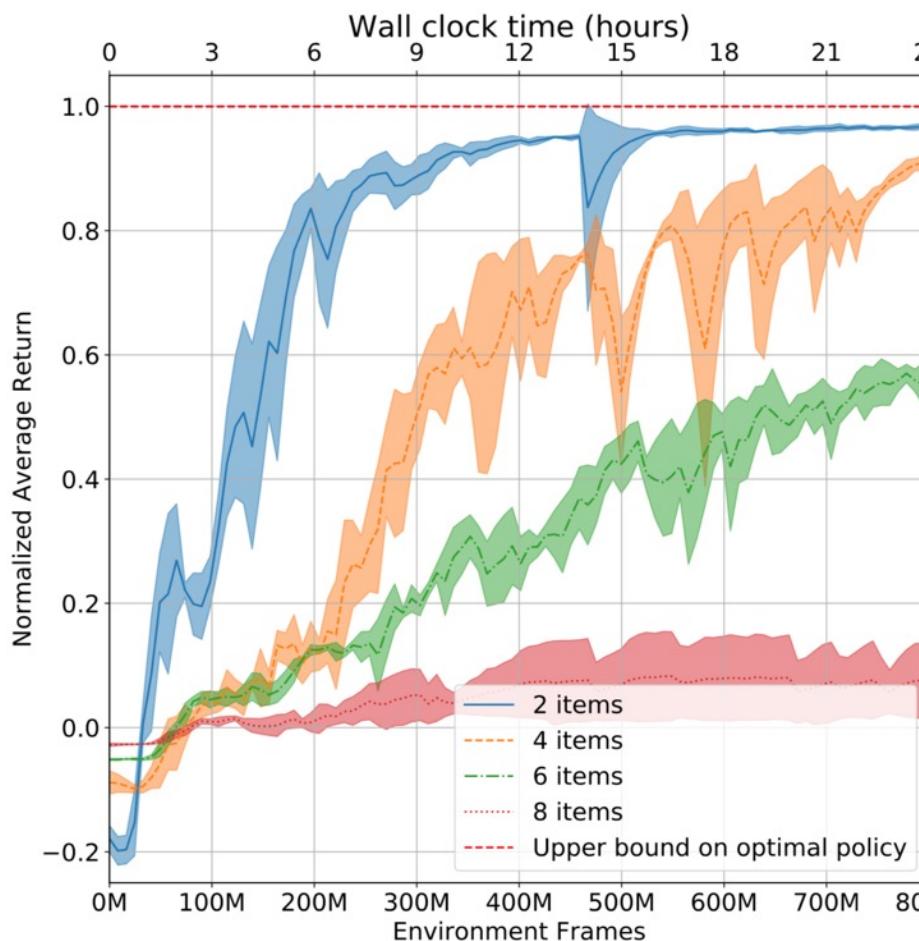
K - Testing configurations



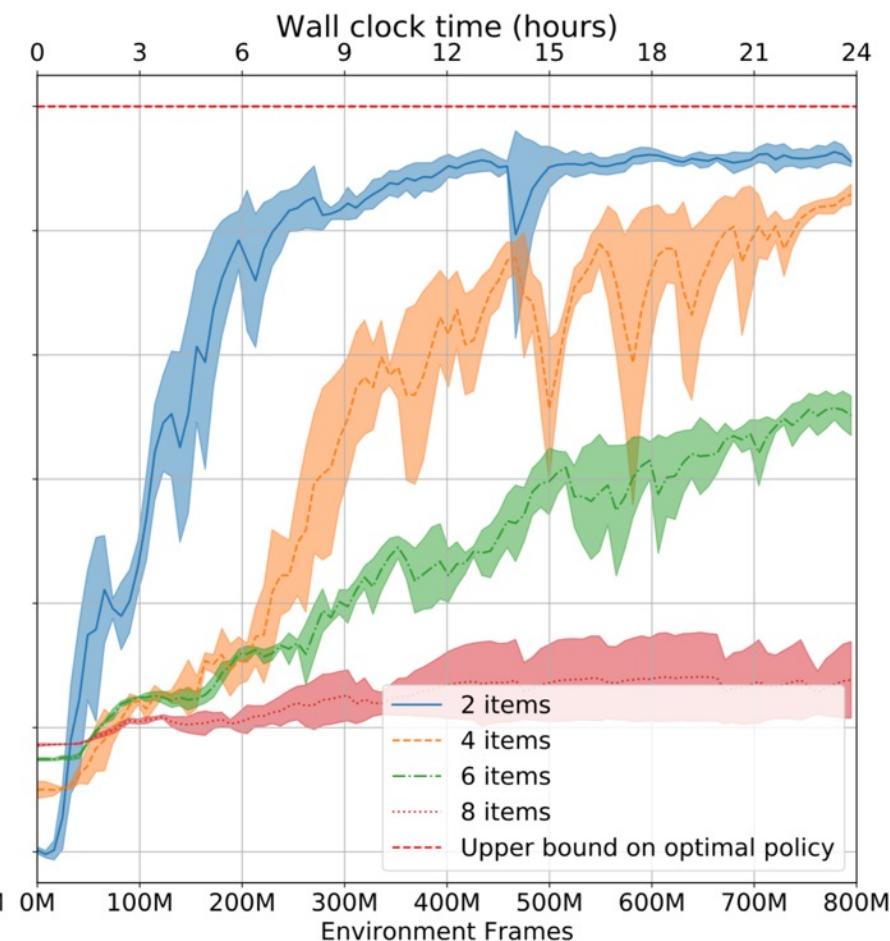
Generalization to unseen mazes



K-item scenario: recurrent agent (A3C)

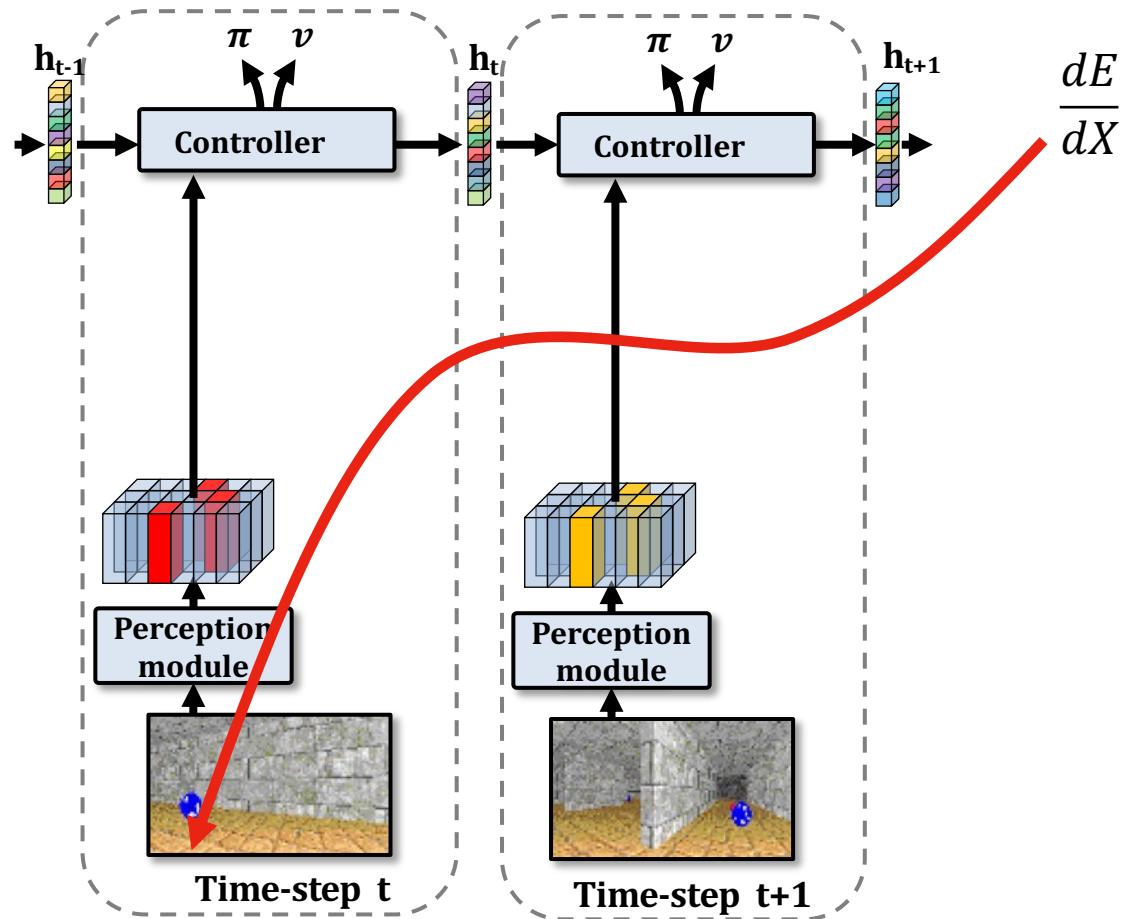


Training scenarios

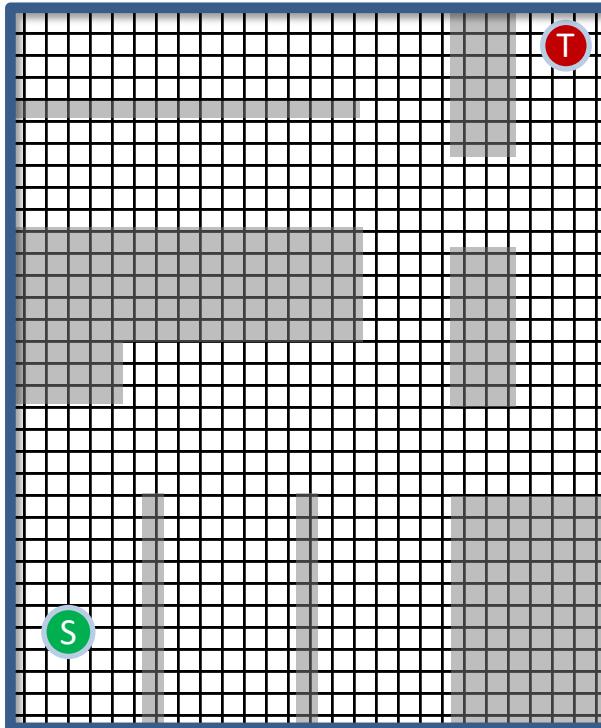


Testing scenarios

The unstructured recurrent baseline

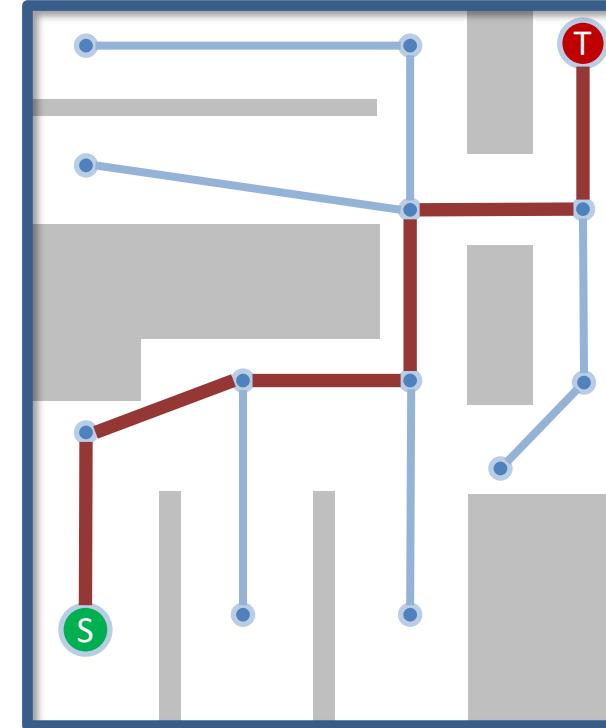


Spatial maps in robotics



Metric map
(=2D or 3D Grid)

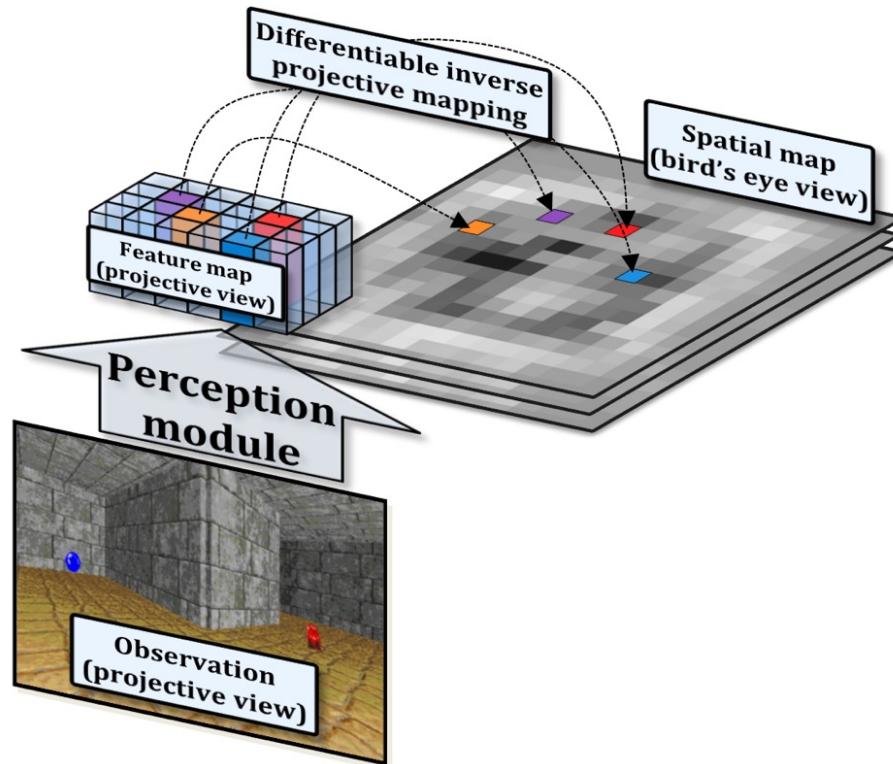
Beeching, Dibangoye, Simonin, Wolf,
*EgoMap: Projective mapping and structured
egocentric memory for Deep RL*,
ECML-PKDD 2020



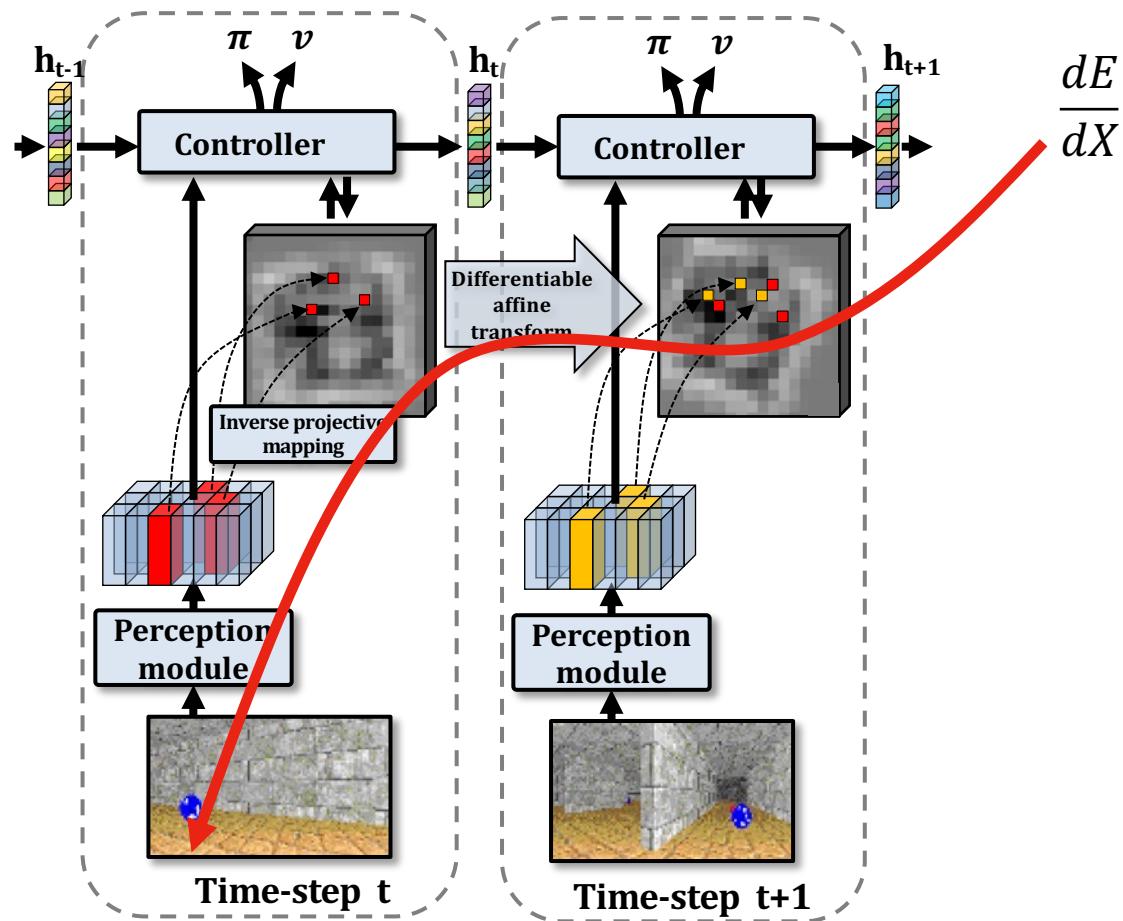
Topological map
(=Graph)

Beeching, Dibangoye, Simonin, Wolf,
*Learning to plan with
uncertain topological maps*,
ECCV 2020

Projective mapping



Learning agent control/behavior



A single forward pass

1. Transform the map to the agent's egocentric frame of reference

$$\hat{M}_t = \text{Affine}(M_{t-1}, dx_t, dy_t, d\phi_t)$$

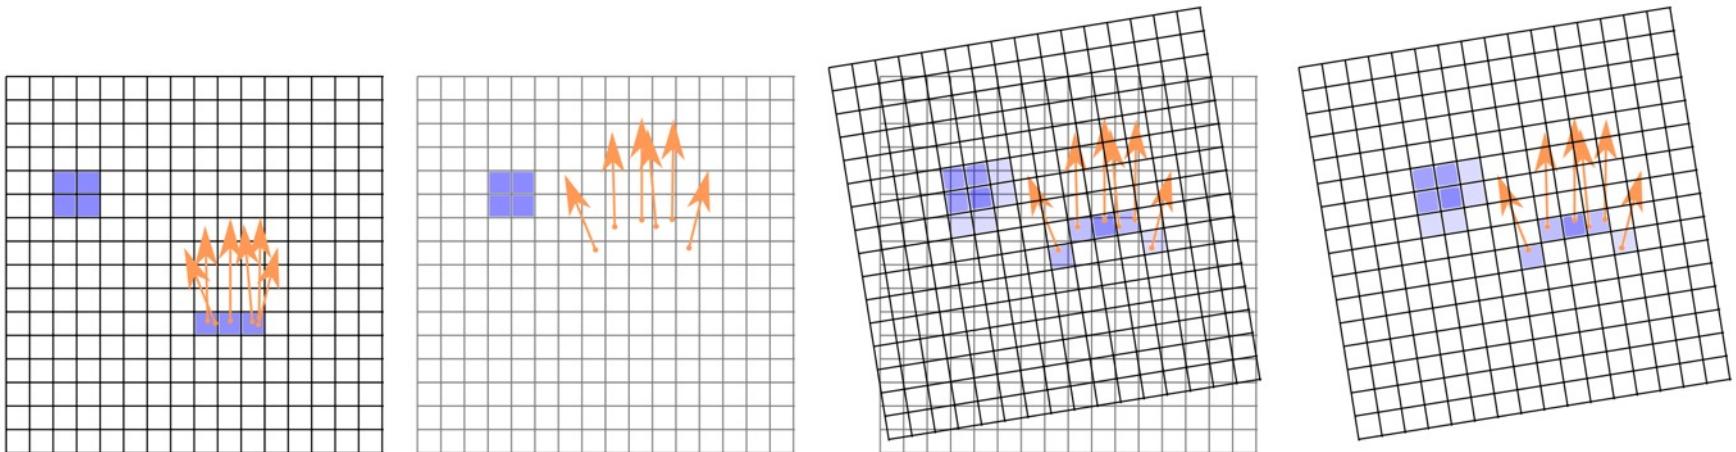
2. Update the map to include new observations

$$\tilde{M}_t = \text{InverseProject}(s_t, D_t) \quad M'_t = \text{Combine}(\hat{M}_t, \tilde{M}_t)$$

3. Perform a global read and attention based read,

$$r_t = \text{Read}(M'_t) \quad c_t = \text{Context}(M'_t, s_t, r_t)$$

Occupancy Grid vs. Egomap



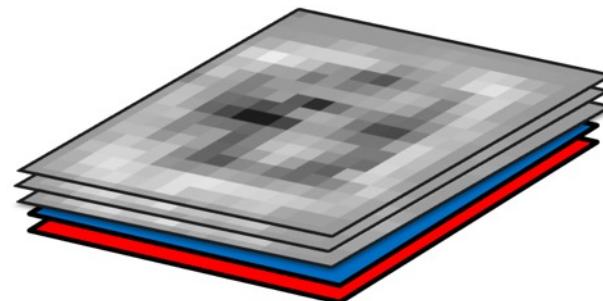
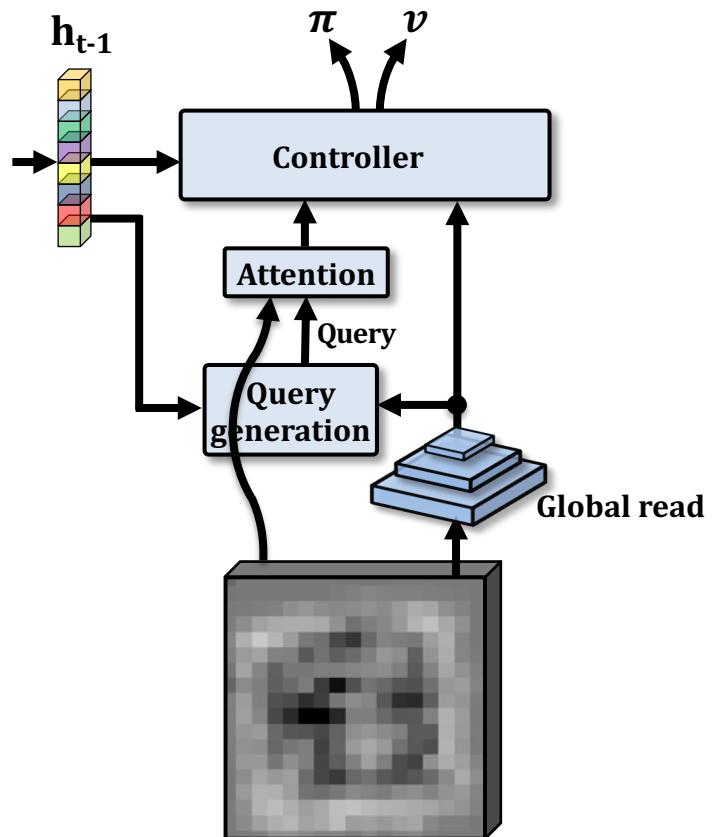
[Rummelhard, Negre,
Laugier, 2015]



[Beeching, Dibangoye,
Simonin, Wolf,
ECML-PKDD 2020]

Querying spatial memory

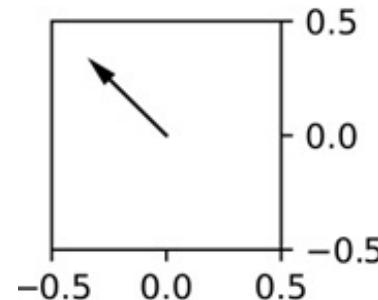
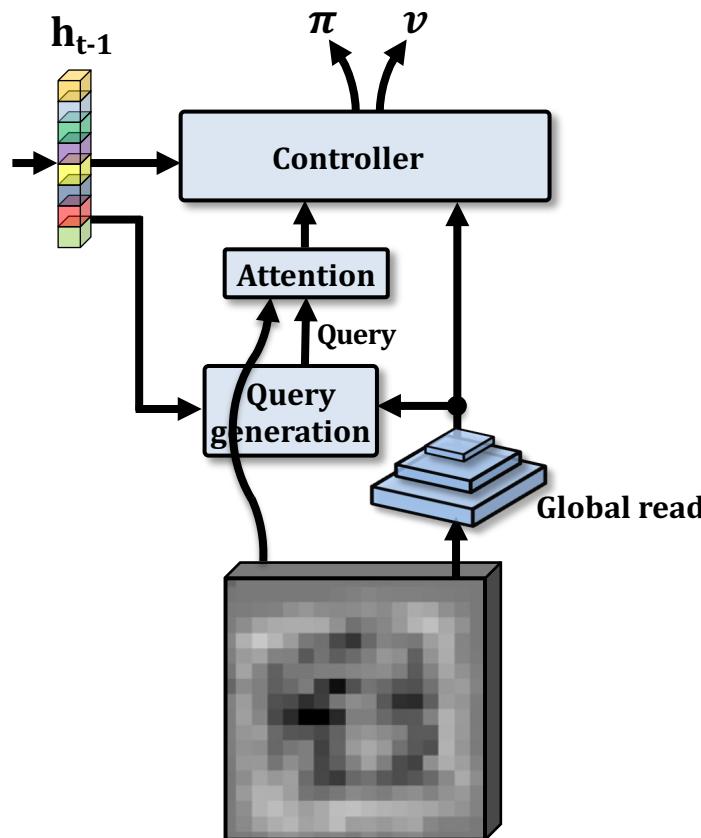
The network learns to query for specific content



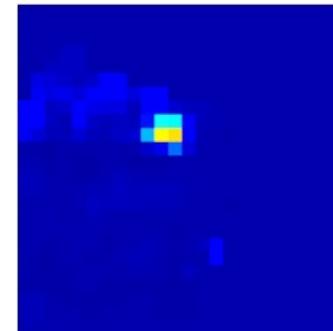
Ease localization with
coordinate planes

Querying spatial memory

The network learns to query for specific content

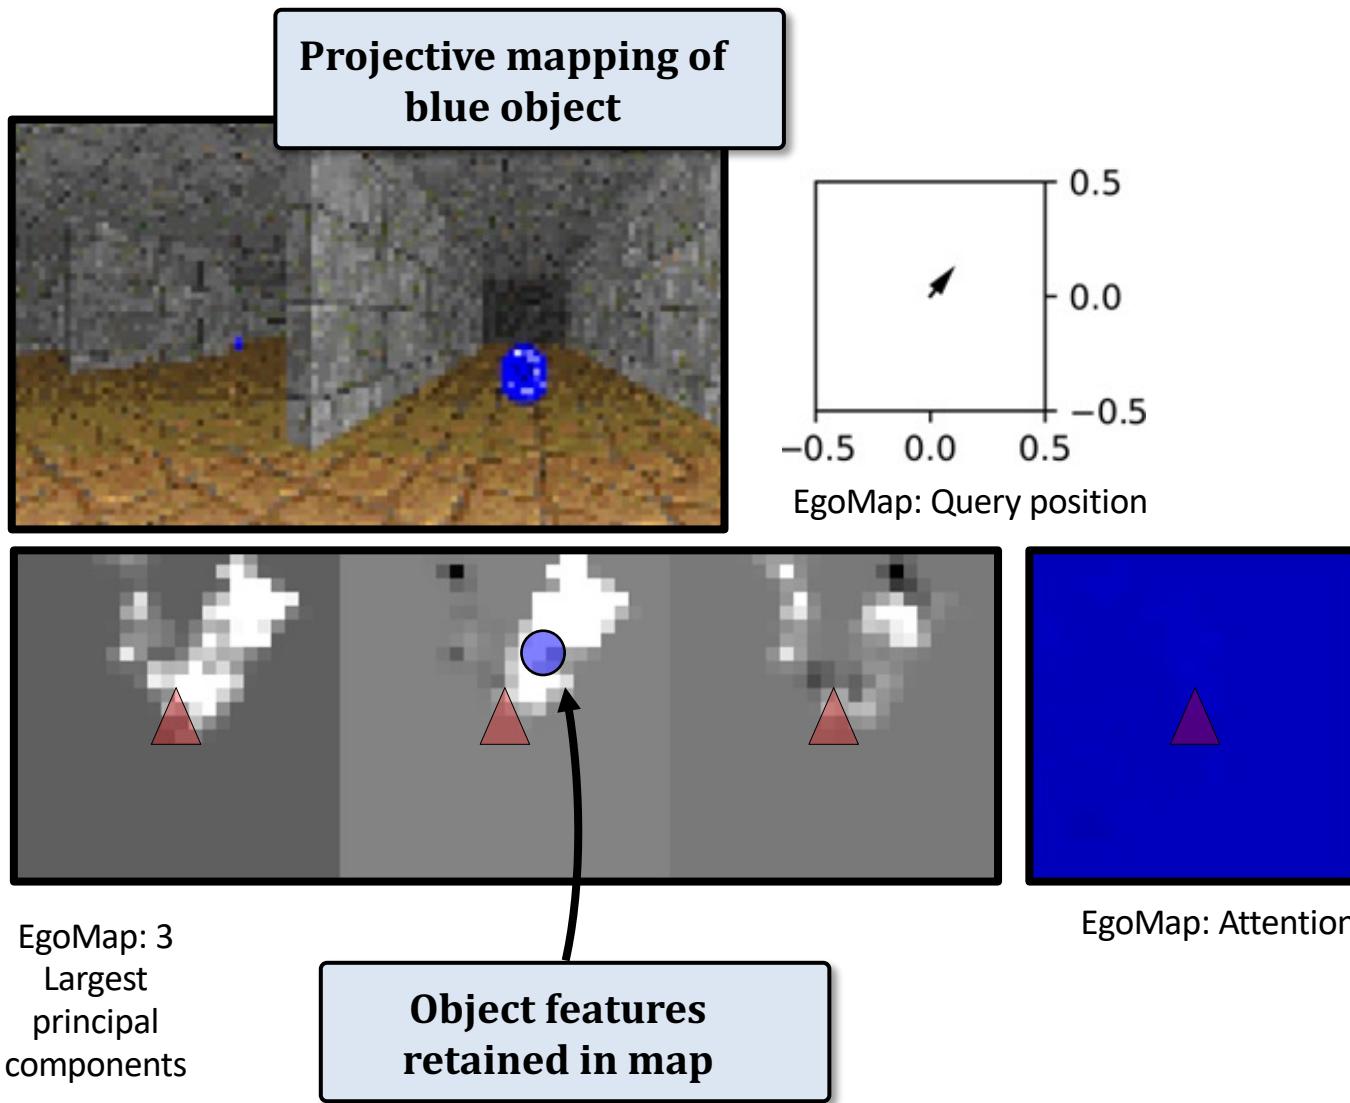


Attention:
Spatial
direction



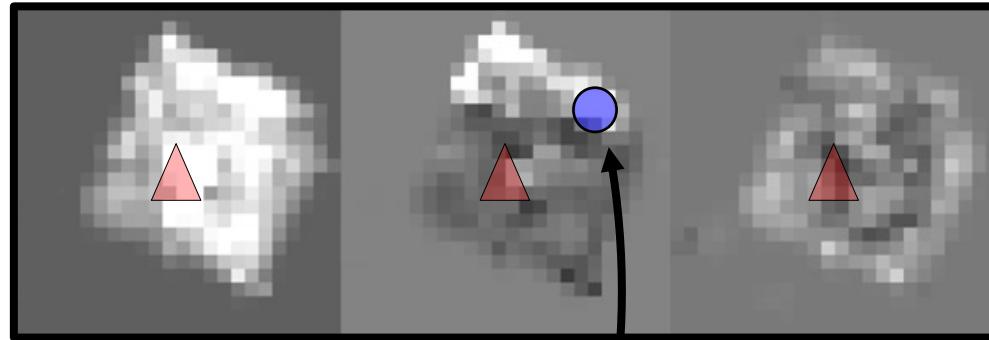
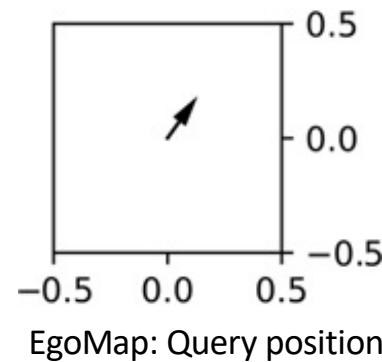
Attention
distribution

6 item scenario: time-step 005

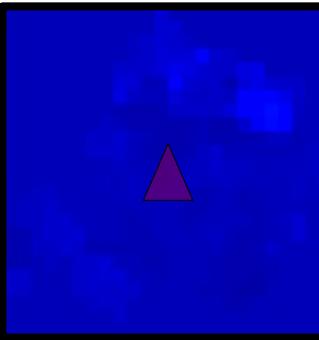


**Object features
retained in map**

6 item scenario: time-step 105

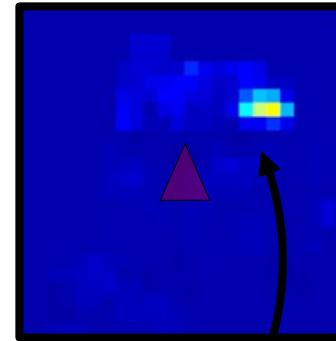
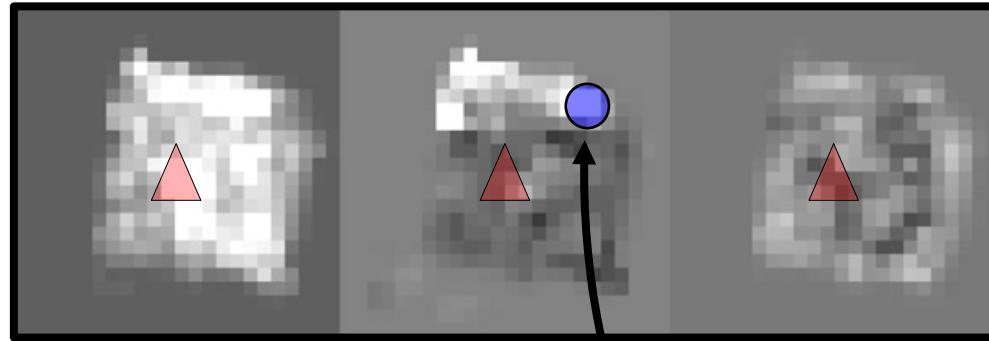
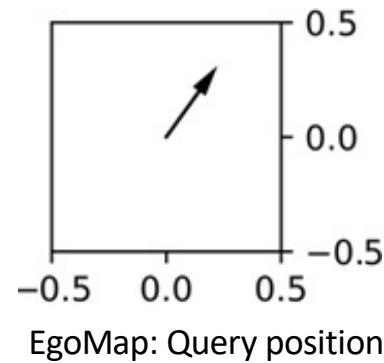


EgoMap: 3
Largest
principal
components



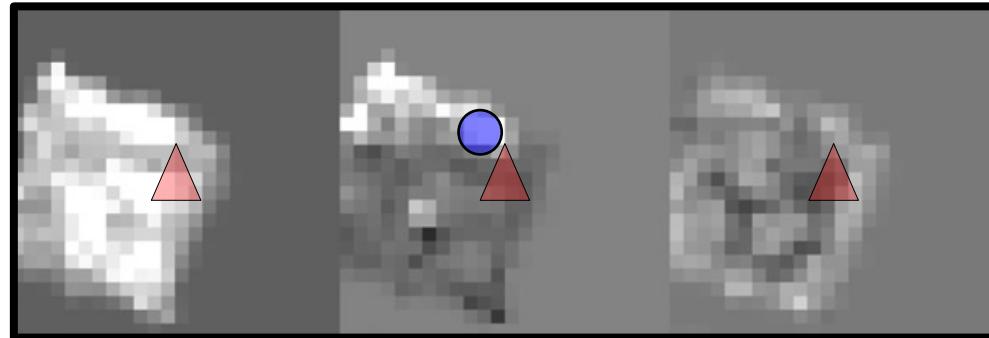
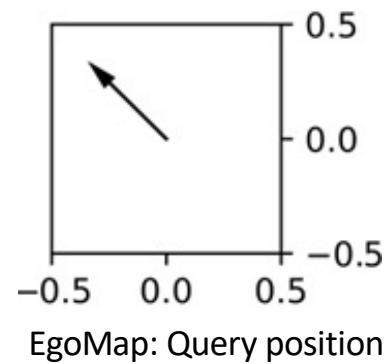
**Object features
retained in map**

6 item scenario: time-step 108

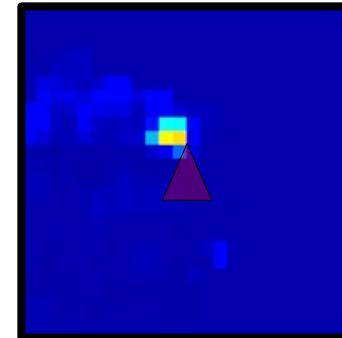


Collection of object n-1 triggers attention to object n

6 item scenario: time-step 134

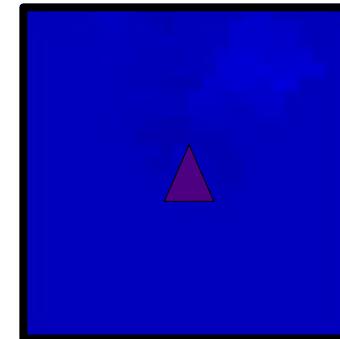
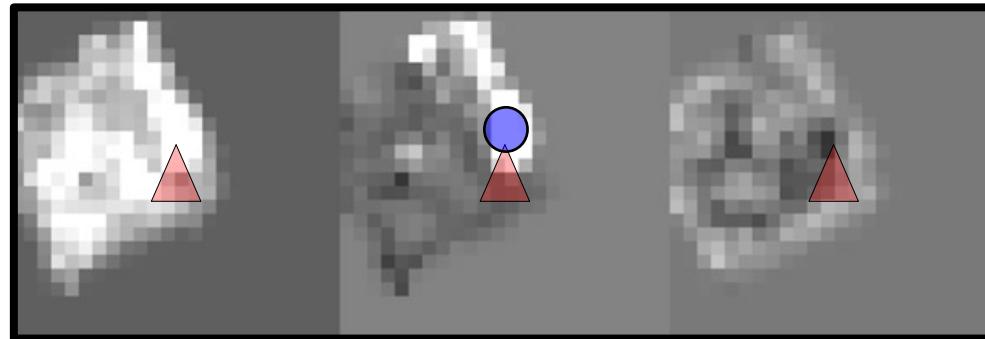
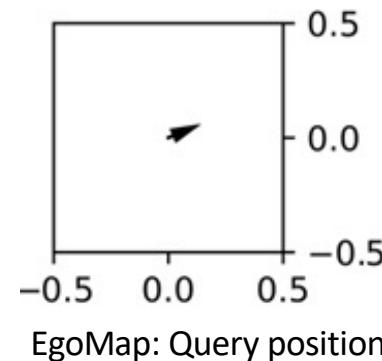


EgoMap: 3
Largest
principal
components



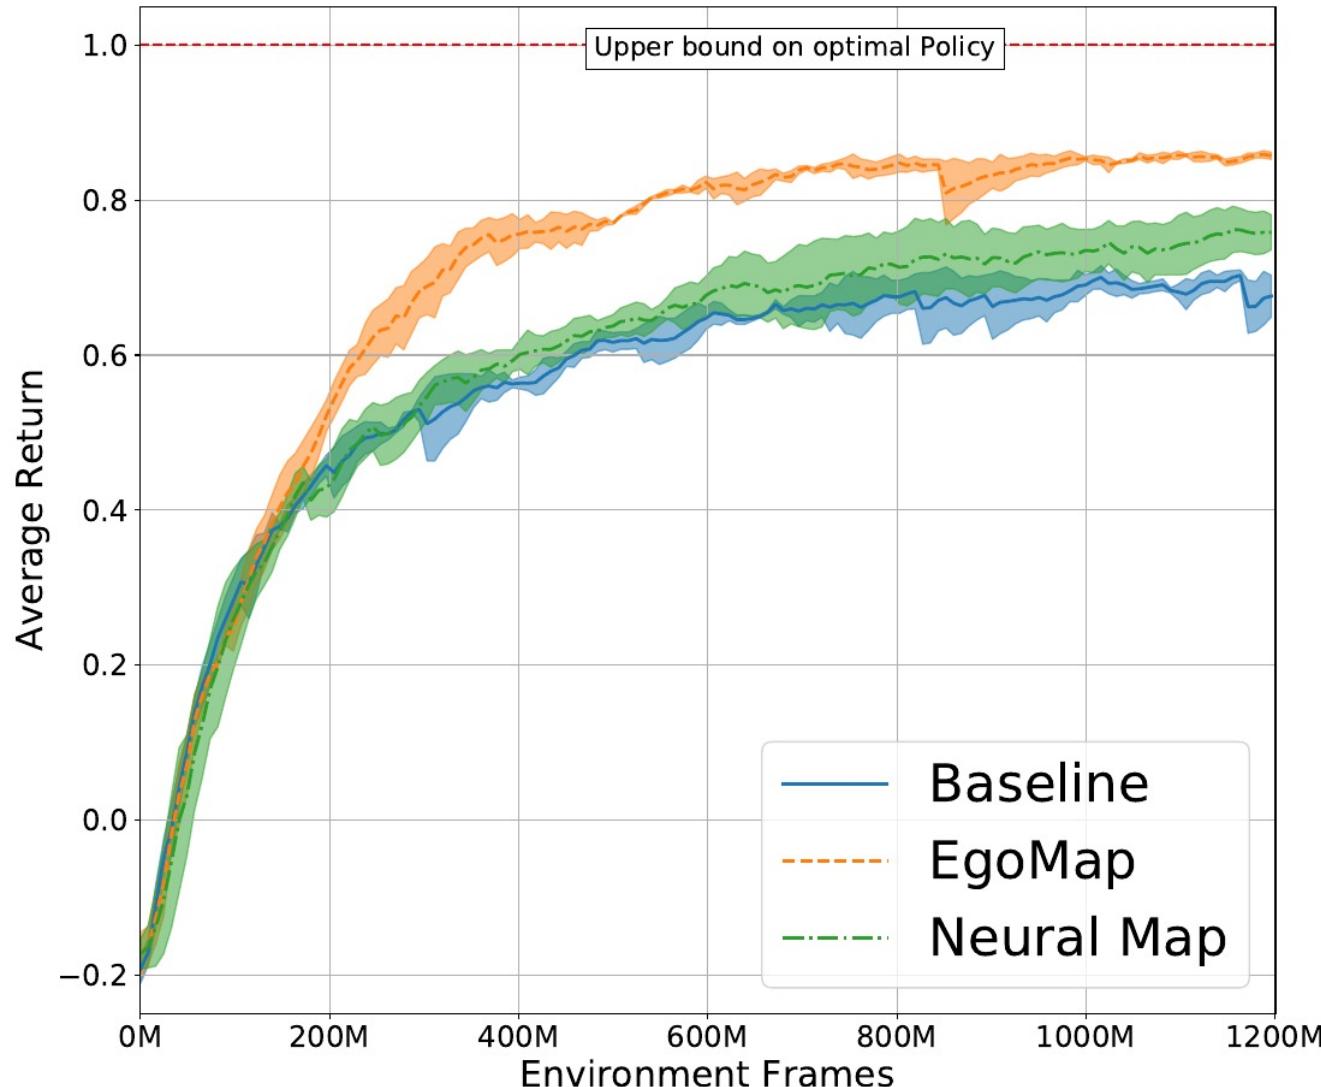
EgoMap: Attention

6 item scenario: time-step 140



**When the object is not
occluded, the agent
does not attend to it**

Results



Quantitative results

Agent	Scenario							
	4 item		6 item		Find and Return		Labyrinth	
	Train	Test	Train	Test	Train	Test	Train	Test
Random	-0.179	-0.206	-0.21	-0.21	-0.21	-0.21	-0.115	-0.086
Baseline	2.341 ± 0.026	2.266 ± 0.035	2.855 ± 0.164	2.545 ± 0.226	0.661 ± 0.003	0.633 ± 0.027	0.73 ± 0.02	0.694 ± 0.009
Neural Map	2.339 ± 0.038	2.223 ± 0.040	2.750 ± 0.062	2.465 ± 0.034	0.825 ± 0.070	0.723 ± 0.026	0.769 ± 0.042	0.706 ± 0.018
EgoMap	2.398 ± 0.014	2.291 ± 0.021	3.214 ± 0.007	2.801 ± 0.048	0.893 ± 0.007	0.848 ± 0.017	0.753 ± 0.002	0.732 ± 0.016
Optimum	2.5	2.5	3.5	3.5	1	1	1	1



Results: ablation study

Ablation	Train	Test
Baseline	0.668 ± 0.028	0.662 ± 0.036
No global read	0.787 ± 0.007	0.771 ± 0.029
No query	0.838 ± 0.003	0.811 ± 0.013
No query temperature	0.845 ± 0.014	0.815 ± 0.019
No query position	0.839 ± 0.007	0.814 ± 0.008
Cosine query	0.847 ± 0.011	0.814 ± 0.017
L1 query	0.851 ± 0.014	0.828 ± 0.011

Quantitative analysis: Ablations

Ablation	Train	Test	Noise std.	Average return
Baseline	0.668 ± 0.028	0.662 ± 0.036	0.00	0.702 ± 0.026
No global read	0.787 ± 0.007	0.771 ± 0.029	0.02	0.686 ± 0.027
No query	0.838 ± 0.003	0.811 ± 0.013	0.04	0.669 ± 0.031
No query temperature	0.845 ± 0.014	0.815 ± 0.019	0.06	0.623 ± 0.018
No query position	0.839 ± 0.007	0.814 ± 0.008	0.08	0.585 ± 0.017
EgoMap	0.847 ± 0.011	0.814 ± 0.017	0.10	0.568 ± 0.023
EgoMap (L1 query)	0.851 ± 0.014	0.828 ± 0.011	0.12	0.575 ± 0.020
			0.14	0.578 ± 0.023
			0.16	0.546 ± 0.017
			0.20	0.537 ± 0.031
			Baseline	0.527 ± 0.047

Self-supervision for Deep-RL

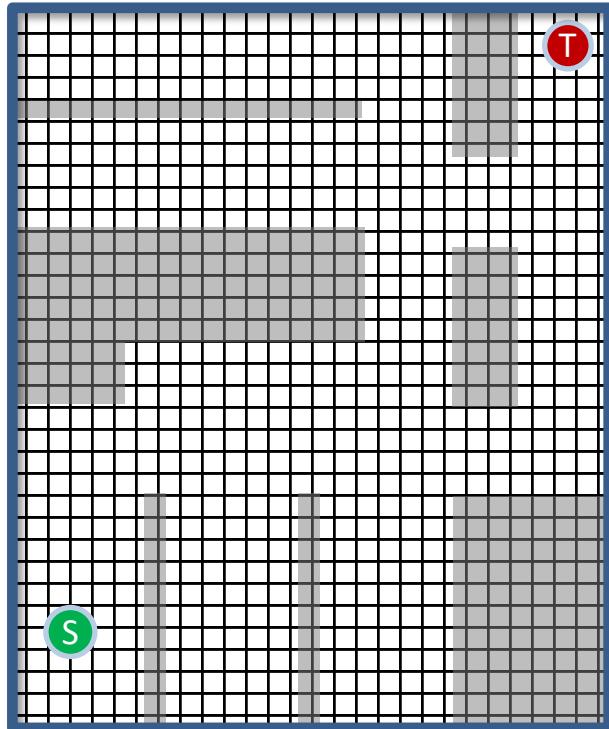
1st ranked at Multi-ON Challenge (CVPR 2021)



Rank	Team	Progress	PPL	Success	SPL
1	Lyon	67	44	55	35
2	SGoLAM	64	38	52	32
3	VIMP	57	36	41	26

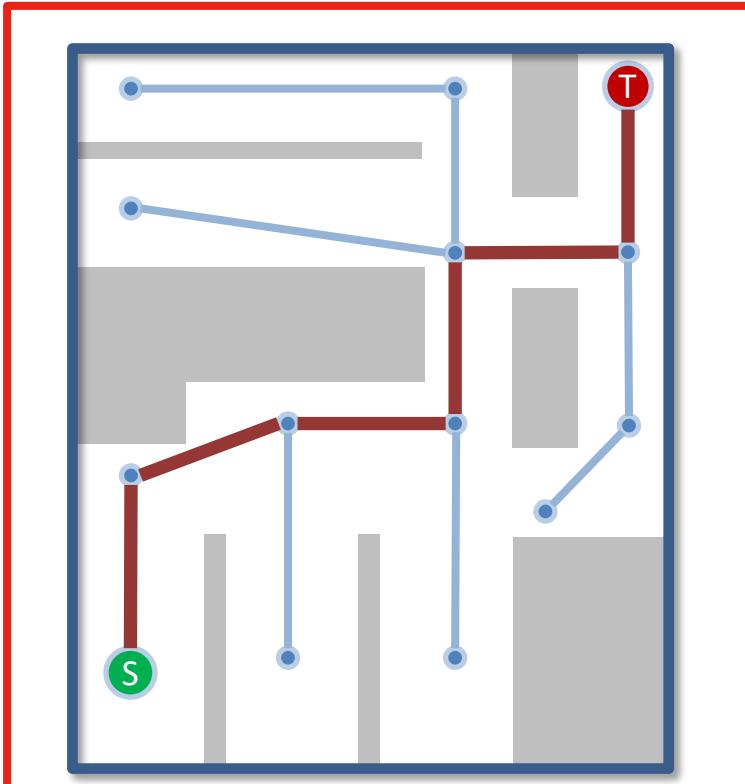
1st place: Team Lyon, Pierre Marza, Laetitia Matignon, Olivier Simonin, Christian Wolf, Learning mapping and spatial reasoning with auxiliary tasks.

Spatial maps in robotics



Metric map
(=2D or 3D Grid)

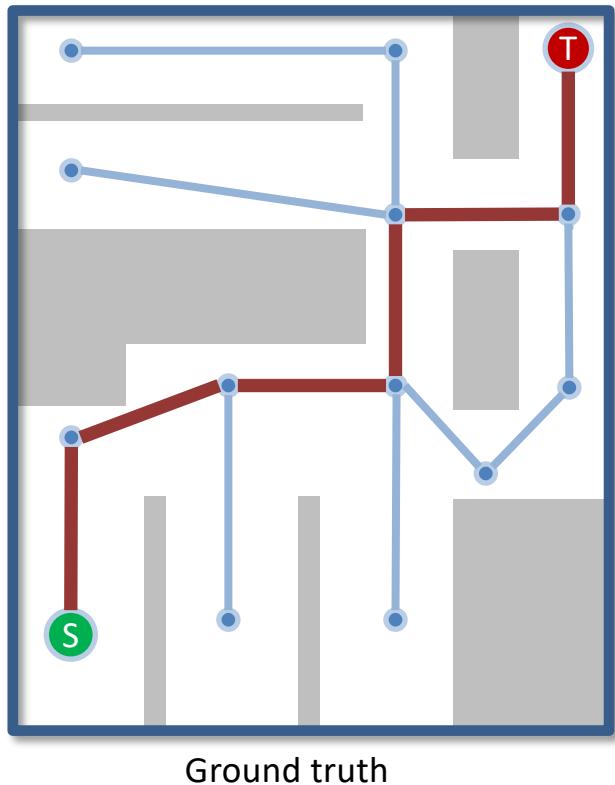
Beeching, Dibangoye, Simonin, Wolf,
*EgoMap: Projective mapping and structured
egocentric memory for Deep RL*,
ECML-PKDD 2020



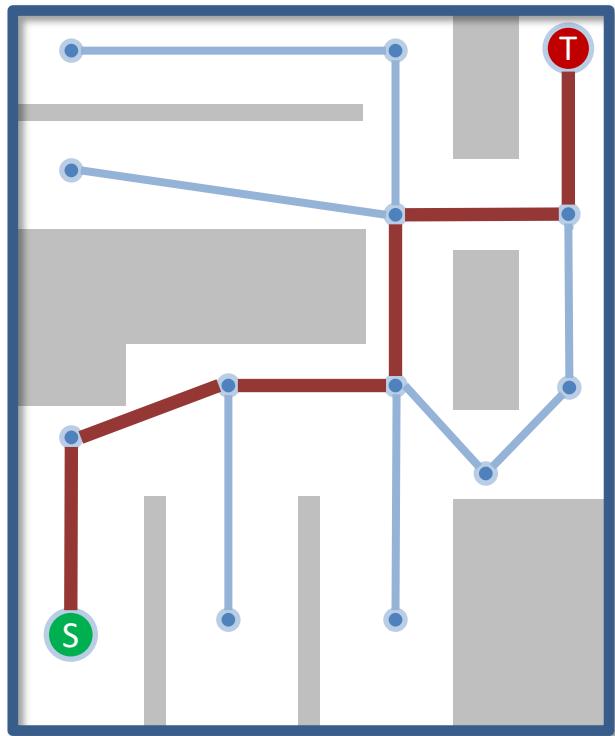
Topological map
(=Graph)

Beeching, Dibangoye, Simonin, Wolf,
*Learning to plan with
uncertain topological maps*,
ECCV 2020

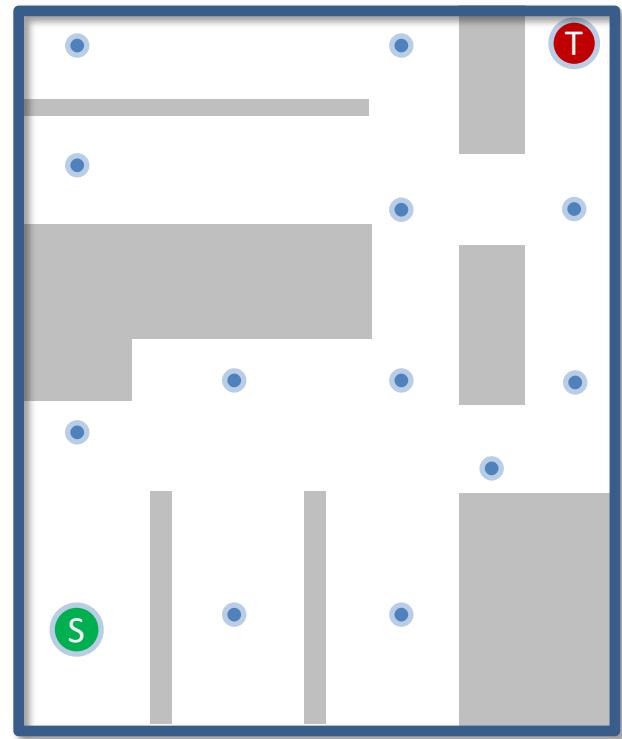
Uncertain topological maps



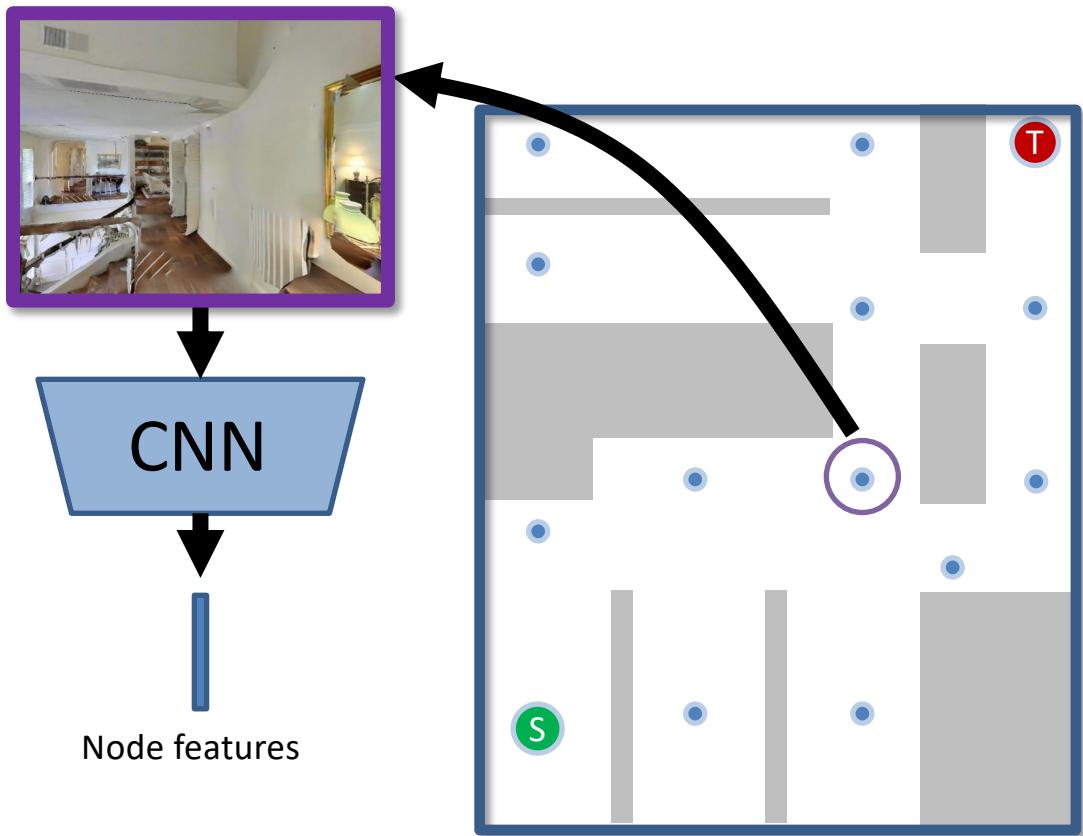
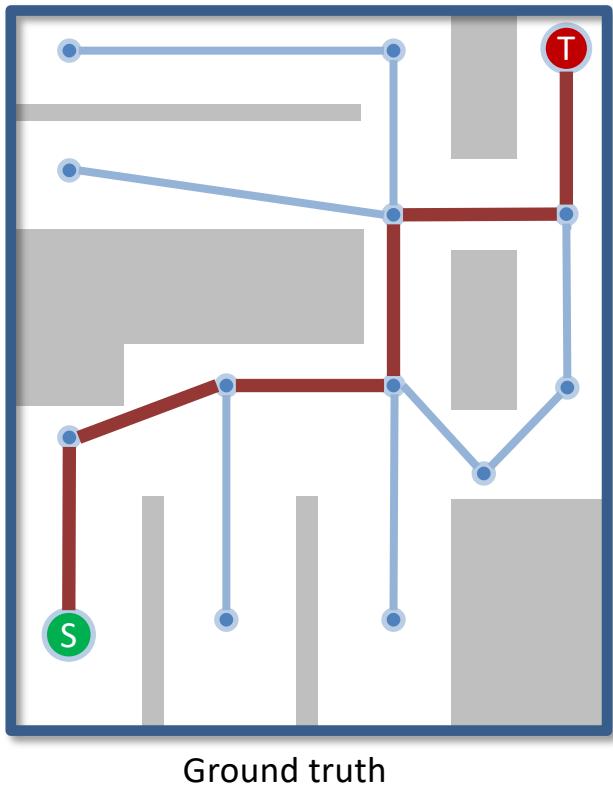
Uncertain topological maps



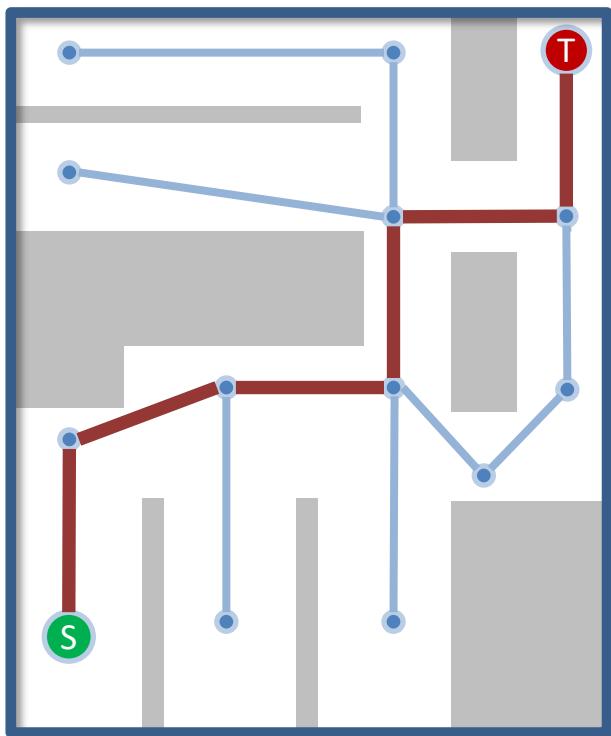
Ground truth



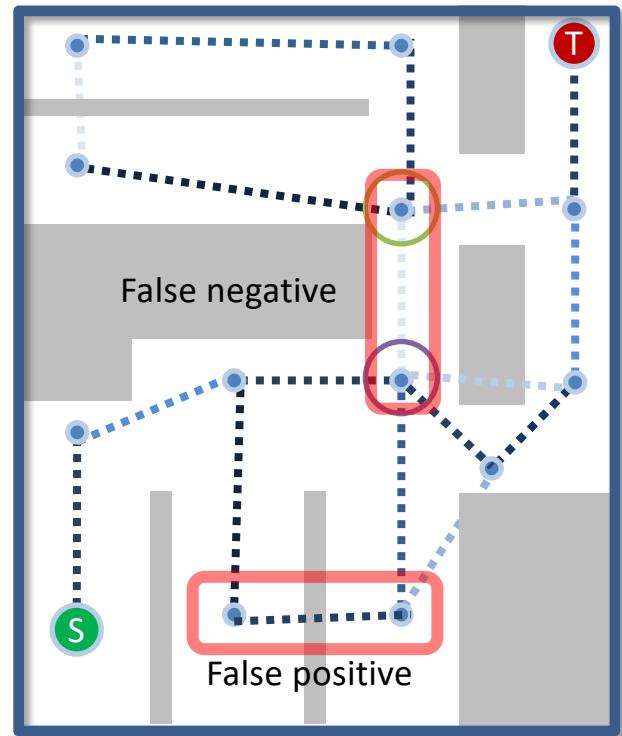
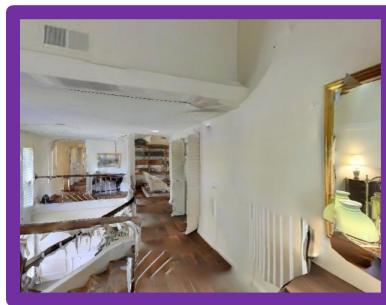
Uncertain topological maps



Uncertain topological maps

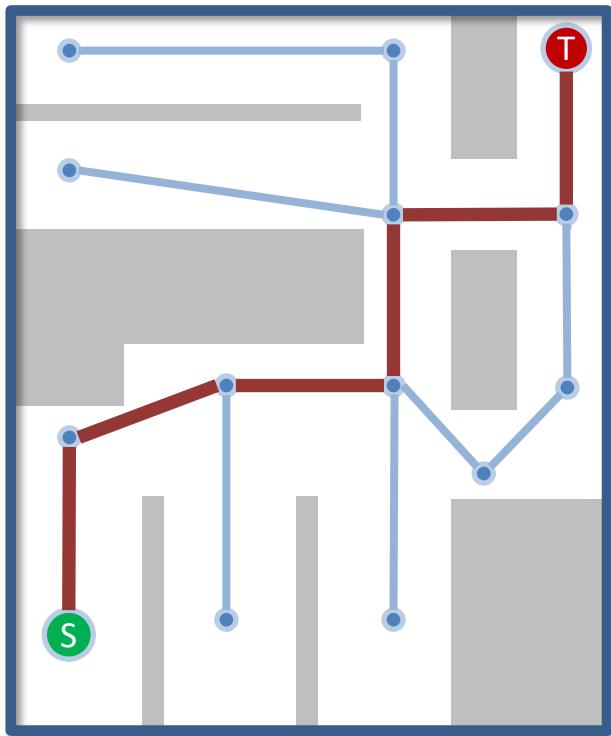


Ground truth

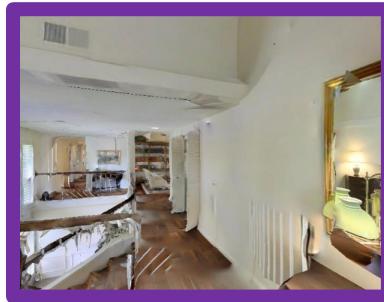


Uncertain graph

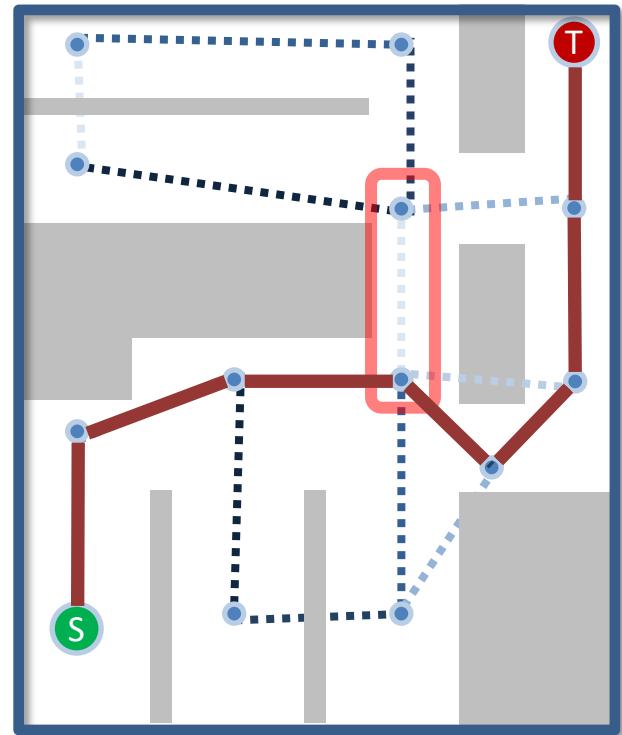
Uncertain topological maps



Ground truth

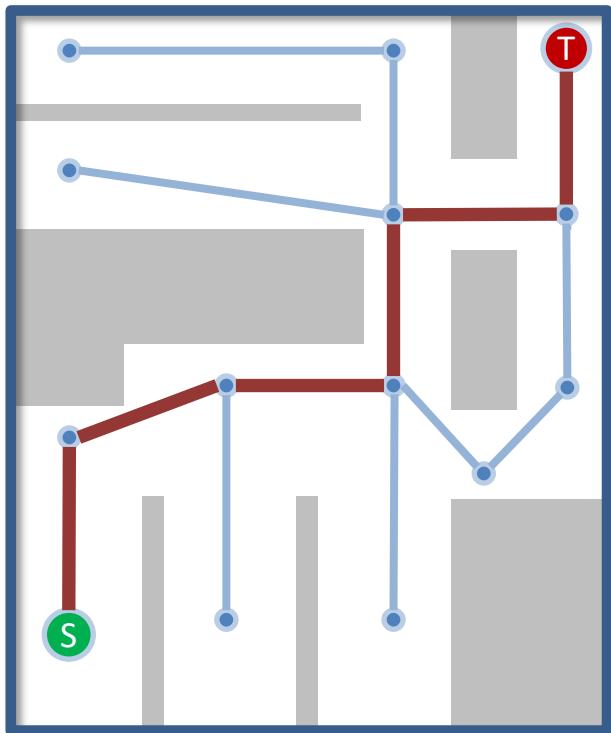


Classical planning:
• Binary connectivity
• Cannot exploit visual information

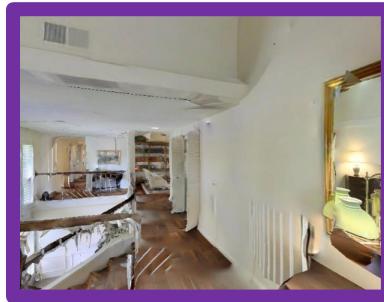


Uncertain graph

Uncertain topological maps



Ground truth

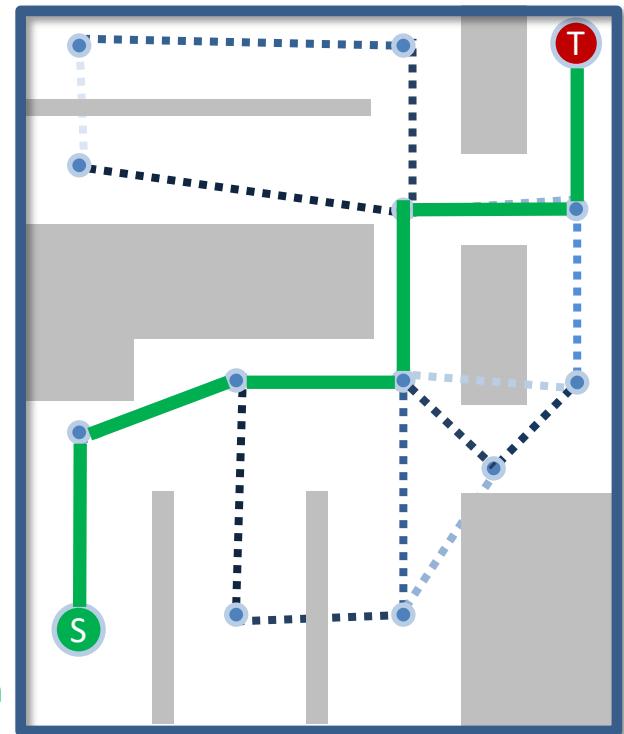


Classical planning:

- Binary connectivity
- Cannot exploit visual information

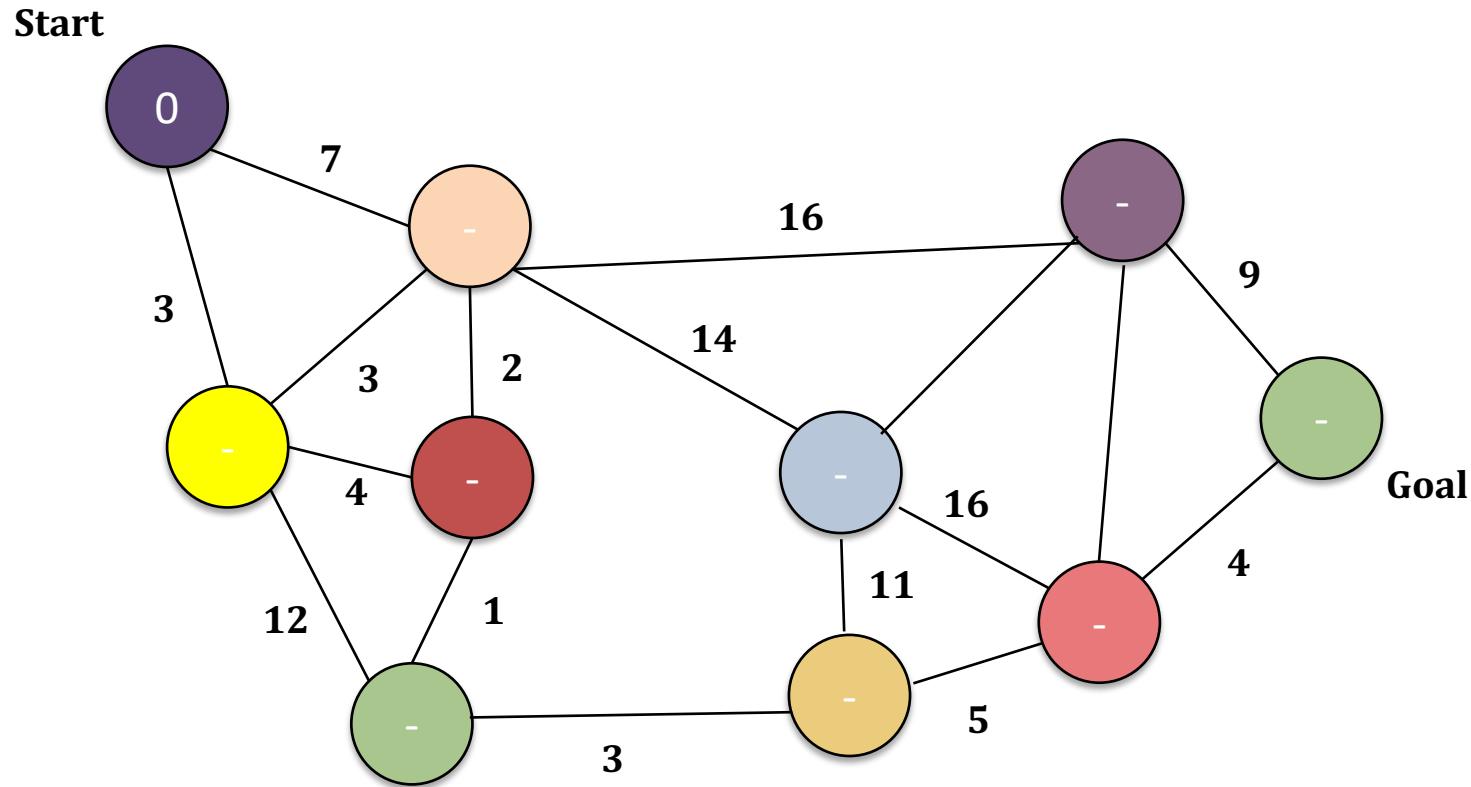
Neural planning:

- Exploits uncertain connectivity
- Uses visual features
- Uses neighbor information

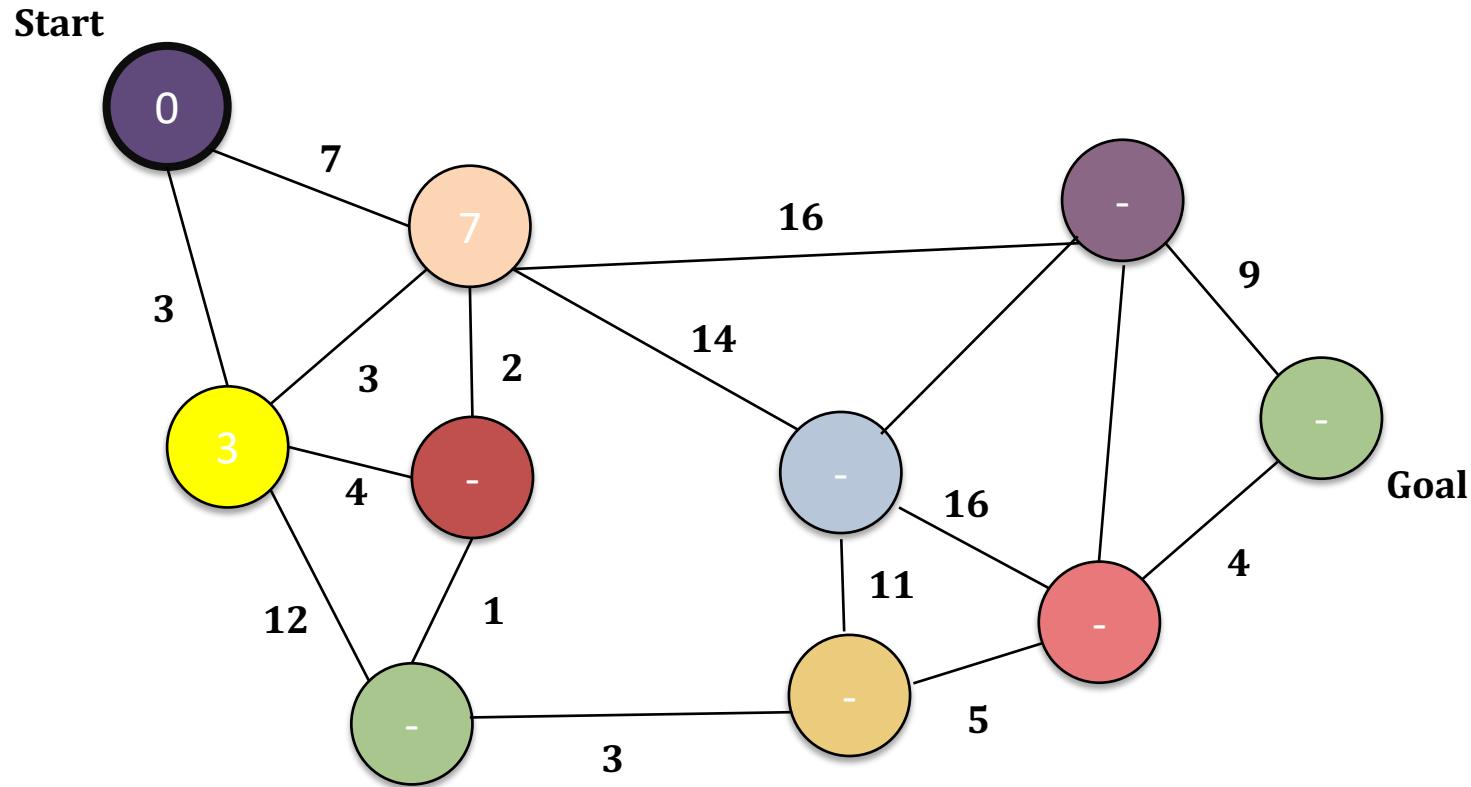


Uncertain graph

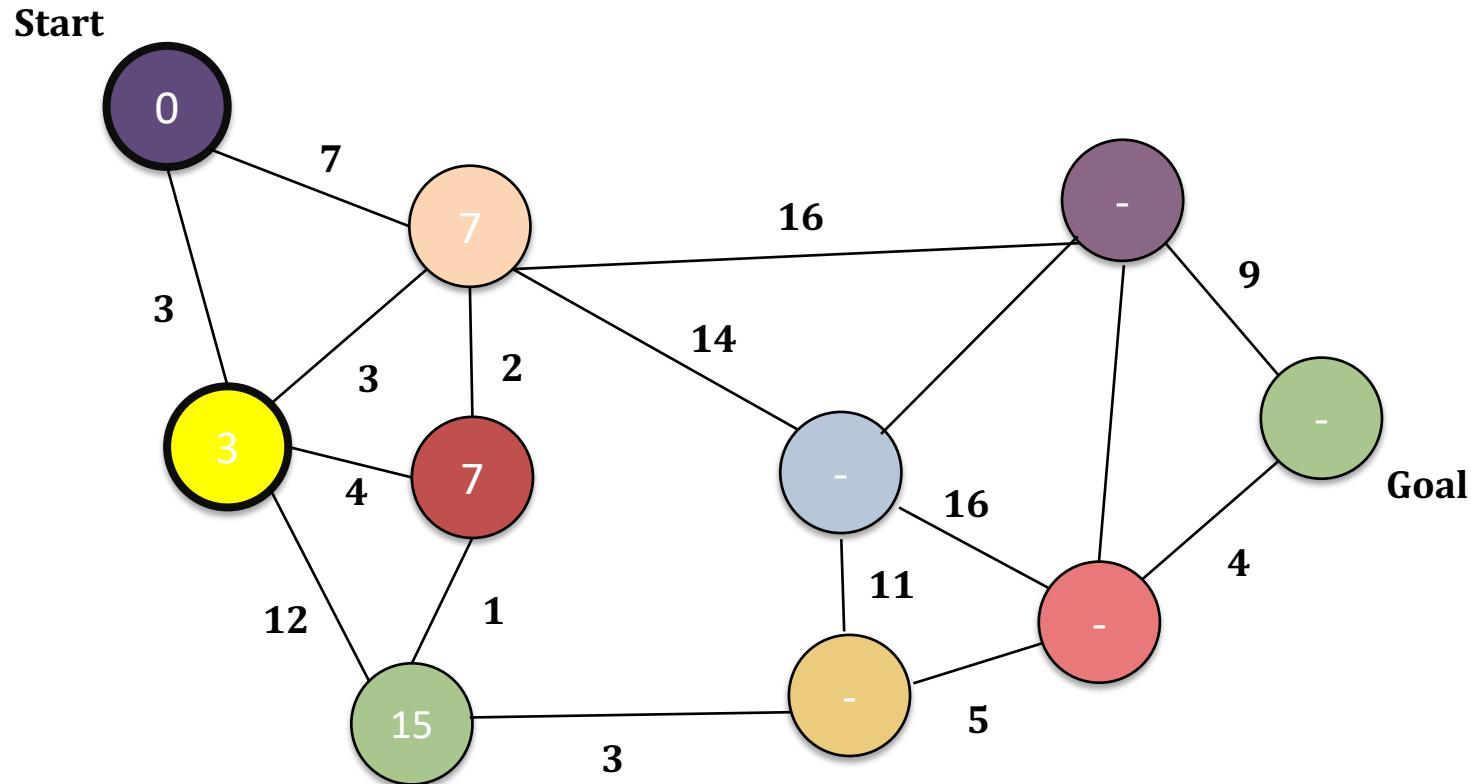
Shortest path problems: Dijkstra's algorithm



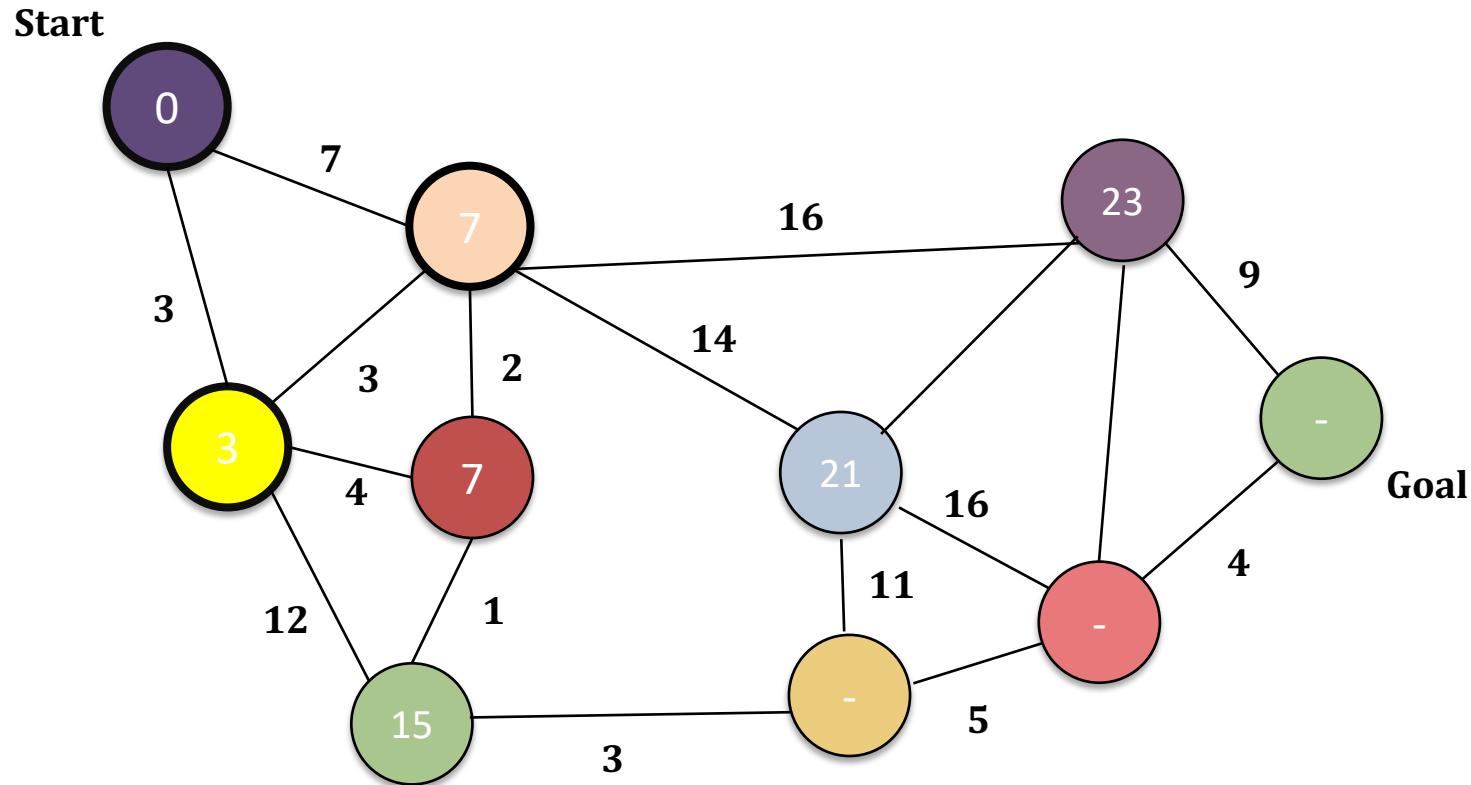
Shortest path problems: Dijkstra's algorithm



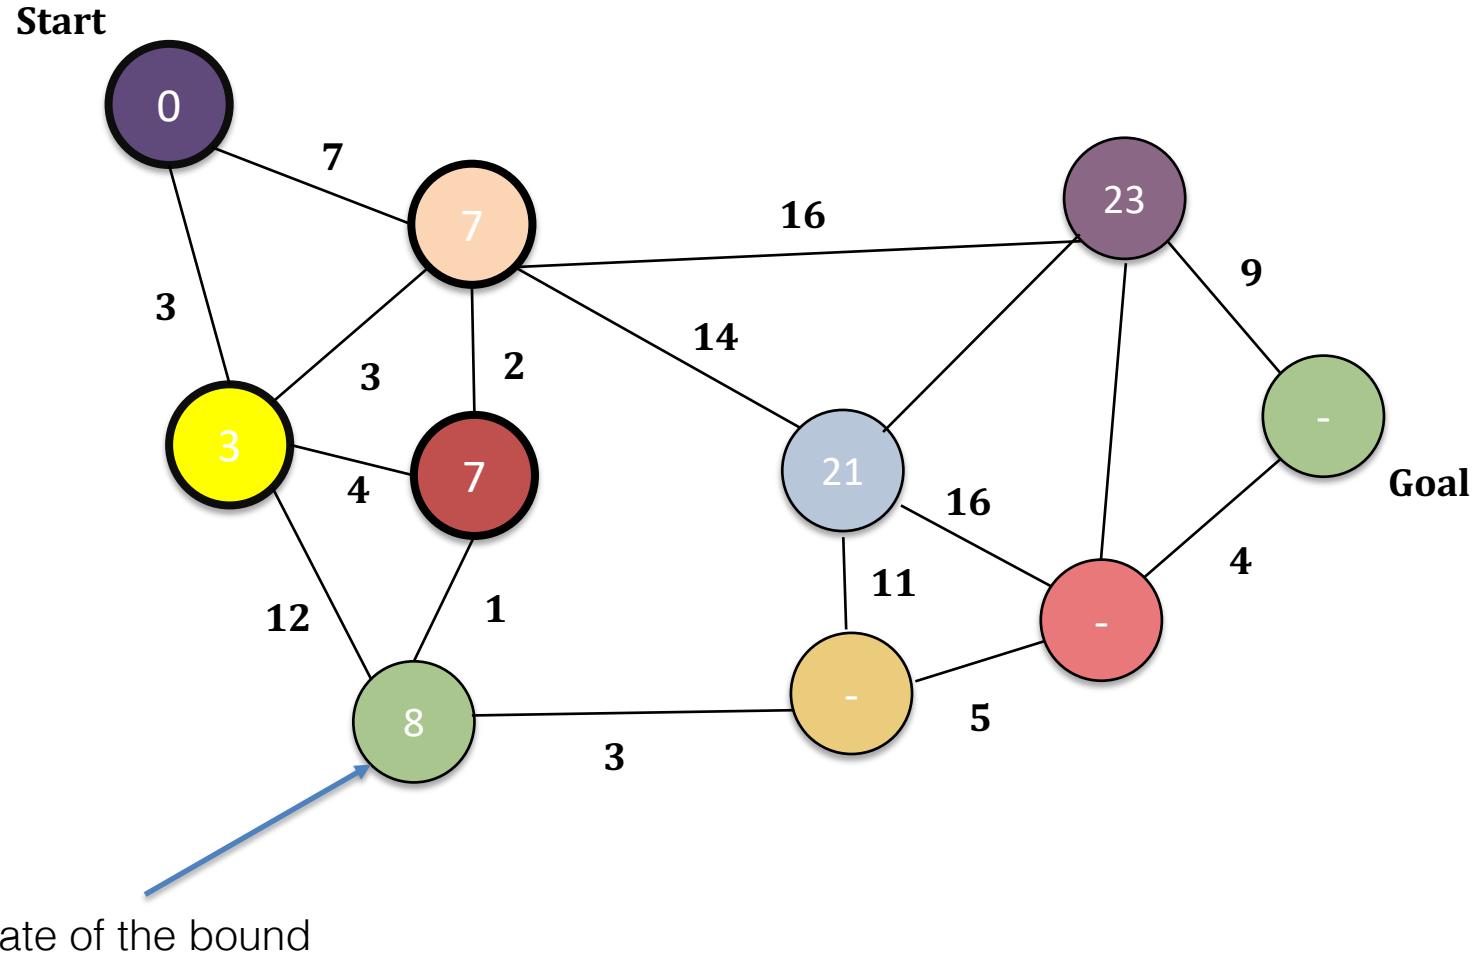
Shortest path problems: Dijkstra's algorithm



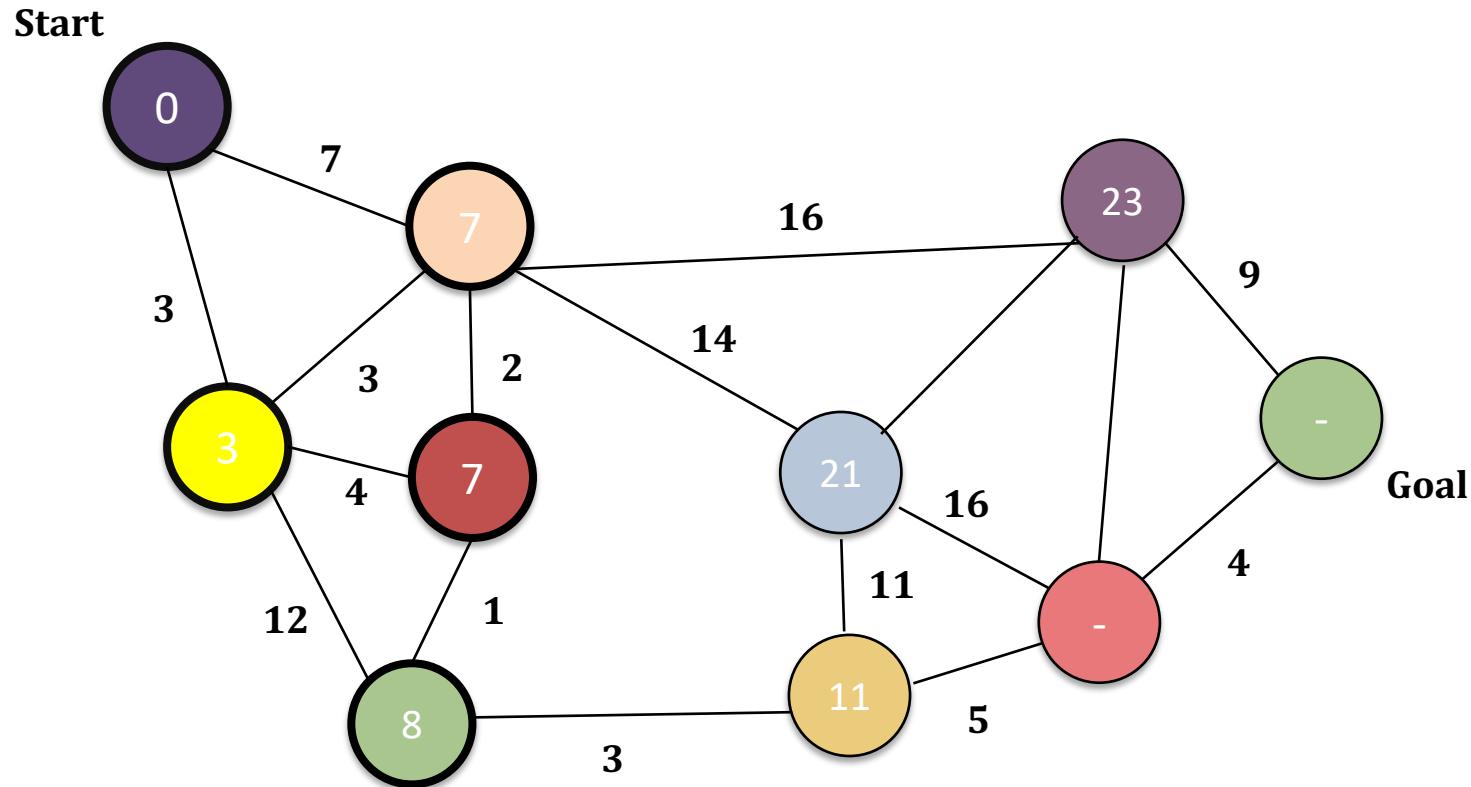
Shortest path problems: Dijkstra's algorithm



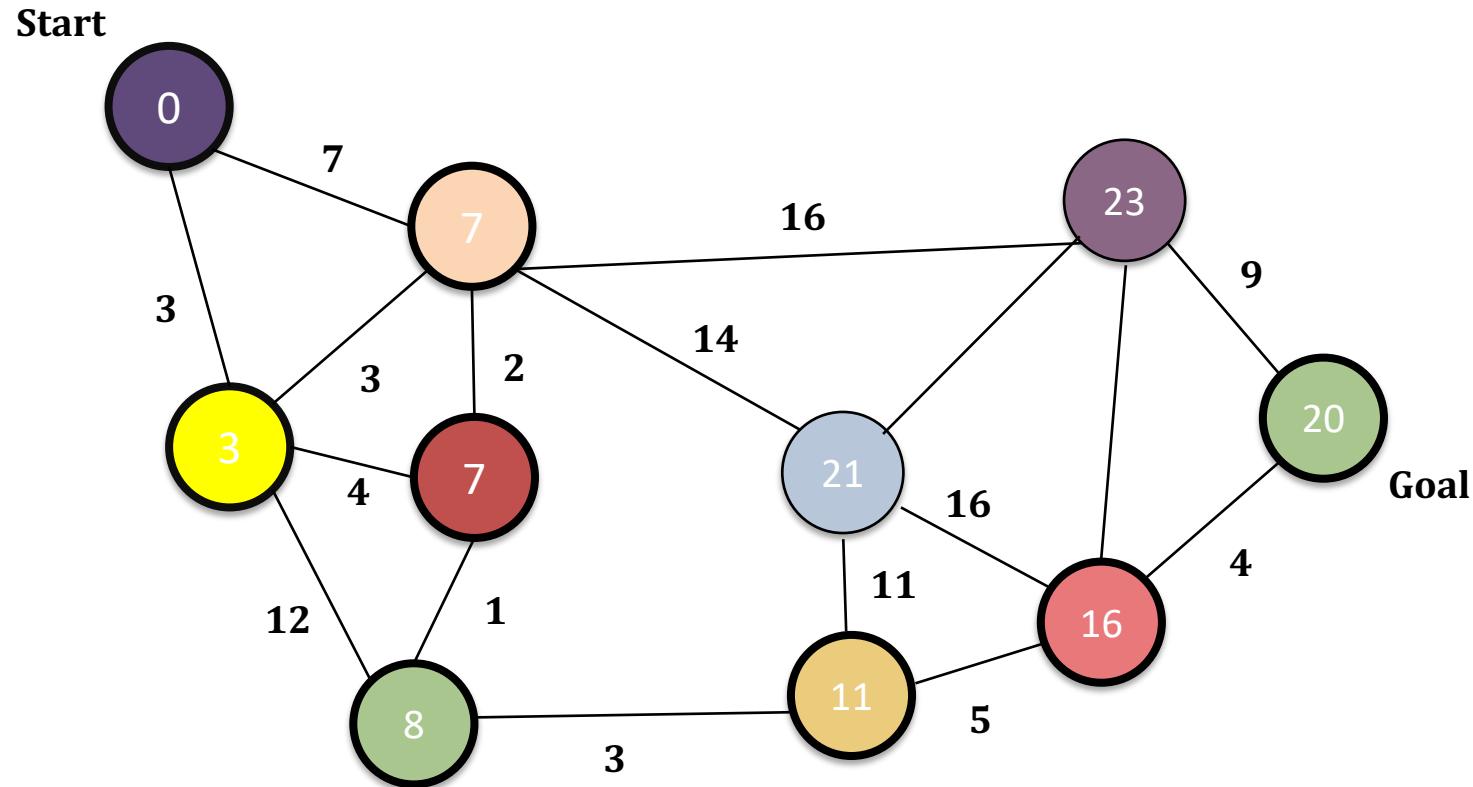
Shortest path problems: Dijkstra's algorithm



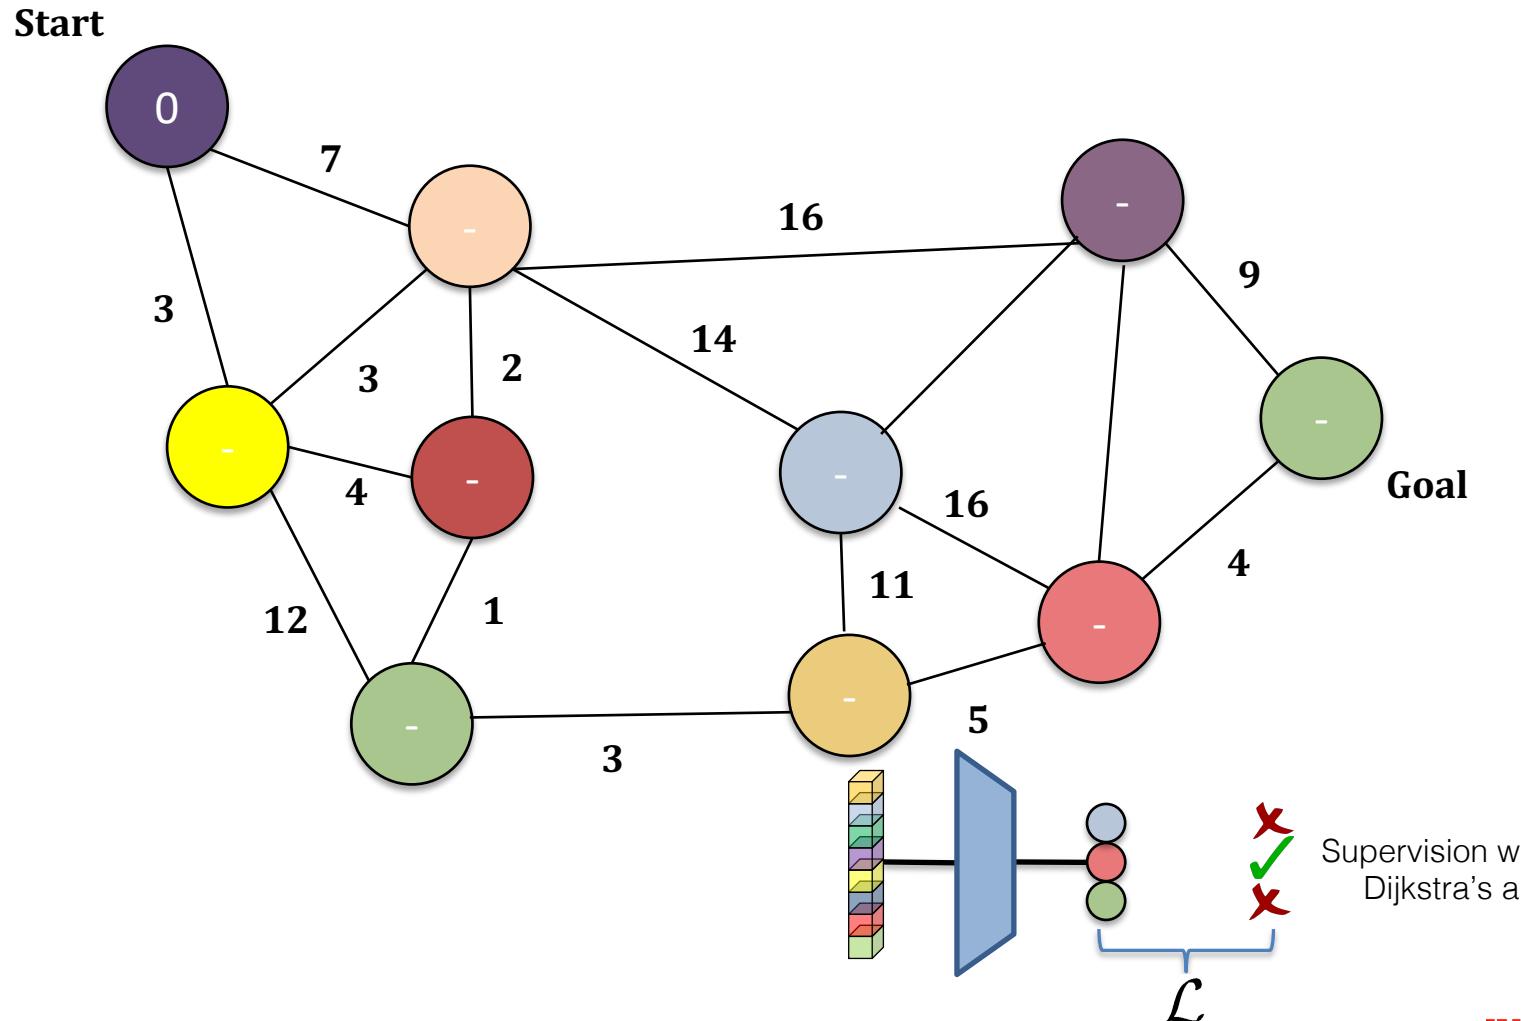
Shortest path problems: Dijkstra's algorithm

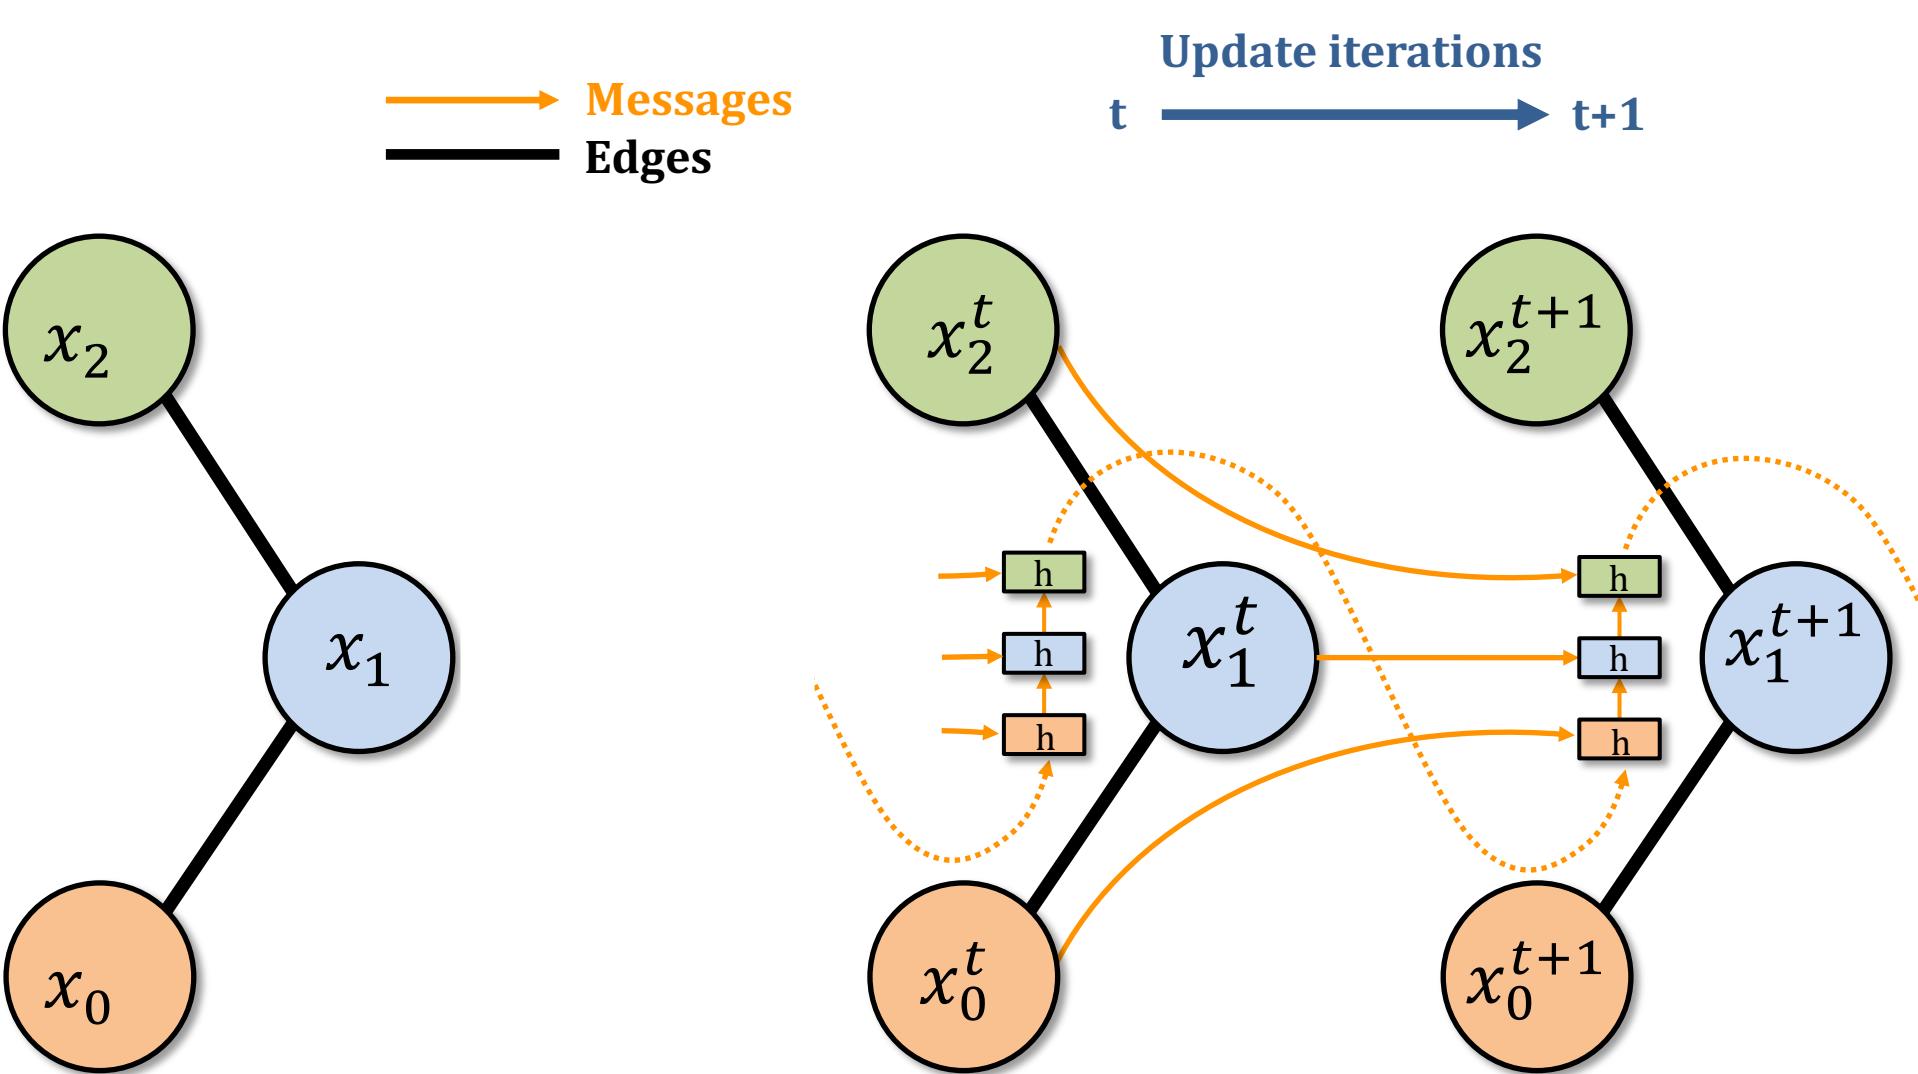


Shortest path problems: Dijkstra's algorithm



Planning as classification with Graph Neural Networks





Results: Neural planner

(a) Uncertain graphs

Method	Acc	H-SPL
Symbolic (threshold)	0.114	0.184
Symbolic (custom cost)	0.115	0.269
Neural (w/o visual)	0.251	0.468
Neural (w visual)	0.262	0.501

(b) Ground truth graphs

Method	Acc	H-SPL
Symbolic (GT)	1.00	1.00
Neural planner (GT)	0.921	0.983

We can beat optimal algorithms ...
... by cheating ... data is King!

Results: Neural planner

(a) Uncertain graphs

Method	Acc	H-SPL
Symbolic (threshold)	0.114	0.184
Symbolic (custom cost)	0.115	0.269
Neural (w/o visual)	0.251	0.468
Neural (w visual)	0.262	0.501

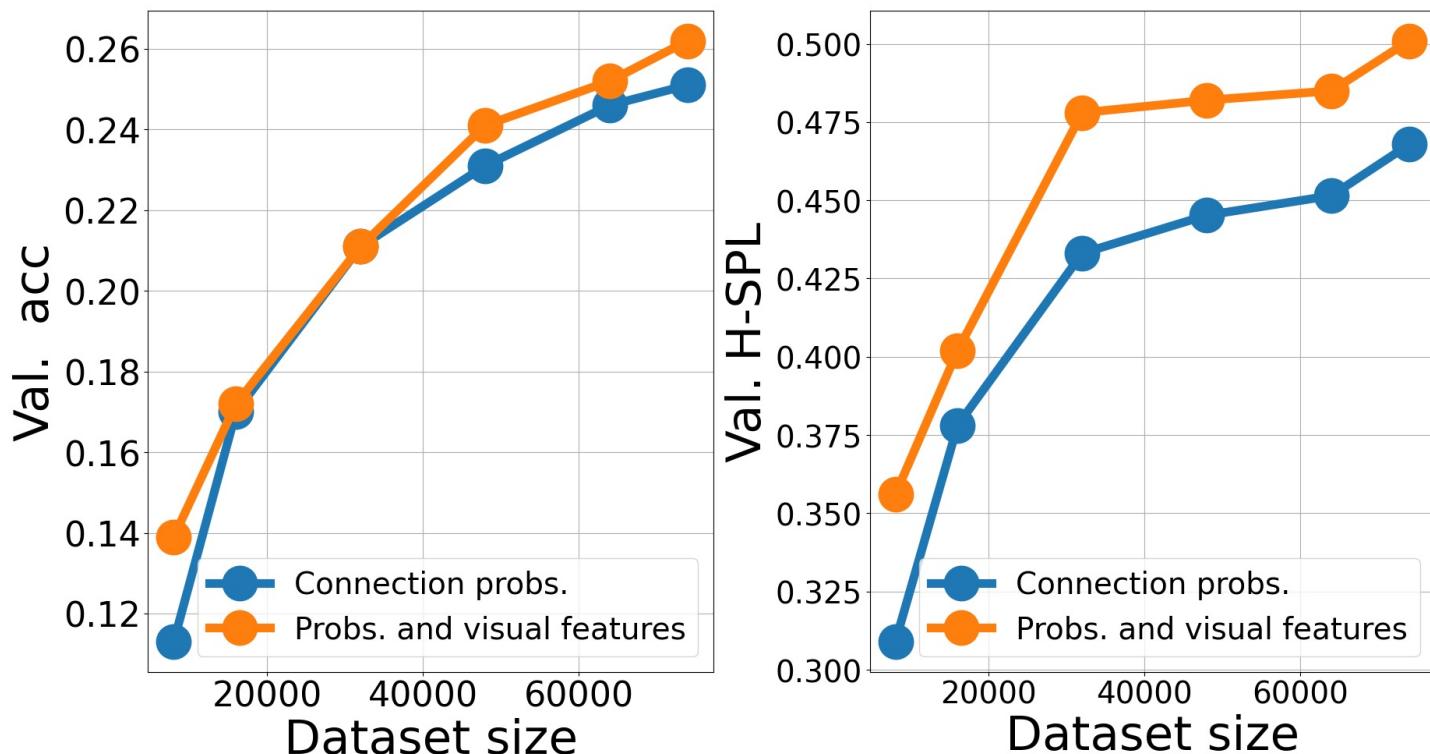


We can beat optimal algorithms ...
... by cheating ... data is King!

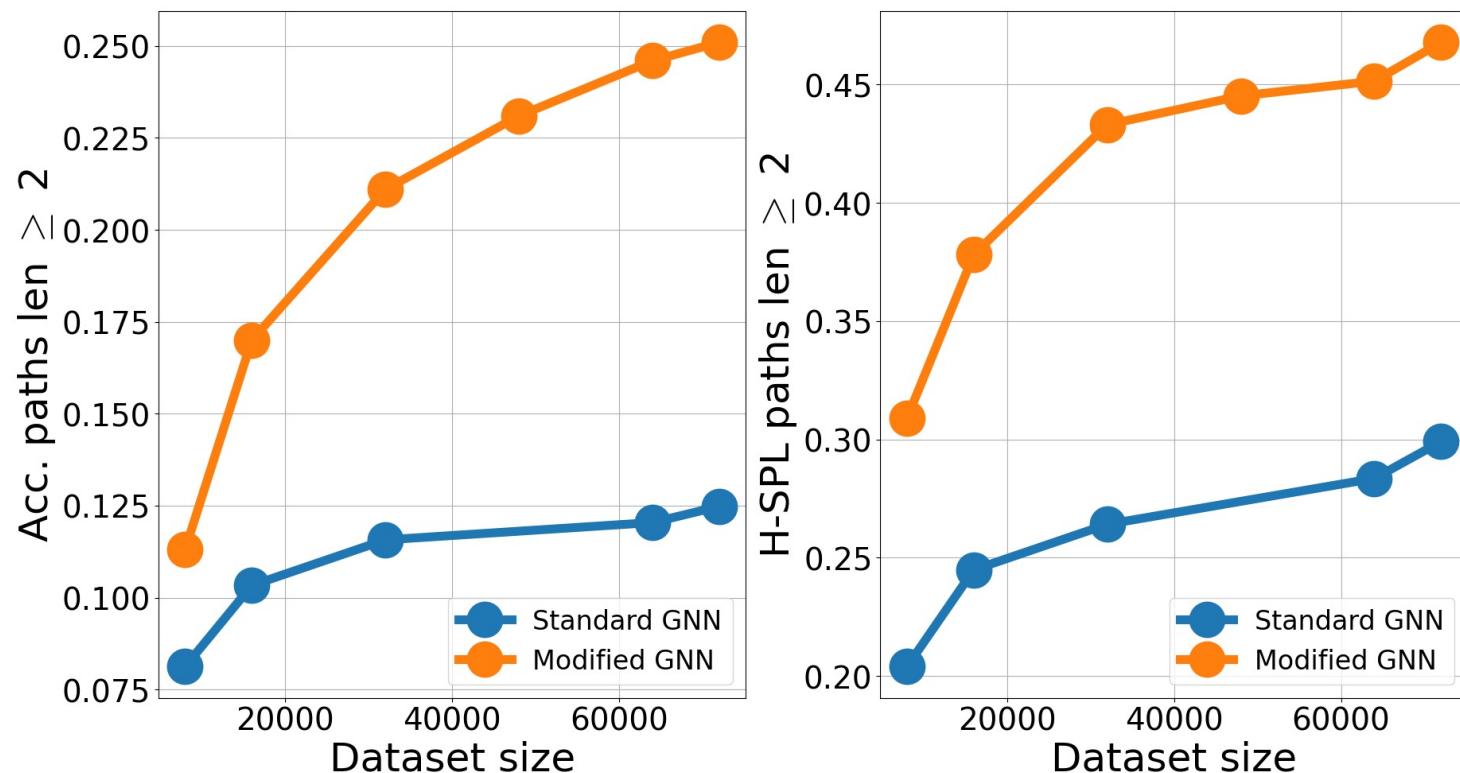
(b) Ground truth graphs

Method	Acc	H-SPL
Symbolic (GT)	1.00	1.00
Neural planner (GT)	0.921	0.983

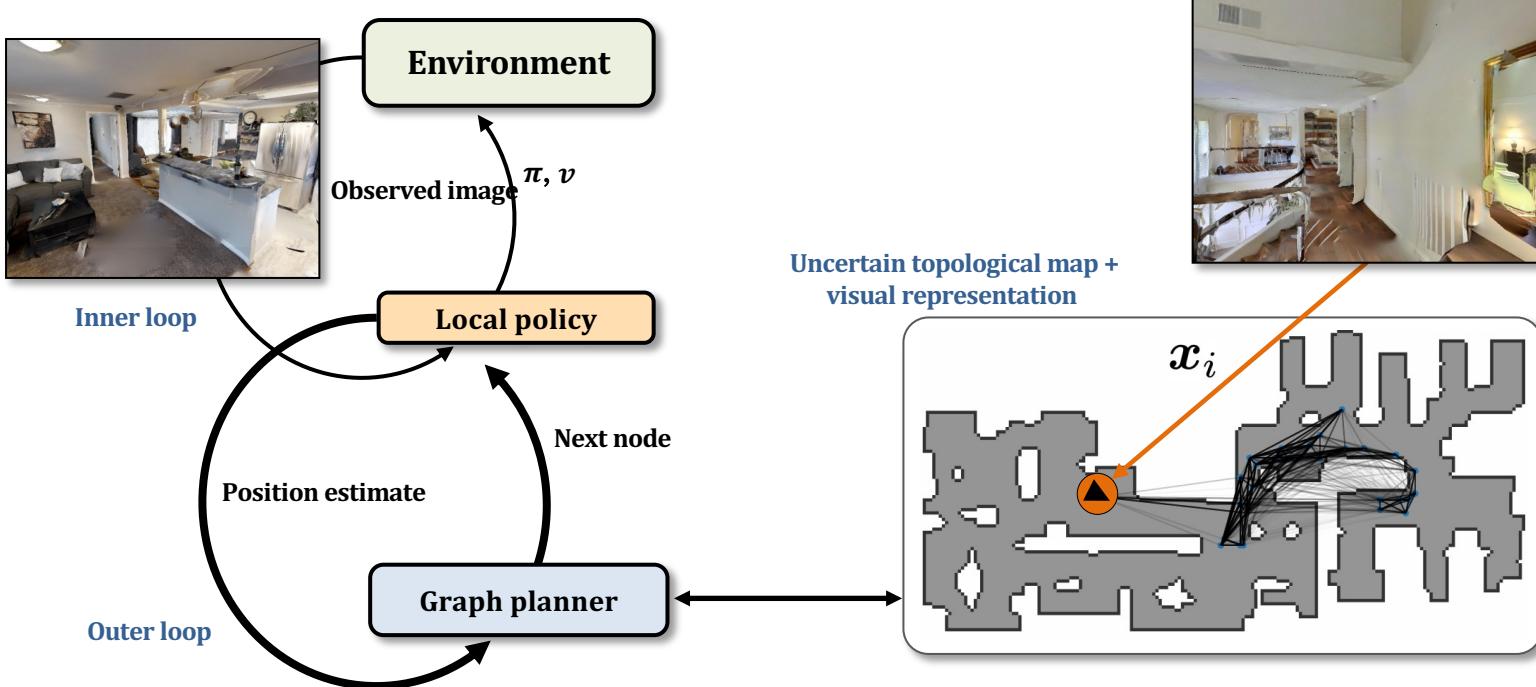
Ablation: Visual features



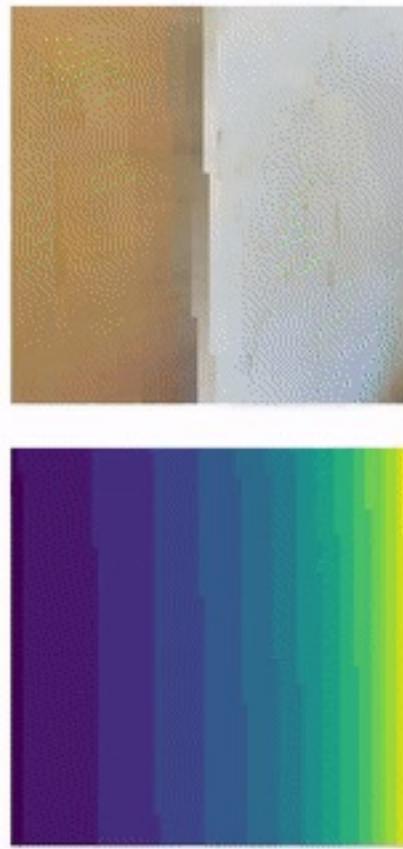
Ablation: GRU for min operation



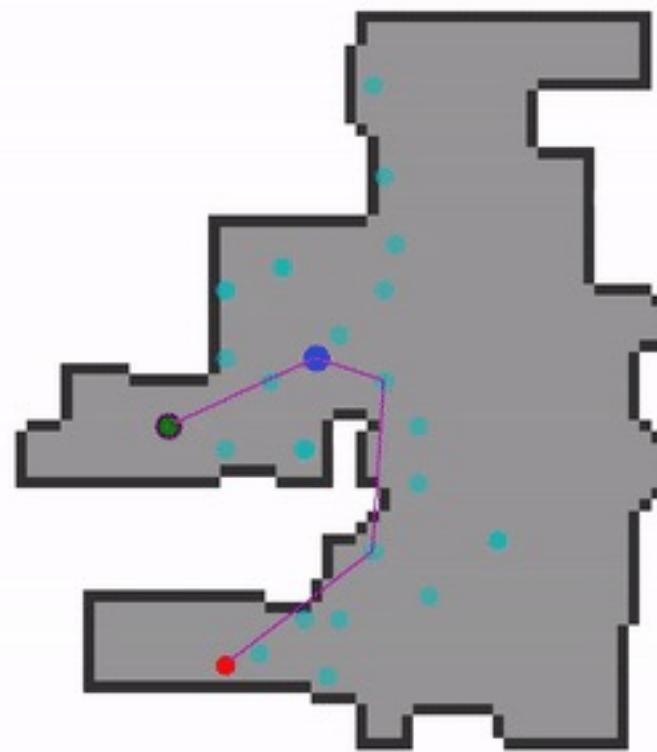
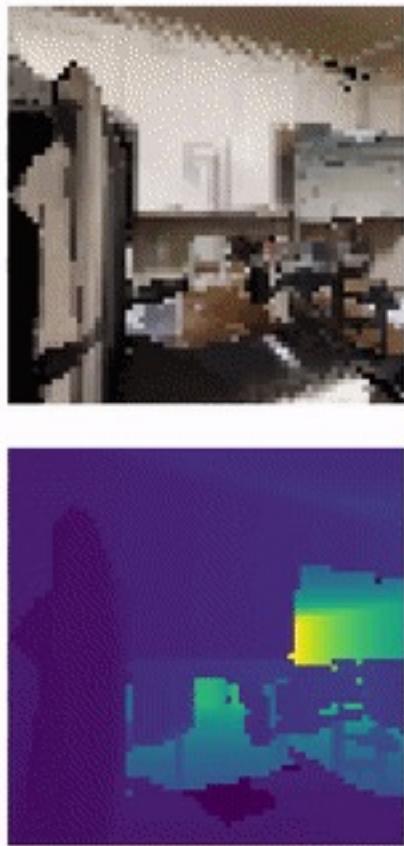
Hierarchical planning



Hierarchical planning and control



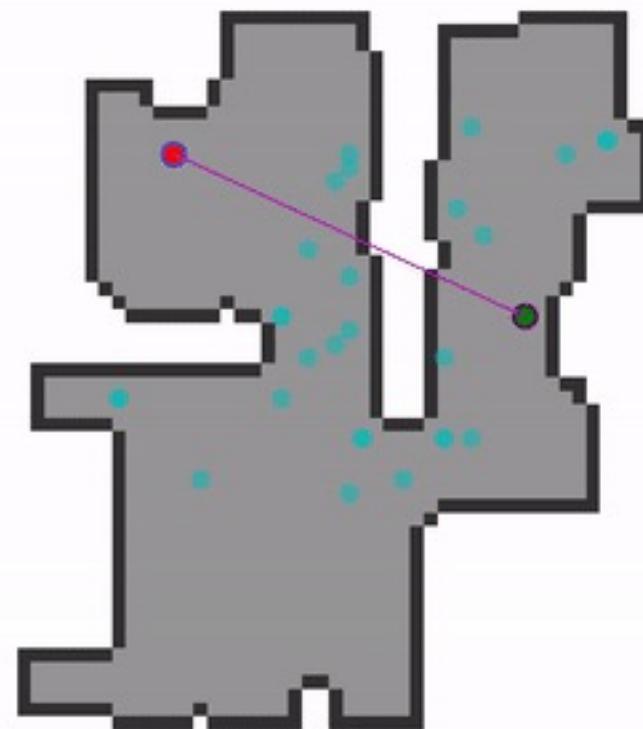
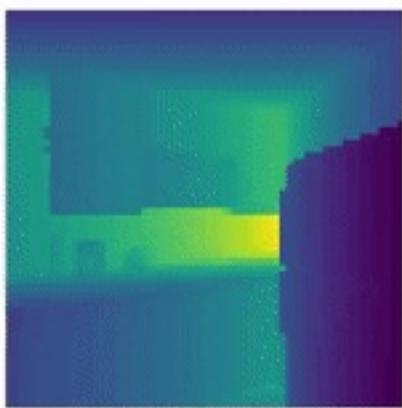
Hierarchical planning and control



Results: Hierarchical planning and control

Method: Planner + Local policy	Success rate	SPL
<i>Graph oracle (optimal point-goals, not comparable)</i>	0.963	0.882
Random	0.152	0.111
Recurrent Image-goal agent	0.548	0.248
Symbolic (threshold)	0.621	0.527
Symbolic (custom cost)	0.707	0.585
Neural planner (sampling)	0.966	0.796
Neural planner (deterministic)	0.983	0.877

Failure Case



Navigation in real environments



LABS
NAVER LABS EUROPE



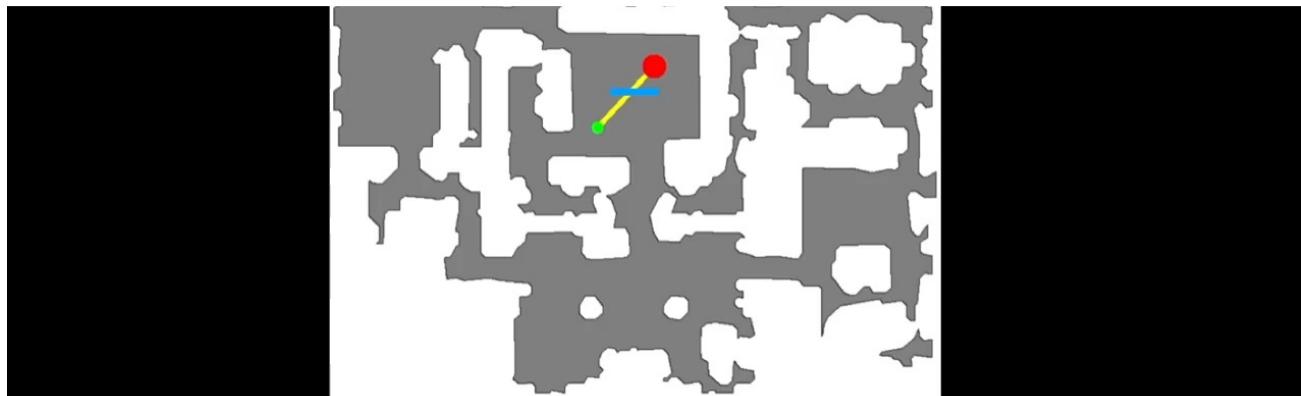
Assem
Sadek

Guillaume
Bono

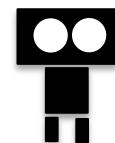
Boris
Chidlovskii

Christian
Wolf

Navigation in real environments



Language as Universal knowledge source



*I am looking for a pillow and
I see a glimpse of a
bedroom there, where
pillows are frequently found.
Let's check at this place first*



Pierre
Marza



Laetitia
Matignon



Olivier
Simonin



Christian
Wolf